

Topology Optimization of Heat Transfer Applications Operating in the Turbulent Regime

Tiebert Lefebure

Thesis submitted for the degree of
Master of Science in Mechanical
Engineering

Supervisors

Prof. Niels Horsten
Prof. Maarten Blommaert

Assessors

Prof. Elke Deckers
Olivier Renders

Assistant-supervisor

Esteban Foronda Obando

© 2026 KU Leuven – Faculty of Engineering Science

Published by Tiebert Lefebure,

Faculty of Engineering Science, Kasteelpark Arenberg 1 bus 2200, B-3001 Leuven

All rights reserved. No part of the publication may be reproduced in any form by print, photoprint, microfilm, electronic or any other means without written permission from the publisher. This publication contains the study work of a student in the context of the academic training and assessment. After this assessment no correction of the study work took place.

Preface

This thesis marks the culmination of my Master of Science in Mechanical Engineering at KU Leuven. The journey of exploring computational fluid dynamics and mathematical optimization has been both intensely challenging and deeply rewarding. Bridging the gap between theoretical mathematics and practical engineering applications required countless hours of coding, debugging, and analyzing flow fields, but seeing the optimizer successfully carve out organic, aerodynamic fluid channels made the entire process worthwhile.

I would like to express my sincere gratitude to my supervisors, Prof. Niels Horsten and Prof. Maarten Blommaert, for providing me with the opportunity to work on this fascinating topic. Their academic guidance, critical insights, and continuous support were instrumental in shaping the direction and rigor of this research. A very special thanks goes to my assistant-supervisor, Esteban Foronda Obando, for his day-to-day patience, technical expertise, and willingness to help me navigate the complex numerical hurdles within FEniCS.

Furthermore, I would like to thank the assessors, Prof. Elke Deckers and Olivier Renders, for their time and valuable feedback in evaluating this work.

Finally, I want to thank my family and friends for their unwavering encouragement, patience, and perspective throughout the course of my studies.

Tiebert Lefebure

Contents

Preface	i
Abstract	iv
List of Figures and Tables	v
List of Abbreviations and Symbols	ix
1 Introduction	1
1.1 Background and Motivation	1
1.2 Conceptual Overview of Density-Based Topology Optimization . . .	1
1.3 The Evolution of Fluid Topology Optimization	2
1.4 Objectives and Research Questions	4
1.5 Outline of the Thesis	5
2 Turbulence Modeling	6
2.1 Introduction to Fluid Dynamics and Turbulence	6
2.2 Reynolds-Averaged Navier-Stokes (RANS) Equations	7
2.3 Boundary Layer Theory and Near-Wall Treatment	8
2.4 The Spalart-Allmaras One-Equation Model	8
2.5 Evaluation of Curvature Effects and the Dean Number	9
2.6 Benchmark Studies in Ansys Fluent	12
2.7 Conclusion	14
3 Implementation of the Spalart-Allmaras Model in FEniCS	16
3.1 Introduction	16
3.2 Numerical Solution Strategy for Fluid Dynamics	17
3.3 Detailed Formulation of the Spalart-Allmaras Model	19
3.4 Variational Formulation and FEniCS Implementation	21
3.5 Conclusion	22
4 Turbulent Topology Optimization Methodology	23
4.1 Introduction	23
4.2 The Fictitious Domain Method and Brinkman Penalization	24
4.3 Penalization of the Governing Equations	25
4.4 Boundary Conditions and Turbulence Initialization	27
4.5 Sensitivity Analysis and the Frozen Turbulence Assumption	27
4.6 Optimization Problem Formulation	29
4.7 Density Filtering and Projection	31

4.8	Algorithmic Workflow and Solver Implementation	32
4.9	Conclusion	34
5	Verification of the Spalart-Allmaras Solver for Turbulent Internal Flows	36
5.1	Introduction	36
5.2	Channel Verification Benchmark	36
5.3	U-Bend Verification Benchmark	39
5.4	Conclusion	42
6	Turbulent Topology Optimization Results	43
6.1	Introduction	43
6.2	Benchmark Cases and Evaluation Metrics	44
6.3	Pipe Bend Benchmark	45
6.4	U-Bend Benchmark	51
6.5	Diffuser Benchmark	55
6.6	Conclusion	59
7	Conclusion and Future Work	61
7.1	Conclusion	61
7.2	Limitations	62
7.3	Future Work	63
A	Configuration and Grid Independence of the Spalart-Allmaras Verification Cases	66
A.1	Verification Benchmark 1: Channel Flow	66
A.2	Verification Benchmark 2: U-Bend Geometry	70
B	Configuration and Grid Independence of the Topology Optimization Benchmarks	75
B.1	Validation and Meshing Strategy	75
B.2	Shared Pipe-Bend and U-Bend Topology Optimization Conventions	76
B.3	Pipe Bend Benchmark	77
B.4	U-Bend Benchmark	79
B.5	Diffuser Benchmark	83
C	Laminar-Design Mesh Independence for Benchmark Topologies	86
C.1	Methodology and Evaluation Pipeline	86
C.2	Pipe Bend Benchmark	87
C.3	U-Bend Benchmark	88
C.4	Diffuser Benchmark	89
	Bibliography	91
	Code of Conduct	95

Abstract

As modern electronic and automotive systems grow increasingly compact and powerful, efficient thermal management has emerged as a fundamental engineering challenge. While topology optimization has proven to be a highly effective computational design tool for the fluid manifolds that drive these cooling systems, traditional methodologies rely heavily on laminar flow approximations. These laminar assumptions severely limit industrial applicability, as they fail to capture the separation, curvature effects, and turbulence-model dependence inherent to high-speed internal flows.

This thesis develops and implements a density-based topology optimization framework for turbulent internal flows as a baseline for future heat-transfer applications. The computational methodology is built upon the open-source FEniCS platform and utilizes the fictitious domain method with Brinkman momentum penalization to conceptualize solid boundaries within a continuous fluid mesh. To capture the complex physics of high-Reynolds flows, the framework integrates the one-equation Spalart-Allmaras (SA) turbulence model. The optimization process is driven by a frozen-turbulence continuous adjoint, meaning that the eddy-viscosity and wall-distance fields are updated in the forward loop but treated as fixed coefficients during sensitivity assembly.

The assessment strategy is deliberately separated into methodology and results. First, the predictive scope of the SA model for curvature-dominated internal flows is established. The FEniCS-SA implementation is then verified against Ansys Fluent on turbulent channel and U-bend benchmarks. Finally, the pipe-bend and U-bend topology optimization benchmarks from Bayat et al. are evaluated, with the pipe-bend derivative checked through a Dilgen-style finite-difference sensitivity verification before the optimized geometries are exported to Ansys Fluent for sharp-geometry validation with SA and SST $k - \omega$ turbulence models. Laminar-versus-turbulent design comparison is retained as supporting context rather than treated as CFD validation in itself.

The resulting thesis is therefore intentionally conservative in its claims. Its central question is not merely whether turbulent topologies look different from laminar ones, but whether SA-driven topology optimization produces low-loss designs whose predicted performance remains credible under independent CFD re-analysis. This establishes a more defensible and transparent baseline for future three-dimensional and conjugate heat-transfer extensions.

List of Figures and Tables

List of Figures

1.1	Iterative density-based topology optimization process: (a) initialization with a uniform field, (b) intermediate continuous distribution where $\gamma \in [0, 1]$, and (c) final discrete fluid topology where $\gamma \in \{0, 1\}$	2
1.2	Topology optimized manifolds with turbulent flow. Adapted from Dilgen et al. [16].	3
2.1	Streamwise velocity profiles in a 90° pipe bend ($R_c = 7D$, $Re_D \approx 34,000$, $De \approx 9,000$). The SA model performs relatively well, closely matching the experimental data and two-equation models for mild curvature. Reprinted from Apalowo [8].	11
2.2	Turbulence model predictions at the 90° bend exit ($X/D = 0$, $R_c = 2.0D$, $Re_D = 60,000$, $De \approx 30,000$). Due to the strong curvature, the standard Spalart-Allmaras model fails to accurately capture the near-wall velocity profile compared to two-equation closures. Adapted from Abuhatira et al. [1].	12
2.3	Flow predictions for a 180° U-bend ($R_c \approx 4.46D$, $Re_D \approx 44,700$, $De \approx 15,000$). The SA model accurately matches the higher-fidelity SST $k - \omega$ contours and directly aligns with the experimental data from Takamasa and Kondo [40].	13
2.4	Flow predictions at the 90° pipe bend exit ($R_c = 2.8D$, $Re_D \approx 43,000$, $De \approx 18,000$). The SA model over-smooths the velocity gradients near the inner wall, deviating significantly from both the SST $k - \omega$ model and the experimental data from Enayet et al. [17].	14
5.1	Verification of the 2D channel flow ($Re_H \approx 20,300$). The comparison includes qualitative velocity contours from Ansys Fluent and FEniCS, plus quantitative velocity profiles extracted at the streamwise mid-plane, $x = 0.5L$, while varying the wall-normal coordinate y	38

5.2	Verification of the 2D U-bend flow ($Re_D \approx 44,700$, $De \approx 15,000$). The comparison includes qualitative velocity contours from Ansys Fluent and FEniCS, showing separation characteristics, plus quantitative velocity profiles extracted at the bend exit across the pipe diameter $D = 0.028$ m.	41
6.1	Geometry and boundary conditions for the pipe-bend benchmark from Bayat et al. [11]. The present evaluation targets the $V_f = 0.25$ case using the SA model.	46
6.2	FEniCS-SA optimization result for the pipe-bend benchmark at $V_f = 0.25$. Final objective function value $J_p = 0.2032$ Pa.	48
6.3	Optimization convergence history for the pipe-bend benchmark. The blue curve shows the normalized pressure objective J/J_0 , while the orange curve shows the relative volume-fraction residual r_V from Eq. (6.4).	49
6.4	Laminar and turbulent topology comparison for the pipe-bend benchmark.	51
6.5	Velocity-contour comparison for the extracted pipe-bend designs under identical turbulent operating conditions (SST $k - \omega$ model).	51
6.6	Geometry and boundary conditions for the U-bend benchmark from Bayat et al. [11].	52
6.7	FEniCS-SA optimization result for the U-bend benchmark at $V_f = 0.27$. Final objective function value $J_p = [TBD]$	53
6.8	Optimization convergence histories for the U-bend benchmark.	53
6.9	Laminar and turbulent topology comparison for the U-bend benchmark.	54
6.10	SST $k - \omega$ velocity-contour comparison for the extracted U-bend designs under identical turbulent operating conditions.	55
6.11	Geometry and boundary conditions for the diffuser benchmark from Yoon [47].	56
6.12	FEniCS-SA optimization result for the diffuser benchmark at $V_f = 0.30$. Final objective function value $J_D = 6.2274 \times 10^5$ W/m.	57
6.13	Optimization convergence histories for the diffuser benchmark.	57
6.14	Laminar and turbulent topology comparison for the diffuser benchmark.	58
6.15	SST $k - \omega$ velocity-contour comparison for the extracted diffuser designs under identical turbulent operating conditions.	59
A.1	FEniCS wall-normal velocity profiles extracted at $x = 0.5L$ for the Coarse, Medium, and Fine grids in the periodic channel flow benchmark. The overlapping profiles demonstrate asymptotic spatial convergence.	69
A.2	Wall-normal velocity profiles extracted at $x = 0.5L$ for the Coarse, Medium, and Fine grids in the Ansys periodic channel flow benchmark, demonstrating spatial convergence of the baseline data.	70
A.3	Velocity profiles at the U-bend exit, extracted across the pipe diameter $D = 0.028$ m, for the Coarse, Medium, and Fine grids in FEniCS. The profiles demonstrate spatial convergence, justifying the selection of the Medium grid for the final verification.	73

A.4	Velocity profiles at the U-bend exit, extracted across the pipe diameter $D = 0.028$ m, for the Coarse, Medium, and Fine Ansys grids, demonstrating spatial convergence of the commercial baseline.	74
-----	---	----

List of Tables

5.1	Quantitative comparison of flow metrics and extracted profile discrepancies for the 2D channel benchmark.	39
5.2	Quantitative comparison of flow metrics and extracted profile discrepancies for the 2D U-bend benchmark.	42
6.1	Benchmark cases evaluated in Chapter 6. Detailed physical and numerical parameters are listed in Appendix B.	44
6.2	Coordinate sensitivity check for the frozen pipe-bend configuration. Values are taken from <code>SensitivityVerificationTable.tsv</code>	47
6.3	Directional Taylor check for the frozen pipe-bend configuration. Values are taken from <code>TaylorCheckLog.tsv</code> . The same five sampled cells as in Table 6.2 are perturbed together.	47
6.4	CFD validation metrics for the extracted pipe-bend topology.	50
6.5	Performance comparison of the extracted pipe-bend topologies under identical turbulent operating conditions.	51
6.6	CFD validation metrics for the extracted U-bend topology.	54
6.7	Performance comparison of the extracted U-bend topologies under identical turbulent operating conditions.	55
6.8	CFD validation metrics for the extracted diffuser topology.	58
6.9	Performance comparison of the extracted diffuser topologies under identical turbulent operating conditions.	59
A.1	Geometric, physical, and boundary-condition parameters for the periodic channel flow benchmark.	67
A.2	Numerical simulation parameters for the FEniCS steady Spalart-Allmaras channel benchmark.	67
A.3	Wall-resolved channel mesh hierarchy used for the FEniCS grid-independence study. The first-layer height is kept fixed to maintain $y^+ \approx 1$, while the streamwise resolution, wall-normal resolution, and inflation-layer resolution are refined.	68
A.4	Wall-resolved channel mesh hierarchy used for the Ansys grid-independence study.	69
A.5	Geometric, physical, and boundary-condition parameters for the U-bend benchmark.	71
A.6	Numerical simulation parameters for the highly convective FEniCS U-bend benchmark.	72

A.7	Wall-resolved U-bend mesh hierarchy used for the FEniCS grid-independence study. The first-layer height is kept fixed to maintain $y^+ \approx 1$, while the bulk element size and inflation-layer resolution are refined.	72
A.8	Wall-resolved U-bend mesh hierarchy used for the Ansys grid-independence study.	74
B.1	Shared wall-distance and SA topology-penalty conventions for the pipe-bend and U-bend benchmarks.	76
B.2	Geometric, physical, and boundary-condition parameters for the pipe-bend benchmark.	77
B.3	Numerical FEniCS solver, topology-optimization, and sensitivity-verification parameters for the pipe-bend benchmark.	78
B.4	Ansys Fluent mesh-independence table for the extracted turbulent pipe-bend benchmark geometry.	79
B.5	Geometric, physical, and boundary-condition parameters for the U-bend benchmark.	80
B.6	Numerical FEniCS solver and topology-optimization parameters for the U-bend benchmark.	81
B.7	Ansys Fluent mesh-independence table for the extracted turbulent U-bend benchmark geometry.	82
B.8	Geometric, physical, and boundary-condition parameters for the diffuser benchmark.	83
B.9	Numerical FEniCS solver and topology-optimization parameters for the diffuser benchmark.	84
B.10	Ansys Fluent mesh-independence table for the extracted turbulent diffuser benchmark geometry.	85
C.1	Ansys mesh-independence table for the extracted laminar pipe-bend design.	88
C.2	Ansys mesh-independence table for the extracted laminar U-bend design.	89
C.3	Ansys mesh-independence table for the extracted laminar diffuser design.	90

List of Abbreviations and Symbols

Abbreviations

CAD	Computer-Aided Design
CEV	Constant Eddy Viscosity
CFD	Computational Fluid Dynamics
CG1	Continuous Galerkin (Linear Basis Functions)
CHT	Conjugate Heat Transfer
DG0	Discontinuous Galerkin (Piecewise Constant Elements)
DNS	Direct Numerical Simulation
FEM	Finite Element Method
FVM	Finite Volume Method
IPCS	Incremental Pressure Correction Scheme
MMA	Method of Moving Asymptotes
PDE	Partial Differential Equation
RAMP	Rational Approximation of Material Properties
RANS	Reynolds-Averaged Navier-Stokes
SA	Spalart-Allmaras
SIMP	Solid Isotropic Material with Penalization
SST $k - \omega$	Shear Stress Transport $k - \omega$ model
TO	Topology Optimization
UFL	Unified Form Language

Symbols

Fluid Dynamics & Turbulence

d, y	Distance to the nearest wall / Hydraulic diameter [m]
De	Dean number [-]
k	Turbulent kinetic energy [m^2/s^2]
$p, \bar{p}, p'$	Instantaneous, mean, and fluctuating static pressure [Pa]
ΔP	Macroscopic total pressure drop [Pa]
R_c	Radius of curvature [m]
Re	Bulk Reynolds number [-]
$S, \tilde{S}$	Scalar vorticity magnitude and modified SA vorticity magnitude [$1/\text{s}$]
t	Time / Pseudo-time [s]
$\mathbf{u}, \bar{u}_i$	Velocity vector / Mean velocity component [m/s]
u'_i	Fluctuating velocity component [m/s]
u_τ	Friction velocity [m/s]
y^+	Non-dimensional wall distance [-]
μ, μ_t	Dynamic and turbulent eddy viscosity [$\text{Pa}\cdot\text{s}$]
ν, ν_t	Kinematic and turbulent eddy kinematic viscosity [m^2/s]
$\tilde{\nu}$	Spalart-Allmaras working variable [m^2/s]
ρ	Fluid density [kg/m^3]
Ω_{ij}	Vorticity tensor [$1/\text{s}$]
κ	von Kármán constant [-]

Numerical Implementation (FEniCS)

G, G_0	Relaxed Eikonal field variable and boundary value [$1/\text{m}$]
h, h_{avg}	Finite element cell diameter and average cell size [m]
$\ \cdot\ _{L_2}$	L_2 (Euclidean) norm [-]
$\llbracket \cdot \rrbracket$	Jump operator across an internal element facet [-]
$\mathbf{n}$	Outward-pointing normal vector [-]
v, ξ	Test functions for pressure and turbulence equations [-]
$\alpha_u, \alpha_{\tilde{\nu}}$	Relaxation factors for velocity and turbulence [-]
Δt	Pseudo-time step [s]
ϵ_{tol}	Convergence tolerance [-]
ϕ	General scalar field variable [-]
σ_w	Eikonal relaxation constant [-]

x

Topology Optimization

$\alpha(\gamma)$	Continuous Brinkman penalization term / inverse permeability
$\alpha_{solid}, \alpha_{fluid}$	Maximum solid and minimum fluid Brinkman penalties
α_{dg}	Artificial diffusion scaling factor for DG0 filter [-]
β	Heaviside projection steepness parameter [-]
γ, γ_{raw}	Continuous and raw design variable (1 for fluid, 0 for solid) [-]
$\tilde{\gamma}, \bar{\gamma}$	Filtered and Heaviside-projected design variables [-]
η	Heaviside projection threshold [-]
J	Generic objective function used in the optimization problem [problem dependent]
J_D	Dissipation objective, evaluated from viscous dissipation [W, or W/m for two-dimensional unit-depth cases]
J_p	Average inlet-pressure objective [Pa]
q	RAMP penalization tuning parameter [-]
r	Continuous Helmholtz filter radius [m]
$\mathbf{s}, \mathbf{s}^*$	Monolithic forward state and continuous adjoint state vectors
V_f, V_{domain}	Physical fluid volume and total design domain volume [m ³]
V_{max}	Maximum allowable fluid volume fraction constraint [-]
$\Omega, \partial\Omega, \Gamma_{int}$	Computational domain, domain boundary, and internal element facets

Chapter 1

Introduction

1.1 Background and Motivation

With the continued miniaturization of components in the electronics and automotive industries, thermal management has become a critical bottleneck in the development of new technologies. While the heat sinks themselves are vital, the efficiency of any cooling system relies heavily on the fluid manifolds that deliver the coolant. If the flow distribution is uneven or if pumping the fluid requires excessive power due to massive pressure drops, the overall efficiency of the system severely degrades. Addressing this fluid delivery problem from an engineering perspective is a challenging task due to the complex aerodynamic behaviors that govern high-speed flow networks.

Innovative design strategies based on mathematical optimization have made it possible to improve the performance of these devices while capturing the complexity of the underlying physical phenomena. Among these strategies, topology optimization has emerged as a highly effective computational design tool for advanced thermal management [18]. For instance, the geometry of a complex microfluidic channel network can be automatically optimized to distribute coolant with specified flow rates without relying on human intuition [50]. Furthermore, the ultimate frontier of this technology lies in the design of conjugate and two-fluid heat exchangers, where the algorithm must simultaneously maximize heat transfer while strictly limiting the fluid pumping power [25]. However, before multi-physics heat transfer can be reliably optimized at an industrial scale, the fundamental algorithms governing the turbulent fluid dynamics must first be stabilized and validated.

1.2 Conceptual Overview of Density-Based Topology Optimization

To understand the challenges of applying optimization to turbulent fluids, it is first necessary to distinguish topology optimization from traditional design methodologies. Historically, computational design relied on size optimization (altering the thickness of predefined components) or shape optimization (moving the boundary

nodes of a predefined geometry). While effective, these methods are inherently limited by their starting topology. A shape optimization algorithm cannot spontaneously create a new hole or split a single fluid channel into two.

Density-based topology optimization, pioneered by Bendsøe and Kikuchi in 1988 for solid mechanics [12], approaches the problem fundamentally differently. Instead of moving boundaries, the design space (the computational domain) is discretized into a fixed grid of finite elements. The algorithm then determines the optimal distribution of material across this grid. In fluid problems, this is achieved by assigning a continuous design variable, $\gamma \in [0, 1]$, to every element in the domain. In the specific convention adopted for this thesis, a value of $\gamma = 1$ represents a pure fluid channel (e.g., the cooling air or water), while $\gamma = 0$ represents a solid, impermeable structure (e.g., a manifold wall).

However, because gradient-based optimizers require continuous and differentiable functions, this discrete binary choice is mathematically relaxed. The design variable is allowed to take intermediate "gray" values between 0 and 1. To give physical meaning to these intermediate states, the fluid equations are modified using a fictitious porous medium approach, where intermediate densities correspond to varying levels of flow permeability. The optimization algorithm iteratively updates the γ field (as conceptualized in Figure 1.1) to minimize a specific objective function (e.g., dissipated power) subject to physical constraints (e.g., a maximum allowable fluid volume fraction). The choice of the initial value for γ is a critical parameter for solver stability. A uniform initial guess of $\gamma = 0.5$ is standard practice to ensure an unbiased algorithmic search across the design space.




FIGURE 1.1: Iterative density-based topology optimization process: (a) initialization with a uniform field, (b) intermediate continuous distribution where $\gamma \in [0, 1]$, and (c) final discrete fluid topology where $\gamma \in \{0, 1\}$.

1.3 The Evolution of Fluid Topology Optimization

While density-based topology optimization became a standard tool in structural engineering in the late 1990s, its application to fluid dynamics is relatively recent. The optimization of fluid manifolds remains an active and highly challenging area of research [18]. The genesis of fluid topology optimization is widely attributed

to Borrvall and Petersson (2003), who successfully applied the technique to minimize power dissipation in creeping Stokes flow [13]. This was quickly followed by Gersborg-Hansen et al. (2005), who extended the framework to steady, laminar Navier-Stokes flows [20]. Over the past decade, the literature surrounding laminar flow optimization expanded extensively, successfully adapting the methodology to standard incompressible Navier-Stokes equations [15, 3, 50].

However, the assumption of laminar flow significantly simplifies the governing equations and the subsequent derivation of the adjoint sensitivities required for gradient-based optimization. Unfortunately, this assumption prevents the application of these methodologies to most industrial and high-performance applications, where flow velocities are high and turbulence dominates the flow physics. Only recently has the computational mechanics community begun to tackle the severe non-linearities and numerical instabilities associated with optimizing geometries directly in the fully turbulent regime [47, 16, 48].

Extending topology optimization to the turbulent regime is highly non-trivial. It requires coupling optimization algorithms with complex Reynolds-Averaged Navier-Stokes (RANS) turbulence models. These models introduce steep velocity gradients, highly non-linear behaviors, and severe numerical instabilities during the iterative optimization process. To mitigate this, gradient-based turbulent optimizations frequently rely on the "frozen turbulence" assumption, where the turbulent viscosity field is held constant during the sensitivity analysis to avoid differentiating the complex turbulence transport equations. While this assumption provides necessary solver stability and has driven massive success in industrial shape optimization for automotive aerodynamics [30], advancing toward fully coupled turbulent sensitivities in topology optimization remains an ongoing and highly complex mathematical challenge.




FIGURE 1.2: Topology optimized manifolds with turbulent flow. Adapted from Dilgen et al. [16].

Pioneering works by Yoon (2016) [47] and Dilgen et al. (2018) [16] have pushed the boundaries of these turbulent frameworks. As seen in Figure 1.2, integrating advanced turbulence models allows the optimizer to generate highly complex, or-

ganic fluid manifolds specifically tailored to manipulate and guide turbulent flow distributions efficiently.

1.4 Objectives and Research Questions

The main objective of this Master’s thesis is to extend the application of topology optimization to the design of fluid manifolds operating strictly in the turbulent regime. To achieve this goal, this work investigates the coupling of fluid dynamics and turbulence modeling within an automated mathematical optimization framework. By strictly isolating the fluid mechanics from complex thermodynamic calculations, this thesis bounds its scope to establish a highly stable and mathematically rigorous baseline for pure aerodynamic manifold design.

Rather than relying solely on black-box commercial Computational Fluid Dynamics (CFD) software, the core of this thesis involves implementing the framework using **FEniCS**, an open-source computing platform for solving partial differential equations [4]. FEniCS, combined with high-level discrete adjoint methods (e.g., `dolfin-adjoint` [5]), provides the mathematical flexibility required to implement, verify, and integrate the one-equation Spalart-Allmaras (SA) turbulence model into a fluid-design loop. For full transparency and reproducibility, the complete FEniCS source code developed for this thesis is open-source and publicly available at: <https://github.com/TiebertLefebure/TurbulentT0.git>.

At this stage, it is important to distinguish two different comparisons that appear later in the thesis. A comparison between laminar-optimized and turbulent-optimized layouts is useful as supporting design context, because it shows how inertia and turbulence alter the preferred topology. However, this comparison is not treated as the principal CFD validation step. The decisive assessment instead consists of re-analyzing the final SA-optimized geometries in Ansys Fluent, first with the same SA model and then with alternative RANS closures.

Based on this methodology, this thesis seeks to answer the following specific research questions:

1. **How accurately does the Spalart-Allmaras turbulence model predict curvature-dominated internal turbulent flows, and where do its physical limitations appear?**
2. **Can a custom standard-Galerkin FEniCS implementation of the Spalart-Allmaras turbulence model converge robustly and achieve quantitative agreement with established commercial CFD solvers?**
3. **To what extent do SA-driven topology optimizations produce final designs whose predicted performance remains consistent under independent Ansys re-analysis, both with the same SA model and with alternative RANS closures?**

1.5 Outline of the Thesis

The remainder of this thesis is structured into the following six chapters:

- **Chapter 2 (Turbulence Modeling):** Provides the theoretical background on fluid dynamics and turbulence. It discusses the Navier-Stokes equations, the principles of RANS modeling, and the physical characteristics of the Spalart-Allmaras model.
- **Chapter 3 (Implementation of the Spalart-Allmaras Model in FEniCS):** Details the numerical implementation of the chosen turbulence model within the FEniCS framework, including fractional-step solvers, standard Galerkin discretization, Picard coupling, and wall-distance computation.
- **Chapter 4 (Turbulent Topology Optimization Methodology):** Introduces the density-based optimization framework. It details the Brinkman momentum penalization, the power-law SA and wall-distance penalties, the reciprocal penalized wall-distance equation, the pressure objective, and the rationale for adopting the frozen turbulence assumption as the baseline design engine.
- **Chapter 5 (Verification of the Spalart-Allmaras Solver for Turbulent Internal Flows):** Verifies the standalone FEniCS-SA forward solver against Ansys Fluent on turbulent channel-flow and U-bend benchmarks.
- **Chapter 6 (Turbulent Topology Optimization Results):** Presents the turbulent topology optimization results for the pipe-bend and U-bend benchmarks from Bayat et al., documents the pipe-bend sensitivity verification, exports the final designs to Ansys Fluent, and evaluates the designs with both SA and SST $k - \omega$ turbulence models.
- **Chapter 7 (Conclusion and Future Work):** Summarizes the main findings, answers the research questions conservatively, and outlines the remaining limitations, with particular emphasis on turbulence-model dependence, geometry extraction, and future multi-physics heat-transfer applications.

Chapter 2

Turbulence Modeling

2.1 Introduction to Fluid Dynamics and Turbulence

Fluid dynamics is governed by the Navier-Stokes equations, which describe the conservation of mass and momentum for a viscous continuum fluid [45]. In this introductory text, these equations are presented using index notation (Einstein summation convention) to clearly illustrate the individual tensor components. In this convention, the indices i and j represent the three spatial dimensions (1, 2, and 3, corresponding to the x , y , and z axes). Consequently, x_i denotes the Cartesian coordinates, and u_i represents the corresponding components of the velocity vector. A repeated index within a single term implies a mathematical summation over all three spatial directions. (Note: In subsequent chapters dealing with finite element implementation, this notation will shift to standard vector calculus to align with FEniCS unified form language syntax). The fundamental principle of mass conservation is expressed by the continuity equation. For a general, compressible fluid, this is defined as:

$$\frac{\partial \rho}{\partial t} + \frac{\partial(\rho u_i)}{\partial x_i} = 0 \quad (2.1)$$

where ρ is the fluid density. However, for most liquid cooling applications and low-speed fluid manifolds, the fluid can be safely assumed to be incompressible. Under the incompressible flow assumption, the density ρ is constant in both space and time. Therefore, the time derivative of density becomes zero, and the constant density can be factored out of the spatial derivative. Dividing the equation by ρ yields the simplified, incompressible continuity equation:

$$\frac{\partial u_i}{\partial x_i} = 0 \quad (2.2)$$

Similarly, for an incompressible, Newtonian fluid with a constant dynamic viscosity μ , the conservation of momentum is expressed as:

$$\rho \left(\frac{\partial u_i}{\partial t} + u_j \frac{\partial u_i}{\partial x_j} \right) = - \frac{\partial p}{\partial x_i} + \frac{\partial}{\partial x_j} \left(\mu \frac{\partial u_i}{\partial x_j} \right) \quad (2.3)$$

where p represents the static pressure of the fluid. In principle, solving Equation 2.3 directly yields the exact flow field. This approach is known as Direct Numerical Simulation (DNS). However, as the Reynolds number (Re) increases, the flow transitions into a chaotic turbulent state characterized by a broad spectrum of spatial and temporal scales [32]. Resolving the smallest turbulent eddies requires extremely fine spatial grids. For industrial applications, including the design of electronics cooling manifolds, DNS is computationally prohibitive. Therefore, engineers rely on turbulence modeling to predict the averaged effects of these fluctuations without resolving the complete turbulent spectrum [43, 6].

2.2 Reynolds-Averaged Navier-Stokes (RANS) Equations

The most common approach to turbulence modeling in engineering is the Reynolds-Averaged Navier-Stokes (RANS) framework, based on the theories originally introduced by Osborne Reynolds [32]. The core principle of RANS is the Reynolds decomposition. This mathematical operation separates the instantaneous flow variables into a time-averaged mean component ($\bar{u}_i, \bar{p}$) and a fluctuating component (u'_i, p'):

$$u_i = \bar{u}_i + u'_i \quad \text{and} \quad p = \bar{p} + p' \quad (2.4)$$

Substituting these decomposed variables into the standard Navier-Stokes equations and taking the time average yields the RANS momentum equations:

$$\rho \left(\frac{\partial \bar{u}_i}{\partial t} + \bar{u}_j \frac{\partial \bar{u}_i}{\partial x_j} \right) = -\frac{\partial \bar{p}}{\partial x_i} + \frac{\partial}{\partial x_j} \left(\mu \frac{\partial \bar{u}_i}{\partial x_j} - \rho \overline{u'_i u'_j} \right) \quad (2.5)$$

The averaging process introduces a new term, $-\rho \overline{u'_i u'_j}$, known as the Reynolds stress tensor. This tensor represents the momentum transfer caused by turbulent fluctuations. The introduction of this term creates a fundamental mathematical closure problem because the RANS equations now contain more unknowns than available physical equations [43].

2.2.1 The Boussinesq Hypothesis and the Isotropic Assumption

To close the system and solve for the mean flow, most RANS models rely on the Boussinesq eddy viscosity hypothesis [32]. This hypothesis assumes that the turbulent momentum transfer is analogous to molecular momentum transfer. It relates the Reynolds stresses linearly to the mean velocity gradients using a proportionality constant called the turbulent eddy viscosity (μ_t):

$$-\rho \overline{u'_i u'_j} = \mu_t \left(\frac{\partial \bar{u}_i}{\partial x_j} + \frac{\partial \bar{u}_j}{\partial x_i} \right) - \frac{2}{3} \rho k \delta_{ij} \quad (2.6)$$

where k is the turbulent kinetic energy and δ_{ij} is the Kronecker delta. The primary task of any Boussinesq-based turbulence model is to calculate this local eddy viscosity. It is crucial to highlight a fundamental physical limitation of Equation 2.6: the Boussinesq hypothesis inherently assumes that μ_t is an *isotropic* scalar quantity. It mathematically enforces that turbulence behaves identically in all spatial directions and aligns perfectly with the mean strain rate. While this assumption holds for simple shear flows, it systematically fails in predicting anisotropic phenomena, such as the strong secondary flows induced by streamline curvature.

2.3 Boundary Layer Theory and Near-Wall Treatment

Before selecting a specific turbulence equation to calculate μ_t , the physical structure of a turbulent boundary layer must be considered, as the model must accurately resolve these near-wall regions. The boundary layer is traditionally divided into three distinct regions based on the non-dimensional wall distance (y^+) and non-dimensional velocity (u^+):

$$y^+ = \frac{yu_\tau}{\nu}, \quad u^+ = \frac{\bar{u}}{u_\tau} \quad (2.7)$$

where y is the absolute distance from the wall, ν is the kinematic viscosity (μ/ρ), and $u_\tau = \sqrt{\tau_w/\rho}$ is the friction velocity derived from the wall shear stress (τ_w) [32].

1. **The viscous sublayer** ($y^+ < 5$): Immediately adjacent to the wall, turbulent fluctuations are damped out by kinematic constraints. Viscous shear dominates, and the velocity profile is strictly linear ($u^+ = y^+$).
2. **The buffer layer** ($5 < y^+ < 30$): A transition region where both viscous and turbulent shear stresses are of equal magnitude.
3. **The log-law region** ($y^+ > 30$): Farther from the wall, turbulence dominates, and the velocity follows a logarithmic distribution predicted by the von Kármán constant [43].

High-Reynolds models (like the standard $k - \epsilon$ model) use empirical "wall functions" to bridge the gap between the wall and the log-law region, requiring the first computational mesh element to be placed at $y^+ > 30$. However, for topology optimization, where solid-fluid boundaries continuously evolve during the optimization loop, empirical wall functions introduce severe mathematical discontinuities and numerical instabilities. Therefore, a "wall-resolved" model that integrates the governing equations all the way through the viscous sublayer ($y^+ \approx 1$) is strictly required for robust gradient-based optimization.

2.4 The Spalart-Allmaras One-Equation Model

While two-equation models (such as $k - \epsilon$ or $k - \omega$) solve separate transport equations for the turbulent kinetic energy and a turbulent length scale, they introduce highly

non-linear source terms. These terms frequently cause numerical instabilities during the adjoint sensitivity analysis required for topology optimization. To achieve a trade-off between physical accuracy and computational efficiency, this thesis utilizes the Spalart-Allmaras (SA) model, originally developed in 1992 for aerodynamic flows with adverse pressure gradients [35]. The SA model is a one-equation model that solves a single transport equation directly for a working variable, $\tilde{\nu}$, which is subsequently linked to the true kinematic eddy viscosity ($\nu_t = \mu_t/\rho$) [36]. The transport equation for the standard SA model is given by:

$$\frac{\partial \tilde{\nu}}{\partial t} + \bar{u}_j \frac{\partial \tilde{\nu}}{\partial x_j} = c_{b1} \tilde{S} \tilde{\nu} - c_{w1} f_w \left(\frac{\tilde{\nu}}{d} \right)^2 + \frac{1}{\sigma} \left[\frac{\partial}{\partial x_j} \left((\nu + \tilde{\nu}) \frac{\partial \tilde{\nu}}{\partial x_j} \right) + c_{b2} \left(\frac{\partial \tilde{\nu}}{\partial x_j} \right)^2 \right] \quad (2.8)$$

The terms on the right-hand side represent the production, destruction, and diffusion of the turbulent working variable, respectively. Crucially, the production term relies on a modified vorticity scalar, $\tilde{S}$. This indicates that turbulence generation in the SA model is fundamentally driven by fluid shear, a physical characteristic that introduces unique boundary-condition sensitivities during advanced adjoint calculations (discussed further in Chapter 4). Furthermore, the destruction term is explicitly a function of d , the absolute distance to the nearest wall. This makes the SA model a highly effective "wall-resolved" low-Reynolds model. It is inherently sensitive to near-wall damping effects and integrates smoothly through the viscous sublayer without requiring complex empirical wall functions or a second length-scale equation [33]. Detailed derivations of the model constants and wall damping functions are deferred to Chapter 3, where the numerical implementation in FEniCS is discussed.

2.5 Evaluation of Curvature Effects and the Dean Number

Despite its computational advantages for topology optimization, the standard SA model relies on a single length scale tied to the nearest wall distance. As it lacks an explicit rotation/curvature correction and remains constrained by the Boussinesq eddy-viscosity assumption, it struggles to capture the anisotropic physics of streamline curvature. When fluid passes through a bend, centrifugal forces induce secondary flows perpendicular to the main flow direction. To quantify the strength of these secondary flows, the dimensionless Dean number (De) is utilized:

$$De = Re_d \sqrt{\frac{d}{2R_c}} \quad (2.9)$$

where d is the characteristic length of the cross-section (e.g., the pipe diameter D or hydraulic diameter D_h), R_c is the radius of curvature, and Re_d is the bulk Reynolds number based strictly on that exact same characteristic length. A higher Dean number indicates a stronger cross-sectional influence of centrifugal forces.

A review of the literature demonstrates that the predictive capability of the SA model is heavily dependent on this Dean number. For mildly curved flows, the SA model performs relatively well. Apalowo [8] evaluated a 90° circular pipe bend with a Dean number of $De \approx 9,000$ (operating at $Re_D \approx 34,000$, based on the pipe diameter). Under these moderate conditions, where the bend radius was seven times larger than the pipe diameter ($R_c = 7D$), the SA model provided a highly acceptable approximation of the streamwise velocity profiles compared to higher-order models, successfully capturing the primary flow physics. This moderate-curvature agreement is illustrated in Figure 2.1, where the SA velocity profiles remain close to both the experimental measurements and the two-equation model predictions.

However, as the curvature tightens and the Dean number increases, the limitations of the isotropic Boussinesq assumption become severe. Abuhatira et al. [1] evaluated the internal turbulent flow through a 90° pipe bend characterized by a tighter curvature ratio of $R_c = 2.0D$. Operating at a bulk Reynolds number of $Re_D = 60,000$, this geometry generated strong centrifugal forces and a corresponding Dean number of $De \approx 30,000$. Unlike extreme benchmark cases where all isotropic models fail uniformly, this configuration distinctly isolates the shortcomings of the Spalart-Allmaras formulation. At the critical bend exit ($X/D = 0$), the SA model drastically mischaracterized the near-wall velocity profiles, completely missing the flow physics close to the inner wall that higher-order two-equation models (such as $k - \epsilon$ and $k - \omega$) resolved with much greater accuracy. This confirms that the failure is not attributed to inertial scaling, but is strictly driven by the model's inability to compensate for strong anisotropic centrifugal effects. The resulting near-wall mismatch is visible in Figure 2.2, which motivates the additional Ansys benchmark cases considered below.

While the literature clearly establishes that the SA model is viable at $De \approx 9,000$ and fails to capture inner-wall physics at $De \approx 30,000$, a bracketed or case-specific threshold between these states must be estimated to justify its use in topology optimization. Therefore, to better evaluate where the model begins to break down, rigorous benchmark studies were subsequently conducted within Ansys Fluent.




FIGURE 2.1: Streamwise velocity profiles in a 90° pipe bend ($Re_c = 7D$, $Re_D \approx 34,000$, $De \approx 9,000$). The SA model performs relatively well, closely matching the experimental data and two-equation models for mild curvature. Reprinted from Apalowo [8].






FIGURE 2.2: Turbulence model predictions at the 90° bend exit ($X/D = 0$, $Re_c = 2.0D$, $Re_D = 60,000$, $De \approx 30,000$). Due to the strong curvature, the standard Spalart-Allmaras model fails to accurately capture the near-wall velocity profile compared to two-equation closures. Adapted from Abuhatira et al. [1].

2.6 Benchmark Studies in Ansys Fluent

To investigate where the SA model’s accuracy deteriorates, benchmark simulations were performed using **Ansys Fluent**. The study focused on two distinct, fully three-dimensional (3D) internal flow geometries: a 180° U-bend and a tighter 90° pipe bend. To ensure a standardized and rigorously tested numerical baseline, the computational meshes, fluid properties, and boundary conditions for these simulations were adopted directly from the official Ansys Fluid Dynamics Verification Manual [7]. Specifically, the 90° bend corresponds to verification case VMFL030, while the 3D U-bend corresponds to case VMFL048. These cases evaluate industrial flow regimes ($Re_D > 40,000$) to ensure the turbulence model is rigorously verified against dominant inertial physics before being implemented into the custom FEn-iCS solver. To evaluate the accuracy of the one-equation model, the results were compared directly against the SST $k - \omega$ model, an industry standard for accurately predicting adverse pressure gradients and flow separation.

2.6.1 Mild Curvature: The U-Bend Case

The first benchmark evaluated a 3D U-bend geometry characterized by a moderate curvature ratio of $Re_c \approx 4.46D$. Operating at a bulk Reynolds number of $Re_D \approx 44,700$, this geometry resulted in a Dean number of $De \approx 15,000$. Consistent with the findings of Apalowo for moderate bends, the SA model performed well under these conditions. As shown in Figure 2.3, the qualitative velocity contours of the SA model accurately replicate the bulk flow distribution of the SST $k - \omega$ model. The velocity profiles in the lower panel are extracted along the pipe-radius line visible on the contour plots above, thereby linking the one-dimensional profile comparison directly to the displayed exit-plane flow field. Furthermore, the quantitative velocity profiles agree well with the experimental data from Takamasa and Kondo [40].



(c) Velocity Profiles



FIGURE 2.3: Flow predictions for a 180° U-bend ($R_c \approx 4.46D$, $Re_D \approx 44,700$, $De \approx 15,000$). The SA model accurately matches the higher-fidelity SST $k - \omega$ contours and directly aligns with the experimental data from Takamasa and Kondo [40].

2.6.2 Strong Curvature: The 90° Bend Case

The second benchmark evaluated a tighter 90° pipe bend characterized by a severe curvature ratio of $R_c = 2.8D$. Despite operating at a very similar bulk Reynolds number of $Re_D \approx 43,000$ compared to the U-bend, this significantly tighter bend radius pushed the Dean number up to $De \approx 18,000$. Because the bulk flow speed and Reynolds number remain effectively constant between these two benchmark cases, this comparison successfully isolates the geometric effect of streamline curvature.

As shown in Figure 2.4, examining the quantitative velocity profile at the bend exit ($\theta = 90^\circ$) clearly demonstrates the breakdown of the SA model. In this case, the profile extraction spans the entire pipe diameter at the bend exit; because the line coincides with the full displayed diameter, it is not as visually distinguishable in the contour panels as the radius marker used in the U-bend case. Driven primarily by the tighter bend radius and amplified anisotropic centrifugal effects, the SA model over-predicts the turbulent viscosity. This causes it to over-smooth the velocity

profile, failing to capture the sharp velocity gradients near the inner wall ($y/D = 0$) when compared against both the SST $k - \omega$ prediction and the experimental data from Enayet et al. [17]. This failure is also visually apparent when comparing the side-by-side velocity contours.



(c) Velocity Profiles

FIGURE 2.4: Flow predictions at the 90° pipe bend exit ($R_c = 2.8D$, $Re_D \approx 43,000$, $De \approx 18,000$). The SA model over-smooths the velocity gradients near the inner wall, deviating significantly from both the SST $k - \omega$ model and the experimental data from Enayet et al. [17].

2.7 Conclusion

This chapter directly addresses the first research question regarding the predictive capabilities and physical limitations of the Spalart-Allmaras model. The theoretical breakdown and Ansys benchmark studies confirm that for internal geometries with moderate curvature (up to $De \approx 15,000$), the SA model performs accurately and successfully captures the primary flow physics. Furthermore, because it is a wall-resolved model, it avoids the numerical pitfalls of empirical wall functions, making

it exceptionally well-suited for boundary-evolving topology optimization. However, its physical limitation lies in the isotropic Boussinesq assumption. When streamline curvature induces strong anisotropic secondary flows (heuristically observed to cause model breakdown in this case beyond a threshold of $De \approx 18,000$), the uncompensated SA model systematically overpredicts eddy viscosity and fails to capture flow separation. Consequently, while the SA model provides the necessary numerical stability to drive the optimization algorithms in FEniCS, the final topological designs must be critically verified. Any complex manifold generated by the SA optimizer must be subsequently evaluated in commercial software using a higher-fidelity two-equation model to confirm its true aerodynamic performance.

Chapter 3

Implementation of the Spalart-Allmaras Model in FEniCS

3.1 Introduction

With the theoretical foundation of the Spalart-Allmaras (SA) model established in Chapter 2, this chapter details its numerical translation into a working FEniCS solver. The focus is methodological: the finite-element weak forms, Picard coupling strategy, pressure-correction flow solver, SA transport equation, and relaxed Eikonal wall-distance calculation are described before the solver is verified separately in Chapter 5.

The numerical simulations and eventual topology optimization in this thesis are built upon FEniCS, an open-source computing platform for solving partial differential equations (PDEs). Unlike commercial computational fluid dynamics (CFD) software such as Ansys, which typically relies on the Finite Volume Method (FVM), FEniCS uses the Finite Element Method (FEM) and requires the user to explicitly define the mathematical formulation of the problem [28].

To solve a PDE in FEniCS, the governing equations (the "strong form") must be mathematically transformed into a "weak formulation" (or variational form). FEniCS utilizes the Unified Form Language (UFL) to express these weak forms in Python code that closely mirrors the mathematical notation [26]. This high-level mathematical scripting is particularly advantageous for topology optimization, as it allows for the automated derivation of the adjoint equations required for gradient-based optimization algorithms.

To ensure complete transparency and reproducibility, the entire custom codebase developed for this thesis is available in the public GitHub repository **TurbulentT0**. To maintain modularity, the repository is divided into two primary directories. The **TurbulenceModels** directory contains the standalone Spalart-Allmaras fluid dynamics engine and the verification benchmarks reported in Chapter 5. The **FluidT0** directory houses the optimization algorithms, penalty formulations, and objective

functions detailed in Chapter 4.

3.1.1 Derivation of a Weak Formulation

Before diving into the complex turbulence equations, it is instructive to demonstrate how a continuous PDE is translated into a discrete FEniCS formulation. To illustrate this methodology, consider the Poisson equation for pressure, which is a fundamental component of incompressible fluid solvers. The strong form of the Poisson equation is given by:

$$-\nabla^2 p = f \quad \text{in } \Omega \quad (3.1)$$

where p is the pressure field, f is a source term (derived from the velocity divergence), and Ω is the computational domain. To obtain the weak form, the equation is multiplied by a sufficiently smooth test function v and integrated over the entire domain Ω , following standard Galerkin finite element procedures [28, 26]:

$$-\int_{\Omega} (\nabla^2 p) v \, dx = \int_{\Omega} f v \, dx \quad (3.2)$$

Applying integration by parts (Green's first identity) to the left-hand side yields:

$$\int_{\Omega} \nabla p \cdot \nabla v \, dx - \int_{\partial\Omega} (\nabla p \cdot \mathbf{n}) v \, ds = \int_{\Omega} f v \, dx \quad (3.3)$$

where $\partial\Omega$ is the boundary of the domain and $\mathbf{n}$ is the outward-pointing normal vector. This integral formulation fundamentally lowers the continuity requirements of the solution, requiring only first-order derivatives to be square-integrable. Furthermore, Neumann boundary conditions (e.g., a prescribed pressure gradient at an outlet) can be naturally substituted directly into the boundary integral $\int_{\partial\Omega} (\nabla p \cdot \mathbf{n}) v \, ds$. This variational approach forms the foundation of all fluid and turbulence equations implemented in this thesis.

3.2 Numerical Solution Strategy for Fluid Dynamics

Solving the full Reynolds-Averaged Navier-Stokes (RANS) equations is numerically challenging due to the non-linear convective terms, the saddle-point nature of the coupled velocity-pressure system, and the intense cross-coupling with the turbulence model. To achieve a stable and efficient steady-state solution in FEniCS that seamlessly translates into the topology optimization framework, the governing equations are rigorously decoupled.

3.2.1 Decoupling the Physics: The Picard Iteration Scheme

Because the momentum equation relies on the turbulent eddy viscosity (ν_t), and the turbulence transport equation relies on the convective velocity ($\mathbf{u}$), solving the entire fluid-turbulence system simultaneously in a massive monolithic matrix frequently leads to ill-conditioning and solver divergence.

To resolve this, the forward physics are decoupled using a **Picard iteration scheme** (successive substitution), a robust approach for segregated RANS solvers [19, 43]. During each Picard step, the flow solver advances using the eddy viscosity from the previous iteration. The newly updated velocity field is then passed to the Spalart-Allmaras solver to calculate a new eddy viscosity field. This alternating sequence forms the outer loop of the solver, iteratively updating the decoupled equations until the fully coupled physical system reaches a steady-state equilibrium.

3.2.2 Incremental Pressure Correction Scheme (IPCS)

Within the Picard loop, the Navier-Stokes equations themselves must be solved. Rather than solving for velocity and pressure simultaneously, this implementation utilizes an Incremental Pressure Correction Scheme (IPCS). This algorithmic splitting is fundamentally based on the fractional-step projection methods originally pioneered by Chorin [14] and Temam [41]. To nurse the non-linear convective momentum terms to a stable steady state, a pseudo-time stepping approach (Δt) is employed strictly for the flow variables. The IPCS algorithm splits the fluid equations into three sequential steps per pseudo-time step:

1. **Tentative velocity step:** The momentum equation is solved using the pressure field from the previous iteration (p^n) to compute an intermediate, non-divergence-free velocity field ($\mathbf{u}^*$). Because this is a RANS framework, the diffusion term utilizes the effective kinematic viscosity, $(\nu + \nu_t)$. Treating the convective term explicitly and the diffusive term implicitly, the semi-discrete equation is:

$$\frac{\mathbf{u}^* - \mathbf{u}^n}{\Delta t} + (\mathbf{u}^n \cdot \nabla) \mathbf{u}^* - \nabla \cdot ((\nu + \nu_t^n) \nabla \mathbf{u}^*) = -\nabla p^n \quad (3.4)$$

2. **Pressure correction step:** A Poisson equation is solved to find the pressure variation that enforces the continuity constraint ($\nabla \cdot \mathbf{u}^{n+1} = 0$). The equation for the new pressure field (p^{n+1}) is derived as:

$$\nabla^2 p^{n+1} = \nabla^2 p^n + \frac{1}{\Delta t} \nabla \cdot \mathbf{u}^* \quad (3.5)$$

3. **Velocity update step:** The tentative velocity is strictly corrected using the gradient of the pressure difference to yield a mass-conserving velocity field ($\mathbf{u}^{n+1}$) for the current iteration:

$$\mathbf{u}^{n+1} = \mathbf{u}^* - \Delta t \nabla (p^{n+1} - p^n) \quad (3.6)$$

The simulation is deemed fully converged to a steady state when the L_2 norm of the residual difference between successive Picard iterations falls below a defined tolerance limit (ϵ_{tol}) for both the fluid and turbulence fields.

3.3 Detailed Formulation of the Spalart-Allmaras Model

As concluded in Chapter 2, the Spalart-Allmaras (SA) one-equation model was selected for its balance of physical accuracy and numerical stability. Unlike the transient flow equations, the SA transport equation is solved in its purely steady-state form within the Picard loop. The steady transport equation for the continuous nonlinear working variable $\tilde{\nu}$ is given by:

$$\mathbf{u} \cdot \nabla \tilde{\nu} = c_{b1} \tilde{S} \tilde{\nu} - c_{w1} f_w \left(\frac{\tilde{\nu}}{y} \right)^2 + \frac{1}{\sigma} [\nabla \cdot ((\nu + \tilde{\nu}) \nabla \tilde{\nu}) + c_{b2} (\nabla \tilde{\nu})^2] \quad (3.7)$$

where ν is the molecular kinematic viscosity, $\sigma = 2/3$ is the turbulent Prandtl number, c_{b1} and c_{w1} are calibrated closure constants, $\tilde{S}$ is a modified vorticity magnitude, f_w is the wall-damping function, and y is the physical distance to the nearest solid boundary [36].

3.3.1 Closure Terms and Calibration Constants

The Spalart-Allmaras model relies on several interconnected closure functions to accurately capture boundary layer physics. The production term relies on a modified vorticity magnitude, $\tilde{S}$, which prevents the destruction of turbulence in regions where the vorticity goes to zero, while maintaining correct scaling near the wall [36]. It is defined as:

$$\tilde{S} = S + \frac{\tilde{\nu}}{\kappa^2 y^2} f_{v2} \quad (3.8)$$

where $S = \sqrt{2\Omega_{ij}\Omega_{ij}}$ is the scalar magnitude of the vorticity tensor, κ is the von Kármán constant, and y is the distance to the nearest wall. The function f_{v2} is a near-wall modifier:

$$f_{v2} = 1 - \frac{\chi}{1 + \chi f_{v1}} \quad \text{where} \quad \chi = \frac{\tilde{\nu}}{\nu} \quad \text{and} \quad f_{v1} = \frac{\chi^3}{\chi^3 + c_{v1}^3} \quad (3.9)$$

Similarly, the wall destruction term is modulated by the non-dimensional function f_w , which accelerates the decay of turbulent viscosity inside the viscous sublayer:

$$f_w = g \left(\frac{1 + c_{w3}^6}{g^6 + c_{w3}^6} \right)^{1/6} \quad (3.10)$$

where the intermediate variables g and r are evaluated as:

$$g = r + c_{w2}(r^6 - r) \quad \text{and} \quad r = \min \left(\frac{\tilde{\nu}}{\tilde{S} \kappa^2 y^2}, 10 \right) \quad (3.11)$$

The upper bound of 10 on r ensures numerical stability during solver initialization when $\tilde{S}$ may momentarily approach zero.

It should be explicitly noted that the original formulation by Spalart and Allmaras [36] includes additional geometric trip functions (f_{t1} and f_{t2}) designed to manually trigger the transition from laminar to turbulent flow at specific user-defined coordinates. However, following the topology optimization methodology established by Yoon [47], these transition functions are deliberately omitted from this thesis. Because the topological boundaries are unknown *a priori* and evolve during optimization, tracking discrete trip points is mathematically unfeasible; therefore, the model assumes a fully turbulent flow regime.

To close the system, the standard empirical calibration constants formulated by Spalart and Allmaras are utilized throughout the FEniCS implementation [33]. Notably, the von Kármán constant is set to 0.41 to perfectly align with the standard OpenFOAM and generalized industrial implementations[cite: 3]:

$$\begin{aligned} c_{b1} &= 0.1355 & \sigma &= 2/3 & c_{b2} &= 0.622 \\ \kappa &= 0.41 & c_{w2} &= 0.3 & c_{w3} &= 2.0 \\ c_{v1} &= 7.1 & c_{w1} &= \frac{c_{b1}}{\kappa^2} + \frac{1 + c_{b2}}{\sigma} \end{aligned}$$

3.3.2 Wall-Distance Calculation via the Relaxed Eikonal Equation

A major computational challenge in the SA model is the calculation of the wall distance y . In a standard static CFD simulation, y is purely geometric and can be pre-computed. However, to ensure this verification solver translates seamlessly into the topology optimization code (where solid boundaries evolve continuously), the implementation abandons discrete geometric searches in favor of a PDE-based approach.

The true geometric wall distance satisfies the standard Eikonal equation, $|\nabla y| = 1$. However, this formulation contains singularities at the medial axes of the domain. To guarantee continuous differentiability, a requirement for later adjoint derivations, a relaxed Eikonal equation proposed by Yoon [47] is solved:

$$\nabla G \cdot \nabla G + \sigma_w G (\nabla \cdot \nabla G) = (1 + 2\sigma_w) G^4 \quad (3.12)$$

Here, G is an intermediate, non-dimensional field variable. The parameter σ_w is a relaxation constant (typically set to 0.1) that controls the smoothness of the field. The PDE is subjected to a Dirichlet boundary condition at the physical walls, $G = G_0$, where G_0 is an arbitrarily large constant (e.g., $G_0 = 20.0$). Once the field G is solved over the domain, the physical wall distance y is recovered algebraically at every node using the inverse relationship:

$$y = \frac{1}{G} - \frac{1}{G_0} \quad (3.13)$$

This system is implemented as a nonlinear weak residual for G . After multiplying by a scalar test function and integrating the $\sigma_w G \nabla^2 G$ term by parts, the residual contains only first derivatives of G . Listing 3.1 shows the weak-form structure used by the solver: the first two terms represent the relaxed Eikonal operator, the third

term is the nonlinear source, and the Dirichlet condition fixes G on physical walls before recovering y algebraically.

```

1      # Unknown reciprocal wall-distance field G, test function z.
2      F_G = (
3          (1.0 - sigma_w) * inner(grad(G), grad(G)) * z * dx
4          - sigma_w * G * inner(grad(G), grad(z)) * dx
5          - (1.0 + 2.0 * sigma_w) * G**4 * z * dx
6      )
7
8      wall_bc = DirichletBC(V_G, G0, physical_walls)
9      solve(F_G == 0, G, wall_bc, J=derivative(F_G, G))
10
11     G = enforce_lower_bound(G, G_min)
12     y = 1.0 / G - 1.0 / G0

```

LISTING 3.1: Pseudocode for the relaxed Eikonal wall-distance solve.

3.4 Variational Formulation and FEniCS Implementation

The standard Galerkin weak formulation is constructed by multiplying the steady transport Equation 3.7 by a test function ξ and applying integration by parts to the diffusive terms. Because the equation is highly non-linear, robust convergence is dependent on the stability of the external velocity field provided by the Picard outer loop.

To manage this non-linearity, the equation is linearized within the Picard iteration scheme by freezing certain parameters using the most recent value of the working variable, $\tilde{\nu}^0$. Specifically, the turbulence production and cross-diffusion terms are treated as explicit sources evaluated using $\tilde{\nu}^0$. Conversely, the wall destruction term is linearized into an implicit reaction sink by scaling the continuous $\tilde{\nu}$ against the frozen variable, taking the form $c_{w1}f_w(\tilde{\nu}^0/y^2)\tilde{\nu}$.

The finite-element implementation mirrors the analytical weak form using the most recent convective velocity field passed from the flow solver. Listing 3.2 shows the residual assembled in each Picard step. The convection, implicit diffusion, and implicit destruction terms are placed on the left-hand side, while production and cross-diffusion are evaluated from the previous SA field and treated as explicit sources.

```

1      # Unknown nu_tilde, test function xi, frozen previous field nu_tilde0.
2      S_tilde0 = modified_vorticity(u, nu_tilde0, y)
3      f_w0 = wall_destruction_function(nu_tilde0, S_tilde0, y)
4
5      reaction = c_w1 * f_w0 * nu_tilde0 / y_safe**2
6      source = (
7          c_b1 * S_tilde0 * nu_tilde0
8          + (c_b2 / sigma) * inner(grad(nu_tilde0), grad(nu_tilde0))
9      )
10

```

```

11     F_SA = (
12         dot(u, grad(nu_tilde)) * xi * dx
13         + inner((nu + nu_tilde0) / sigma * grad(nu_tilde), grad(xi)) * dx
14         + reaction * nu_tilde * xi * dx
15         - source * xi * dx
16     )
17
18     solve(lhs(F_SA) == rhs(F_SA), nu_tilde, sa_boundary_conditions)
19     nu_t = eddy_viscosity_from(nu_tilde)

```

LISTING 3.2: Pseudocode for the Picard-linearized SA weak-form assembly.

3.5 Conclusion

This chapter has established the numerical formulation of the FEniCS-SA flow solver used throughout the remainder of the thesis. The key ingredients are the weak-form finite-element implementation, Picard decoupling of the velocity-pressure and turbulence fields, IPCS pseudo-time stepping for the incompressible flow solve, the steady SA transport equation, and the relaxed Eikonal wall-distance formulation.

Separating these methods from their verification keeps the thesis logic clean: this chapter explains how the solver is constructed, Chapter 4 extends the same finite-element machinery to density-based topology optimization and adjoint sensitivities, and Chapter 5 then verifies the standalone SA solver against Ansys Fluent benchmark cases before the topology optimization results are interpreted.

Chapter 4

Turbulent Topology Optimization Methodology

4.1 Introduction

Building upon the Spalart-Allmaras (SA) fluid dynamics solver [35, 36] developed in Chapter 3, this chapter introduces the mathematical framework required to perform Topology Optimization (TO) in the turbulent regime. Chapter 3 explained how the forward turbulent flow problem is solved; the present chapter explains how that forward solver is embedded in a density-based design loop and differentiated for gradient-based optimization. The forward solver is verified separately in Chapter 5. The core objective of this methodology is to algorithmically determine the optimal distribution of fluid channels within a given design domain to reduce hydraulic or mechanical energy loss under a prescribed flow condition.

This chapter details the density-based fictitious-domain formulation used to represent solid and fluid regions on a fixed computational mesh [13, 20]. It then explains the specific turbulent-flow ingredients retained for the remainder of the thesis: a Brinkman-type momentum resistance, a power-law topology penalty for the SA and wall-distance equations, and the reciprocal penalized wall-distance formulation used by the benchmark computations.

To solve this optimization problem in a computationally tractable manner, continuous adjoint sensitivity analysis is required [31, 30]. The optimization results reported in this thesis use the frozen-turbulence, or Constant Eddy Viscosity (CEV), assumption [46, 34]. This assumption is stated explicitly because it is both the practical reason the turbulent optimization loop remains stable and the main limitation of the sensitivity model.

Following the derivation of these sensitivities, the continuous optimization problem is formally defined, establishing the available loss objectives and the strict fluid volume constraints over explicitly defined design regions. Finally, the mathematical techniques for active-domain density filtering [27] and Heaviside projection [44] are detailed to ensure the generation of manufacturable designs.

4.2 The Fictitious Domain Method and Brinkman Penalization

In standard shape optimization, the solid-fluid boundary is explicitly defined by the mesh, and the nodes are physically moved to improve the design. In density-based topology optimization, however, the computational domain (Ω) remains strictly fixed. To create solid walls and fluid channels within this fixed grid, the Fictitious Domain Method is employed.

Originally applied to Stokes flow by Borrvall and Petersson [13] and extended to the Navier-Stokes equations by Gersborg-Hansen et al. [20], this method assigns a continuous design variable to the spatial coordinates in the domain. In the convention used here, a physical topology value of $\bar{\gamma} = 1$ represents a pure **fluid** channel, while $\bar{\gamma} = 0$ represents a **solid**, impermeable structure. To force the fluid velocity to zero inside the solid regions without mathematically removing the finite elements, the fluid is treated as a fictitious porous medium.

This flow resistance is quantified by an inverse permeability term, or Brinkman penalty function, denoted as $\alpha(\bar{\gamma})$. To ensure the optimization algorithm (which relies on continuous gradients) can transition smoothly between fluid and solid states, α must be a continuous interpolation function. The formulation uses a convex Rational Approximation of Material Properties (RAMP) interpolation [37] scaled between a maximum solid penalty (α_{solid}) and a minimum fluid penalty (α_{fluid}):

$$\alpha(\bar{\gamma}) = \alpha_{solid} + (\alpha_{fluid} - \alpha_{solid}) \frac{\bar{\gamma}(1+q)}{\bar{\gamma}+q} \quad (4.1)$$

where α_{solid} is a sufficiently large penalty factor that completely stagnates the flow when $\bar{\gamma} = 0$, and q is a tuning parameter that controls the convexity of the penalization curve. Equation 4.1 is evaluated using the projected effective density $\bar{\gamma}$, rather than the raw optimizer variable γ , because the physical flow equations are formulated on the filtered, projected, and domain-bounded topology. During continuation, q is increased to sharpen the fluid-solid interpolation while keeping the underlying RAMP form differentiable. Unlike traditional structural optimization which relies on the SIMP power-law [12], fluid optimization generally avoids SIMP for momentum because its derivative vanishes at $\bar{\gamma} = 0$. By utilizing the RAMP formulation, a strictly non-zero gradient is guaranteed across the entire active design domain.

4.2.1 The Density Chain and Domain Bounding

While conceptually represented by the physical topology field in the governing equations, the layout is managed through a distinct chain of density representations to handle design and non-design domains. To explicitly avoid mathematical confusion with the fluid mass density (ρ) utilized in the Reynolds-Averaged Navier-Stokes equations, this thesis exclusively utilizes the γ notation for all topological variables:

- **Raw design density** (γ): The baseline parameter manipulated by the optimizer.

- **Filtered density** ($\tilde{\gamma}$): The spatially smoothed design field.
- **Projected effective density** ($\bar{\gamma}$): The bounded, physical topology field evaluated by the governing equations.

The raw design density, γ , is defined as a Discontinuous Galerkin (DG0) [23] function over the entire mesh. The optimization algorithm strictly manipulates the degrees of freedom corresponding to the actively optimized design space, while passive non-design cells are subsequently restored to their prescribed bounding values.

This raw variable is subjected to a spatial PDE filter [27] to yield a smoothed field $\tilde{\gamma}$. To recover sharp boundaries and strictly preserve required non-design regions, the filtered density is then projected and mapped into the bounded effective density ($\bar{\gamma}$) [22]:

$$\bar{\gamma} = \gamma_{\text{lower}} + (\gamma_{\text{upper}} - \gamma_{\text{lower}})\Pi_{\beta}(\tilde{\gamma}) \quad (4.2)$$

where Π_{β} is the projection operator, and γ_{lower} and γ_{upper} define the absolute permissible bounds for any given cell. This rigorous bounding guarantees that predefined features, such as non-design fluid inlet extensions ($\gamma_{\text{lower}} = \gamma_{\text{upper}} = 1$) or explicitly constrained structural walls ($\gamma_{\text{lower}} = \gamma_{\text{upper}} = 0$), remain rigidly fixed without interfering with the mathematical behavior of the active optimization space.

The computational mesh is therefore partitioned into an active design domain and passive non-design domains. The active design domain contains the cells whose raw densities are updated by MMA. Fluid non-design domains are fixed as permanent fluid, typically to preserve inlet and outlet extensions, buffer channels, or measurement regions. Solid non-design domains are fixed as permanent solid, typically to impose external walls, blocked regions, or geometric features that must not be modified. These passive cells still participate in the forward flow, turbulence, and wall-distance solves, but they are excluded from the optimization masks: the volume constraint is evaluated only over the prescribed volume mask Ω_V , and any domain-integrated post-processing metric is evaluated only over its prescribed integration mask. Consequently, prescribed inlet/outlet buffers and mandatory solid walls do not count toward the optimized fluid fraction and do not directly bias the reported domain-integrated metrics.

4.3 Penalization of the Governing Equations

To effectively optimize turbulent manifolds, the respective penalties must be rigorously integrated into the turbulence model and the wall-distance calculations to ensure the physics remain consistent as the solid boundaries evolve. All physical penalties strictly utilize the projected effective density, $\bar{\gamma}$.

4.3.1 Modified Reynolds-Averaged Navier-Stokes Equations

The primary application of the Brinkman penalty is the modification of the steady-state, incompressible RANS momentum equation. The RAMP-interpolated inverse

permeability $\alpha(\bar{\gamma})$ is added as a linear sink term (Darcy friction) to the momentum conservation equation:

$$\rho(\mathbf{u} \cdot \nabla)\mathbf{u} = -\nabla p + \nabla \cdot [(\mu + \mu_t)(\nabla\mathbf{u} + (\nabla\mathbf{u})^T)] - \alpha(\bar{\gamma})\mathbf{u} \quad (4.3)$$

It is critical to note the retention of the full symmetric rate-of-strain tensor. Here, ρ strictly denotes the fluid mass density. Because the turbulent eddy viscosity (μ_t) varies drastically across the spatial domain [32], the effective fluid viscosity is no longer constant, and the transposed term exerts a non-zero physical influence on the momentum balance. In regions where the optimizer places solid material ($\bar{\gamma} \rightarrow 0$), the enormous magnitude of $\alpha_{solid}\mathbf{u}$ dominates all inertial and viscous forces, forcing the local velocity mathematically to zero.

4.3.2 Modified Spalart-Allmaras Transport Equation

As the Fictitious Domain Method forces velocity to zero in the solid regions, the turbulence generation inside those regions should theoretically cease. However, due to numerical diffusion, turbulent eddy viscosity ($\tilde{\nu}$) can inadvertently bleed into the solid regions. This unphysical behavior destabilizes the solver and heavily distorts the adjoint sensitivity analysis.

To strictly enforce a laminar, zero-turbulence state inside the solid structures, a corresponding penalty sink term is added to the Spalart-Allmaras transport equation [47, 16]:

$$\begin{aligned} \mathbf{u} \cdot \nabla \tilde{\nu} = & c_{b1} \tilde{S} \tilde{\nu} - c_{w1} f_w \left(\frac{\tilde{\nu}}{y} \right)^2 + \frac{1}{\sigma} [\nabla \cdot ((\nu + \tilde{\nu}) \nabla \tilde{\nu}) + c_{b2} (\nabla \tilde{\nu})^2] \\ & - \alpha_{\tilde{\nu}} (1 - \bar{\gamma})^{N_{\tilde{\nu}}} \tilde{\nu} \end{aligned} \quad (4.4)$$

The last term is the power-type topology penalty used to suppress the SA working variable in solid material. Because $\bar{\gamma} = 1$ denotes fluid and $\bar{\gamma} = 0$ denotes solid, the penalty is applied to the local solid fraction $(1 - \bar{\gamma})$. The coefficient $\alpha_{\tilde{\nu}}$ sets the damping strength and the exponent $N_{\tilde{\nu}}$ controls how sharply the damping grows as a cell becomes solid. In the final benchmark computations these two parameters are held constant throughout the optimization. The term is therefore zero in pure fluid, reaches its full strength in solid cells, and acts only as an artificial damping device for the working variable, not as an additional turbulence-model closure.

4.3.3 Reciprocal Penalized Wall-Distance Equation

The Spalart-Allmaras model is highly sensitive to the wall distance, y . While differential equations for wall distance computation are well-established [42], topology optimization introduces a unique challenge. The location of the wall is not a static mesh boundary, but rather the dynamically evolving $\bar{\gamma} = 0$ contour. Therefore, the relaxed Eikonal equation must be penalized to dynamically recognize the fictitious solid as a physical boundary [47]:

$$\nabla G \cdot \nabla G + \sigma_w G (\nabla \cdot \nabla G) - (1 + 2\sigma_w) G^4 - \alpha_G (1 - \bar{\gamma})^{N_G} (G - G_0) = 0 \quad (4.5)$$

The wall-distance equation uses the same power-law topology penalty. The coefficient α_G sets the strength of the artificial wall-distance forcing and the exponent N_G controls its growth through gray material. As for the SA penalty, these parameters are held constant in the final benchmark computations. In fluid regions the term vanishes, while in topology-created solid regions it pulls G toward G_0 . The algebraic recovery function

$$y = \frac{1}{G} - \frac{1}{G_0} \quad (4.6)$$

then collapses the SA wall distance toward zero inside topology-created solid regions while retaining the physical wall distance in fluid regions.

4.4 Boundary Conditions and Turbulence Initialization

To stabilize the topology optimization and ensure physical accuracy, boundary constraints are strictly managed across both fluid and turbulence regimes. Non-design fluid buffer regions are strictly mandated at inlets and outlets to isolate the active optimization domain from artificial numerical shear. The framework robustly handles boundary interactions by incorporating pressure outlets coupled with optional velocity-component constraints, as well as pressure pin controls to fully constrain the discrete PDE problem.

At the domain inlets, establishing accurate turbulent boundary conditions natively bridges physical intent with numerical stability. The prescribed turbulence intensity and the associated turbulence length scale are systematically converted into an intermediate eddy-viscosity ratio. This ratio directly calculates the requisite boundary value for the Spalart-Allmaras working variable $\tilde{\nu}$. Throughout the computation, to protect the solver from chaotic geometric fluctuations during intermediate design cycles, rigorous numerical floors, ceilings, and internal clipping functions are imposed to bind the working variable safely within physical limits.

4.5 Sensitivity Analysis and the Frozen Turbulence Assumption

With the forward state equations fully defined and penalized, the topology optimization algorithm requires the gradient (sensitivity) of the objective function with respect to the design variable. Because evaluating finite differences for every mesh element is computationally impossible, the continuous adjoint method is employed [30].

However, calculating the fully coupled, exact adjoint for turbulent flows is mathematically formidable. As demonstrated by Zymaris et al. [51], formulating the exact continuous adjoint for the Spalart-Allmaras model requires differentiating the highly non-linear production and destruction terms, accounting for the variations in the distance function, and managing the intense cross-coupling with the momentum equations. The complexity of these exact adjoints frequently introduces severe numerical stiffness and divergence.

To make the gradient calculation both numerically robust and computationally feasible, the primary optimization workflows rely on the **Frozen Turbulence Assumption**, frequently referred to in the literature as the Constant Eddy Viscosity (CEV) assumption [46, 34].

4.5.1 Mechanism of Frozen Turbulence

The frozen turbulence assumption simplifies the sensitivity analysis by decoupling the turbulence transport from the design update. During the adjoint formulation, the turbulent eddy viscosity field (ν_t) generated by the forward SA solver is assumed to be locally independent of infinitesimally small changes in the physical topology ($\bar{\gamma}$). Mathematically, this implies that during the calculation of the adjoint sensitivities:

$$\frac{\partial \nu_t}{\partial \bar{\gamma}} \approx 0 \quad (4.7)$$

By treating the eddy viscosity field as frozen (a spatially varying but constant parameter), the complex adjoints of the Spalart-Allmaras and Eikonal equations do not need to be solved. The optimizer only solves the standard fluid adjoint equations (utilizing the effective viscosity $\mu + \mu_t$) to determine how a change in topology affects the selected loss objective.

4.5.2 Justification and Limitations

The primary advantage of the frozen turbulence approach is not that it is universally better, but that it makes the optimization problem solvable within a reasonable numerical and implementation budget. Relative to a fully coupled turbulent adjoint, it offers greater numerical stability, a smaller implementation burden, and a materially lower computational cost per design iteration. This is why frozen turbulence remains common in industrial adjoint workflows and why it is used as the baseline in this thesis [30, 34].

The price of this robustness is that the gradients are approximate reduced derivatives. Because the algorithm ignores how an infinitesimal topology change would alter future turbulence production, the sensitivity can be inaccurate when the objective is strongly controlled by separated shear layers, stall-like behavior, or other regimes where the eddy viscosity changes sharply with geometry. Wu et al. [46] show that the CEV assumption can still achieve useful optimization goals in benign flows, but its sensitivity accuracy and adjoint convergence can deteriorate near stall and in unsteady settings. Zymaris et al. [51] further demonstrate, for SA-based continuous adjoints, that including the turbulence-equation adjoint improves derivative accuracy for loss functionals. The present thesis therefore defends the frozen assumption as a practical and verifiable approximation, not as an exact turbulence sensitivity model. Attempts to include the SA turbulence equation in the adjoint were explored during the implementation phase, but the resulting solves were not sufficiently stable for the thesis optimization campaign.

4.6 Optimization Problem Formulation

Using the frozen-turbulence sensitivities described above, the topology optimization task is formulated as a mathematical programming problem, where the optimizer must find a spatial distribution that minimizes a specific performance metric subject to physical and geometric constraints.

4.6.1 Objective Formulation

The distinction between active and passive domains introduced in Section 4.2 is essential for the objective definition. Only the active design region is meant to be optimized, while passive fluid extensions preserve inlet/outlet flow development and passive solid regions preserve fixed walls or separators. These non-design regions remain in the forward CFD problem, but they are excluded from domain-integrated optimization quantities when they would otherwise bias the design.

Dissipation-type objectives are common in density-based flow topology optimization because they penalize mechanical energy loss directly. In Brinkman-penalized formulations, the corresponding power-loss objective is naturally written as the sum of physical viscous dissipation and Darcy/Brinkman drag over the optimized design domain [47, 16]:

$$J_D(\mathbf{u}, \bar{\gamma}) = \int_{\Omega_D} \left[\frac{1}{2} \mu_{\text{eff}} (\nabla \mathbf{u} + \nabla \mathbf{u}^T) : (\nabla \mathbf{u} + \nabla \mathbf{u}^T) + \alpha(\bar{\gamma}) \mathbf{u} \cdot \mathbf{u} \right] d\Omega. \quad (4.8)$$

Here Ω_D denotes the optimized design domain used for the objective integration, excluding the fixed non-design inlet and outlet extensions, and $\mu_{\text{eff}} = \mu + \mu_t$ under the frozen-turbulence flow solve. The first term measures viscous deformation losses in the design region, while the Brinkman term penalizes residual velocity in low-density regions where $\alpha(\bar{\gamma})$ is large. Unlike a pure post-processing metric, J_D therefore contains an explicit local design-dependent contribution through $\alpha(\bar{\gamma})$.

For the pipe-bend and U-bend benchmarks considered in this thesis, the final engineering goal is instead expressed as a global pressure-loss objective: guide a prescribed inflow from inlet to outlet with the lowest possible hydraulic resistance. Because the outlet pressure is fixed to zero, this loss is represented directly by the average pressure on the physical inlet boundary, where the inlet velocity is prescribed,

$$J_p(p) = \frac{1}{|\Gamma_{\text{in}}|} \int_{\Gamma_{\text{in}}} p d\Gamma. \quad (4.9)$$

Here Γ_{in} is the inlet boundary of the fixed non-design extension, not an internal design-domain pressure plane. When the outlet pressure is used as the gauge reference, minimizing J_p is equivalent to minimizing the pressure loss required to drive the imposed flow rate through the topology. This is a boundary integral rather than a volume integral over Ω_D , and it does not include an explicit Brinkman-drag term in the objective. The Brinkman penalization still enters the optimization through the state equation in Eq. 4.3; changing $\bar{\gamma}$ changes the resistance field, the forward pressure solution, and the adjoint sensitivity. The pressure objective is therefore

less explicitly penalized than J_D , but it remains a valid hydraulic-loss objective for fixed-flow problems. Its sensitivity is mapped back only to the active design degrees of freedom through the same active-mask filter and transpose-filter machinery; consequently, the prescribed non-design inlet extension is used only to impose the velocity and measure the pressure and is not itself optimized.

Both objectives are used in the final benchmark chapter. The Bayat pipe-bend and U-bend cases use the inlet-pressure objective J_p , while the Yoon diffuser benchmark uses the dissipation objective J_D from Eq. 4.8. Physical viscous dissipation is therefore both an optimization objective for the diffuser and a validation metric for all extracted sharp geometries in Chapter 6.

4.6.2 Volume Constraint

Because the goal is to design a targeted manifold or pipe system, a strict volume constraint must be imposed. The volume fraction constraint is evaluated only over the prescribed volume mask Ω_V , rather than over the entire computational mesh:

$$g(\bar{\gamma}) = \frac{\int_{\Omega_V} \bar{\gamma} d\Omega}{V_{\text{mask}}} - V_{\text{max}} \leq 0 \quad (4.10)$$

where V_{mask} is the total volume of the masked region and V_{max} is the maximum allowable fluid volume fraction inside that region. Fixed non-design fluid domains, such as inlet and outlet extensions, are therefore not included in either the numerator or denominator of the volume constraint. They remain fluid for the flow solve, but they do not count as optimized fluid volume and cannot make the active design region satisfy the volume constraint artificially.

4.6.3 The Complete Continuous Optimization Problem

Combining the selected objective function, the governing equations, and the constraints, the continuous topology optimization problem for turbulent flow is formally stated as:

$$\begin{aligned} & \underset{\gamma(\mathbf{x}) \in [0,1]}{\text{minimize}} && J(\mathbf{u}, p, \bar{\gamma}) \\ & \text{subject to} && \begin{aligned} & \text{Modified RANS Equations (Eq. 4.3)} \\ & \text{Modified Spalart-Allmaras Equation (Eq. 4.4)} \\ & \text{Penalized Wall-Distance Equation (Eq. 4.5)} \\ & g(\bar{\gamma}) \leq 0 \end{aligned} \end{aligned} \quad (4.11)$$

where J denotes the selected loss objective, either the dissipation objective J_D in Eq. 4.8 or the inlet-pressure objective J_p in Eq. 4.9. The wall-distance constraint used for the final benchmark computations is the reciprocal penalized relaxed Eikonal equation in Eq. 4.5, with the wall distance recovered by Eq. 4.6.

To solve this constrained minimization problem, the Method of Moving Asymptotes (MMA) [39] is used. MMA is specifically designed for handling large-scale,

gradient-based topology optimization problems. It approximates the highly non-linear design space using convex asymptotes, relying on strict move limits which cap the maximum allowable change in the active raw design variable (γ) during any single optimization step.

4.7 Density Filtering and Projection

A well-known mathematical complication in density-based topology optimization is the lack of existence of a classical solution. If the design space is left unconstrained, the optimizer will tend to create infinitely thin micro-structures (checkerboarding). To guarantee a mesh-independent and manufacturable design, a minimum length scale must be enforced.

4.7.1 The Discontinuous Galerkin Helmholtz Filter

To prevent checkerboarding, a density filter is applied to the raw design variable generated by the optimizer. Because a continuous nodal definition causes the Brinkman penalty to interpolate smoothly across the interior of boundary elements, creating a semi-permeable gray gradient that allows mass leakage, the density field is discretized using piecewise constant elements (Discontinuous Galerkin, DG0) [9]. By assigning a single, uniform density value to the entire element, the structural boundaries remain perfectly sharp at the element edges.

Consequently, the standard continuous Helmholtz filter cannot be applied. The filter is mathematically adapted to account for the discontinuous boundaries utilizing an interior penalty method across the element facets (Γ_{int}). Standard volume diffusion is mathematically replaced by interior facet jump penalties:

$$\int_{\Gamma_{int}} r^2 \frac{\alpha_{dg}}{h_{avg}} \llbracket \tilde{\gamma} \rrbracket \cdot \llbracket v \rrbracket dS + \int_{\Omega} \tilde{\gamma} v dx = \int_{\Omega} \gamma v dx \quad (4.12)$$

where v is the test function, $\llbracket \cdot \rrbracket$ denotes the jump operator across a facet, h_{avg} is the average cell size, and $r = \frac{R}{2\sqrt{3}}$ is the scaled filter radius.

Active-Mask Normalization

To appropriately handle the severe boundary interactions near designated inlet and outlet buffers, this filtering is strictly active-mask-normalized. Rather than applying the filter blindly across the entire domain, the filtered output evaluates as:

$$\tilde{\gamma} = \frac{H(M\gamma)}{H(M)} \quad (4.13)$$

where H represents the Helmholtz PDE operator and M represents the isolated active design indicator mask. Crucially, during sensitivity analysis, the mathematical transpose of this normalized filter must be explicitly applied to map objective gradients back into the active DG0 design space. Listing 4.1 summarizes the forward and transpose operations without reproducing implementation details.

```

1      # Forward filter: only active cells influence the smoothed field.
2      denominator = helmholtz_solve(active_mask)
3      numerator = helmholtz_solve(active_mask * gamma_raw)
4      gamma_filtered = safe_divide(numerator, denominator)
5
6      # Projection and hard restoration of prescribed non-design cells.
7      gamma_projected = heaviside_projection(gamma_filtered, beta, eta)
8      gamma_effective = gamma_lower + (gamma_upper - gamma_lower) *
gamma_projected
9      gamma_effective[passive_fluid_cells] = 1.0
10     gamma_effective[passive_solid_cells] = 0.0
11
12     # Transpose filter: map adjoint sensitivities back to active cells.
13     dJ_dgamma_projected = dJ_dgamma_effective * (gamma_upper - gamma_lower)
14     dJ_dgamma_filtered = dJ_dgamma_projected * projection_slope(gamma_filtered
, beta, eta)
15     filter_rhs = active_mask * safe_divide(dJ_dgamma_filtered, denominator)
16     dJ_dgamma_raw = helmholtz_solve(filter_rhs)
17     dJ_dgamma_raw[inactive_cells] = 0.0

```

LISTING 4.1: Pseudocode for active-mask-normalized density filtering and transpose-gradient filtering.

4.7.2 Heaviside Projection

While the Helmholtz filter successfully imposes a minimum length scale, it inherently smooths the design field, resulting in a large gray transition zone between the solid and fluid regions. To recover a crisp, discrete topology, a continuous Heaviside projection is applied to the filtered variable [44, 22]:

$$\Pi_{\beta}(\tilde{\gamma}) = \frac{\tanh(\beta\eta) + \tanh(\beta(\tilde{\gamma} - \eta))}{\tanh(\beta\eta) + \tanh(\beta(1 - \eta))} \quad (4.14)$$

where $\eta \in [0, 1]$ is the projection threshold (typically 0.5), and β is a steepness parameter. As $\beta \rightarrow \infty$, the function approaches a strict Heaviside step function.

4.8 Algorithmic Workflow and Solver Implementation

Executing the optimization requires a highly robust discrete, iterative algorithm. To support the severe non-linearities of the turbulent regime, the solver architecture natively manages the segregated Spalart-Allmaras turbulence model, explicit physics decoupling via Picard iterations, and deeply hybridized forward flow routines.

4.8.1 Boundary Conditions and Turbulence Initialization

At the domain inlets, establishing accurate turbulent boundary conditions natively bridges physical intent with numerical stability. The prescribed turbulence intensity and the associated turbulence length scale are systematically converted into an intermediate eddy-viscosity ratio. This ratio directly calculates the requisite boundary

value for the Spalart-Allmaras working variable $\tilde{\nu}$. Throughout the computation, to protect the solver from chaotic geometric fluctuations during intermediate design cycles, rigorous numerical floors, ceilings, and internal clipping functions are imposed to bind the working variable safely within physical limits.

Furthermore, the solver dynamically handles pressure outlets coupled with optional velocity-component constraints, allowing the flow field to adapt to the evolving topology without artificial boundary blockages.

4.8.2 SA Numerical Controls and Hybrid Forward Flow Recovery

The SA transport solve is controlled through Picard-level coupling, carefully bounded turbulence variables, fixed topology-dependent SA and wall-distance penalties, and the hybrid forward-flow recovery strategy described in this subsection. Rather than relying solely on simple fixed-point decoupled sweeps, the forward fluid solver utilizes an advanced hybrid architecture integrating both Incremental Pressure Correction Scheme (IPCS) methods and monolithic SNES workflows.

During the optimization loop, a Picard iteration scheme strictly decouples the fluid and turbulence physics. The flow solver advances the momentum field using the eddy viscosity from the previous iteration, and this updated velocity is then passed back to the SA solver. During volatile intermediate design steps, the fluid framework initiates convection continuation strategies and leverages Stokes-Brinkman warm starts to slowly reintroduce inertial physics. If convergence stalls, the SNES solver autonomously steps through multiple recovery schedules (including dynamic tolerance reductions and altered line-search methods).

The optimization architecture supports non-converged acceptance, allowing the algorithm to passively absorb intermediate errors and maintain topological momentum. However, prior to the continuous adjoint assembly, stricter flow solves can be executed. When used, these strict solves confirm that the sensitivities, and specifically the frozen turbulent fields, are mathematically derived from a fully equilibrated state. Otherwise, relying on the accepted residual criteria serves as a necessary pragmatic compromise to maintain solver momentum in highly separated regimes.

4.8.3 Algorithm and Continuation Strategy

To ensure convergence to a physically meaningful binary design, the algorithm iterates over a strict continuation schedule. As the optimization progresses, the Brinkman penalty parameter (q) and the projection steepness (β) are iteratively increased, while the MMA move limits are proportionally tightened. The SA and wall-distance power penalties do not use this continuation in the final benchmark runs; their amplitudes and exponents remain fixed while the local penalty magnitude changes only through the evolving topology field.

The global optimization routine is outlined in Listing 4.2. The listing keeps the algorithmic sequence visible: continuation, density processing, wall-distance update, Picard-decoupled turbulent flow solve, frozen-turbulence adjoint assembly, transpose filtering, and MMA update.

```
1     for stage in continuation_schedule:
2         ramp_convexity = stage.ramp_convexity
3         beta = stage.projection_steepness
4         move_limit = stage.move_limit
5         set_fixed_sa_and_wall_penalties()
6
7         while not mma_converged and inner_iter < stage.max_iterations:
8             gamma_filtered = density_filter(gamma_raw, active_mask)
9             gamma_effective = project_and_restore_bounds(gamma_filtered, beta)
10            alpha = brinkman_ramp(gamma_effective, solid_resistance,
11                                   fluid_resistance, ramp_convexity)
12
13            wall_distance = update_wall_distance(
14                gamma_effective,
15                wall_penalty_strength,
16                wall_penalty_exponent,
17            )
18
19            for picard_iter in range(max_picard_iterations):
20                velocity, pressure = solve_flow(alpha, nu_t)
21                nu_tilde = solve_spalart_allmaras(
22                    velocity,
23                    wall_distance,
24                    sa_penalty_strength,
25                    sa_penalty_exponent,
26                )
27                nu_t = eddy_viscosity_from(nu_tilde)
28                if residuals_are_accepted():
29                    break
30
31                nu_t_frozen = freeze_eddy_viscosity(nu_t)
32                velocity, pressure = solve_final_state(alpha, nu_t_frozen)
33                adjoint = solve_frozen_turbulence_adjoint(velocity, pressure,
34                    nu_t_frozen)
35                dJ_dgamma_effective = assemble_design_sensitivity(adjoint,
36                    gamma_effective)
37                dJ_dgamma_raw = transpose_filter(dJ_dgamma_effective, active_mask)
38                gamma_raw[active_design_cells] = MMA.update(
39                    gamma_raw[active_design_cells], dJ_dgamma_raw, move_limit
40                )
```

LISTING 4.2: Pseudocode for the frozen-turbulence topology optimization loop.

4.9 Conclusion

This chapter has established the mathematical and algorithmic framework used for the turbulent topology optimization results in this thesis. By employing the Fictitious Domain Method with Brinkman momentum resistance [13, 20], the evolving structural topology is integrated into the fixed computational grid. The SA working-variable and reciprocal wall-distance equations are coupled to the topology through the power-law penalty, while the DG0 Helmholtz interior penalty filter [27] and

continuous Heaviside projection [44] enforce a minimum length scale and reduce mesh-dependent numerical artifacts.

The optimization is driven by the frozen-turbulence continuous adjoint [46, 34]. This assumption is defensible as a stable engineering approximation, but the thesis explicitly treats it as an approximation whose limitations must be checked through finite-difference sensitivity verification and independent CFD validation. Embedded within a Picard-decoupled flow architecture using hybrid IPCS/SNES recovery schemes, this methodology transforms the fluid dynamics engine from Chapter 3 into the automated design tool whose forward behavior is verified in Chapter 5 and whose optimized designs are evaluated in Chapter 6.

Chapter 5

Verification of the Spalart-Allmaras Solver for Turbulent Internal Flows

5.1 Introduction

Chapter 3 introduced the numerical implementation of the Spalart-Allmaras (SA) model in FEniCS. Before using this solver inside the topology optimization framework, the forward fluid solver is verified against independent Ansys Fluent calculations on two benchmark flows. The first benchmark is a fully developed turbulent channel, which isolates wall damping, turbulence production, and Picard coupling in a simple geometry. The second benchmark is a curved U-bend, which tests the same implementation under strong convection, streamline curvature, and separation.

This chapter focuses on the verification evidence rather than the full simulation setup. For the channel benchmark, the physical setup is listed in Table A.1, the FEniCS solver settings in Table A.2, and the mesh studies in Tables A.3 and A.4 of Section A.1. For the U-bend benchmark, the corresponding geometry, numerical settings, and mesh studies are documented in Tables A.5, A.6, A.7, and A.8 of Section A.2. The purpose here is to show that the FEniCS-SA implementation reproduces the relevant Ansys-SA flow behavior well enough to serve as the forward model for the topology optimization studies in Chapter 6.

5.2 Channel Verification Benchmark

To verify the fundamental correctness of the custom fluid and turbulence solvers, the code was first tested on a standard 2D turbulent channel flow. To cleanly isolate and verify turbulence production, wall damping, and Picard coupling without the artifacts of a developing boundary layer, this benchmark was formulated as a pressure-driven periodic channel rather than a conventional inlet-outlet duct.

The velocity field ($\mathbf{u}$) and Spalart-Allmaras working variable ($\tilde{\nu}$) were constrained periodically between the streamwise boundaries, while the top and bottom bound-

aries were treated as standard no-slip walls ($\mathbf{u} = 0, \tilde{\nu} = 0$). Flow was driven entirely by a fixed pressure drop from the upstream plane ($p = 2$ Pa) to the downstream plane ($p = 0$ Pa). Based on a reference bulk velocity of $U_{\text{ref}} = 18.5$ m/s and using the full channel height ($H = 2.0$ m) as the characteristic length scale, the solver was targeted to resolve a fully turbulent flow regime at a channel Reynolds number of $Re_H \approx 20,300$. The complete definition of the geometric, physical, and boundary-condition parameters is provided in Table A.1 of Section A.1.1 in Appendix A.

5.2.1 Mesh Resolution Requirements

Because the SA formulation is a low-Reynolds model that does not use empirical wall functions, it requires mathematical integration directly through the viscous sublayer. To satisfy this strict boundary-layer requirement, a custom structured inflation mesh was generated for the channel flow. Because the domain is a simple periodic rectangle, this anisotropic grid was constructed programmatically using NumPy arrays and converted into a FEniCS-compatible format via the `meshio` library. Based on the target bulk Reynolds number, this approach enforced a physical first-cell height of approximately 1.7×10^{-3} m, aiming for a non-dimensional wall distance of $y^+ \approx 1$ across the entire domain.

To verify that the numerical results were not dominated by spatial discretization, a grid-independence study was conducted. The full spatial verification is documented in Section A.1.3 of Appendix A. The mesh hierarchy is detailed in Table A.3, and visual confirmation of convergence is provided by the velocity profiles in Figure A.1.

5.2.2 FEniCS vs. Ansys Comparison


Prior to this comparative analysis, the Ansys baseline itself was verified for grid independence to ensure a robust reference (see Table A.4 and Figure A.2 in Section A.1.3 of Appendix A). The channel flow was then simulated in FEniCS using the decoupled Picard iteration method and compared against the Ansys Fluent baseline.

As shown in Figure 5.1, the qualitative velocity contours of the custom FEniCS implementation show good agreement with the Ansys Fluent baseline. For the quantitative comparison, velocity profiles are extracted from both solvers across the middle of the channel, i.e. at $x = 0.5L$, with y varied from the lower wall to the upper wall. The custom solver resolves the logarithmic region of the turbulent boundary layer and predicts the turbulent diffusion toward the channel center.


Because the channel is driven by a prescribed pressure difference, $\Delta p = 2.0$ Pa is identical by construction and is not a useful comparison metric. The retained global metric is therefore the bulk velocity,

$$U_{\text{bulk}} = \frac{1}{H} \int_0^H u_x(0.5L, y) dy, \quad (5.1)$$


5. VERIFICATION OF THE SPALART-ALLMARAS SOLVER FOR TURBULENT INTERNAL FLOWS



(A) Ansys Fluent velocity magnitude contour



(B) FEniCS velocity magnitude contour



(C) Velocity profiles

FIGURE 5.1: Verification of the 2D channel flow ($Re_H \approx 20,300$). The comparison includes qualitative velocity contours from Ansys Fluent and FEniCS, plus quantitative velocity profiles extracted at the streamwise mid-plane, $x = 0.5L$, while varying the wall-normal coordinate y .

and the corresponding plane-channel skin-friction coefficient,

$$C_f = \frac{\tau_w}{\frac{1}{2}\rho U_{\text{bulk}}^2} = \frac{\Delta p H}{\rho L U_{\text{bulk}}^2}. \quad (5.2)$$

Here u_x is the streamwise velocity component, H is the channel height, L is the streamwise periodic length, τ_w is the mean wall shear stress, ρ is the fluid density, and Δp is the imposed pressure difference. When the solver pressure is treated as kinematic pressure, the same expression is evaluated consistently with $\rho = 1$. Since Δp is fixed, C_f is not independent of U_{bulk} , but it is still useful because it converts the imposed forcing into the wall friction implied by the predicted flow rate.

The two error metrics are cross-code discrepancies, not separate errors for each solver. After interpolating the FEniCS and Ansys profiles onto the same line coor-

denates y_i , the pointwise difference is

$$e_i = u_{x,\text{FEniCS}}(y_i) - u_{x,\text{Ansys}}(y_i). \quad (5.3)$$

The maximum pointwise velocity discrepancy is normalized by the maximum Ansys profile velocity to avoid singular relative errors at the no-slip walls,

$$E_\infty = \frac{\max_i |e_i|}{\max_i |u_{x,\text{Ansys}}(y_i)|} \times 100\%, \quad (5.4)$$

and the relative profile error is evaluated as a line-integrated norm,

$$E_{L_2} = \left(\frac{\int_0^H |u_{x,\text{FEniCS}} - u_{x,\text{Ansys}}|^2 dy}{\int_0^H |u_{x,\text{Ansys}}|^2 dy} \right)^{1/2} \times 100\%. \quad (5.5)$$

The reported error metrics are therefore cross-code verification discrepancies between the FEniCS and Ansys profiles. They are assessed relative to the grid-independence evidence documented in Appendix A and to the accuracy required before using the FEniCS solver in the topology optimization loop. The quantitative errors and macroscopic metrics are summarized in Table 5.1.

TABLE 5.1: Quantitative comparison of flow metrics and extracted profile discrepancies for the 2D channel benchmark.

Metric	Ansys	FEniCS	Rel. diff.
Bulk velocity (U_{bulk}) [m/s]	[TBD]	18.472	[TBD]%
Friction coefficient (C_f)	[TBD]	5.857×10^{-3}	[TBD]%
Max pointwise velocity discrepancy (E_∞)	–	–	[TBD]%
Relative L_2 velocity discrepancy (E_{L_2})	–	–	[TBD]%

This agreement supports the verification that the core Navier-Stokes solver, relaxed Eikonal wall-distance field, and SA production/destruction terms behave correctly in a canonical wall-bounded turbulent flow.

5.3 U-Bend Verification Benchmark

Following the channel flow verification, the solver was tested on the highly convective separating U-bend flow evaluated in Chapter 2. This case is closer to the curved internal-flow topologies studied later in the thesis because the flow experiences strong streamline curvature and pressure-gradient-driven redistribution.

The verification used a 2D geometry corresponding to the symmetry mid-plane of the full 3D VMFL048 case. Since the FEniCS domain is mathematically a 2D planar approximation, it cannot resolve true 3D Dean secondary vortices. Instead, this benchmark validates the 2D symmetry-plane approximation and the solver’s ability to capture curvature-induced separation and momentum redistribution. The straight inlet and outlet legs are shorter than in the original VMFL048 Ansys mesh

to reduce meshing and computational effort; both the FEniCS and Ansys verification runs use this same short-leg U-bend geometry.

To drive the flow, a uniform prescribed velocity ($\mathbf{u}_{in} = 1.42$ m/s) was imposed at the inlet alongside a fully turbulent boundary value ($\tilde{\nu}_{in} = 9.34 \times 10^{-5}$ m²/s), while a zero relative static pressure ($p = 0$ Pa) was assigned to the outlet. While the domain is 2D, the 3D physical pipe diameter ($D = 0.028$ m) and radius of curvature ($R_c = 0.125$ m) were retained as the characteristic length scales. This preserves the target Reynolds number ($Re_D \approx 44,700$) and Dean number ($De \approx 15,000$) of the 3D Ansys reference case. The physical setup and boundary conditions are listed in Table A.5 of Section A.2.1 in Appendix A.

5.3.1 Mesh Resolution Requirements

The U-bend geometry features curved inner and outer arcs, complex corner transitions, and anisotropic boundary-layer inflation around non-orthogonal walls. To manage this geometric complexity while satisfying the $y^+ \approx 1$ resolution requirement, a custom mesh was generated using Gmsh [21]. The first element height adjacent to the wall was controlled to 1.20×10^{-5} m across the entire domain. The independence studies are documented in Section A.2.3 of Appendix A: the inflation parameters and element densities of the Gmsh hierarchy are listed in Table A.7, and the extracted velocity profiles confirming grid convergence at the bend exit are illustrated in Figure A.3.

5.3.2 FEniCS vs. Ansys Comparison

Having validated the spatial resolution of the FEniCS mesh, an equivalent grid-independence study was executed for the commercial Ansys Fluent baseline (detailed in Table A.8 and Figure A.4 of Section A.2.3 in Appendix A). With both solvers operating on spatially converged grids, an independent Ansys Fluent simulation was executed on the shortened 2D domain.

As shown in Figure 5.2, the global flow fields and separation regions observed in the qualitative velocity contours demonstrate visual agreement. To provide a more robust assessment, velocity data are extracted along a cross-sectional line probe positioned at the bend exit, spanning the full pipe diameter $D = 0.028$ m from the inner wall to the outer wall. The quantitative velocity profile computed by the custom FEniCS implementation traces the commercial Ansys results closely, adequately capturing the boundary-layer growth and the momentum shift toward the outer wall.

For the U-bend, the inlet velocity is prescribed rather than produced by a fixed pressure gradient, so pressure drop remains a meaningful global comparison metric. A channel-style friction coefficient is not retained because the curved geometry does not have a single fully developed wall-shear relation equivalent to the periodic channel. In FEniCS, the pressure-drop diagnostic is evaluated from boundary-measure-




FIGURE 5.2: Verification of the 2D U-bend flow ($Re_D \approx 44,700$, $De \approx 15,000$). The comparison includes qualitative velocity contours from Ansys Fluent and FEniCS, showing separation characteristics, plus quantitative velocity profiles extracted at the bend exit across the pipe diameter $D = 0.028$ m.

weighted mean pressures on the inlet and outlet facets,

$$\Delta p_{\text{FEniCS}} = \rho \left[\frac{1}{|\Gamma_{\text{in}}|} \int_{\Gamma_{\text{in}}} p_k \, d\Gamma - \frac{1}{|\Gamma_{\text{out}}|} \int_{\Gamma_{\text{out}}} p_k \, d\Gamma \right], \quad (5.6)$$

where p_k is the kinematic pressure solved by the IPCS formulation, $\rho = 997 \text{ kg/m}^3$, and $|\Gamma| = \int_{\Gamma} 1 \, d\Gamma$. The inlet and outlet integrals correspond to the marked FEniCS facets $\mathbf{ds}(2)$ and $\mathbf{ds}(3)$, respectively, so the reported value is the area-averaged inlet static pressure minus the area-averaged outlet static pressure, converted to pascals. The retained local comparison metrics are the same bend-exit profile discrepancies used for the channel,

$$E_{\infty} = \frac{\max_i ||\mathbf{u}_{\text{FEniCS}}(s_i)|| - ||\mathbf{u}_{\text{Ansys}}(s_i)||}{\max_i ||\mathbf{u}_{\text{Ansys}}(s_i)||} \times 100\%, \quad (5.7)$$

and

$$E_{L_2} = \left(\frac{\int_{\Gamma_D} (|\mathbf{u}_{\text{FEniCS}}| - |\mathbf{u}_{\text{Ansys}}|)^2 ds}{\int_{\Gamma_D} |\mathbf{u}_{\text{Ansys}}|^2 ds} \right)^{1/2} \times 100\%, \quad (5.8)$$

where Γ_D denotes the line segment across the bend-exit diameter and s is the corresponding arclength coordinate. Table 5.2 summarizes the quantitative error margins and the macroscopic pressure drop (Δp) across the 2D U-bend domain for both solvers. Since the FEniCS formulation solves for kinematic pressure, its pressure drop is converted to physical static pressure using $\rho = 997 \text{ kg/m}^3$ before comparison; equivalently, the fine Ansys pressure drop of 797.972 Pa corresponds to $\Delta p/\rho = 0.8004 \text{ m}^2/\text{s}^2$.

TABLE 5.2: Quantitative comparison of flow metrics and extracted profile discrepancies for the 2D U-bend benchmark.

Metric	Ansys	FEniCS	Rel. diff.
Pressure drop (Δp) [Pa]	797.972	[TBD]	TBD%
Max pointwise velocity discrepancy (E_∞)	–	–	[TBD]%
Relative L_2 velocity discrepancy (E_{L_2})	–	–	[TBD]%

5.4 Conclusion

The two verification cases establish the practical operating range of the custom FEniCS-SA solver before it is embedded in the topology optimization loop. The periodic channel confirms that the solver reproduces canonical wall-bounded turbulent behavior, while the U-bend checks the same implementation in a curved, separation-prone internal flow. The localized velocity-gradient agreement and interim pressure-drop metric indicate that the framework provides a credible forward-flow basis for the pressure-drop minimization objectives in the subsequent topology optimization campaigns.

These results do not remove the known limitations of a 2D SA model: secondary Dean vortices, anisotropic Reynolds stresses, and full three-dimensional separation physics remain outside the present solver scope. They do, however, show that the implemented SA equations, wall-distance calculation, and segregated coupling strategy are sufficiently consistent with Ansys Fluent to support the topology optimization methodology and results presented in Chapters 4 and 6.

Chapter 6

Turbulent Topology Optimization Results

6.1 Introduction

Building upon the FEniCS-SA solver verified in Chapter 5 and the topology optimization methodology developed in Chapter 4, this chapter evaluates the turbulent topology optimization solver on the pipe-bend and U-bend internal-flow benchmarks from Bayat et al. [11] and the diffuser benchmark from Yoon [47]. The Bayat cases use reciprocal wall-distance penalization, topology-dependent SA and wall-distance reaction terms, an average inlet-pressure objective, and a frozen-turbulence adjoint. The diffuser benchmark uses the same frozen-SA topology optimization machinery, but it minimizes the dissipation objective J_D .

The chapter first verifies the implemented frozen sensitivity on the same pipe-bend configuration used for optimization with a Dilgen-style coordinate finite-difference check on selected design cells and a directional Taylor check. The optimized pipe-bend, U-bend, and diffuser layouts are then interpreted through their projected density fields, velocity fields, convergence histories, and independent sharp-geometry CFD validation. This ordering keeps the derivative evidence close to the optimization result that uses it, while leaving the standalone forward-flow verification in Chapter 5.

Detailed pipe-bend geometry and FEniCS settings are listed in Tables B.2 and B.3 of Section B.3; the corresponding U-bend setup is listed in Tables B.5 and B.6 of Section B.4. The diffuser benchmark setup is listed in Tables B.8 and B.9 of Section B.5. The Ansys grid-independence studies for the extracted turbulent designs are reported in Tables B.4, B.7, and B.10. The corresponding laminar-design grid studies are retained in Appendix C, while the topology and SST $k - \omega$ metric comparisons are reported directly in the benchmark sections below.

6.2 Benchmark Cases and Evaluation Metrics

The three benchmarks cover curved internal flows and a diffuser with a prescribed outlet contraction. The pipe bend targets the $V_f = 0.25$, $Re = 5,000$ case from Bayat et al. [11], while the U-bend targets their $V_f = 0.27$, $Re = 5,000$ case with a 180° flow reversal. The diffuser targets the $Re = 3,000$ turbulent diffuser from Yoon [47], where the original solid mass usage of 70% corresponds to a maximum fluid fraction of $V_f = 0.30$ in the present density convention. Table 6.1 summarizes the operating point and analyses reported for each case.

TABLE 6.1: Benchmark cases evaluated in Chapter 6. Detailed physical and numerical parameters are listed in Appendix B.

Case	V_f	Re	Analyses reported
Pipe bend	0.25	5,000	FEniCS optimization; finite-difference and Taylor sensitivity verification; Ansys-SA and Ansys-SST $k - \omega$ CFD validation; laminar-design comparison
U-bend	0.27	5,000	FEniCS optimization; Ansys-SA and Ansys-SST $k - \omega$ CFD validation; laminar-design comparison
Diffuser	0.30	3,000	FEniCS optimization with J_D ; Ansys-SA and Ansys-SST $k - \omega$ CFD validation; laminar-design comparison

6.2.1 Optimization and Validation Metrics

Throughout the optimization process, the pipe-bend and U-bend benchmarks minimize the average inlet-pressure objective J_p defined in Eq. 4.9. This pressure is averaged over the physical inlet boundary of the fixed non-design extension, where the inlet velocity is prescribed. The outlet pressure is fixed to zero, so reducing J_p directly reduces the pressure level required to drive the prescribed inlet flow through the evolving topology. The optimization problem is still solved on a diffuse Brinkman domain, but once a design is extracted for Ansys Fluent, the Brinkman penalty is removed and the solid-fluid interface becomes a sharp wall.

The diffuser benchmark instead minimizes the dissipation objective J_D defined in Eq. 4.8. This matches the objective used by Yoon [47]: physical viscous dissipation plus Brinkman drag is integrated over the design domain and minimized subject to the volume constraint. This distinction matters because the diffuser uses prescribed parabolic velocity profiles at both inlet and outlet, so the most direct optimization loss is the domain-integrated energy loss rather than an inlet-pressure boundary average.

All three benchmark cases use the power-law SA sink, the power-law reciprocal wall-distance penalty, and the frozen-turbulence adjoint introduced in Chapter 4.

Bayat et al. [11] report corresponding final average inlet-pressure objectives for their standard $k - \varepsilon$ computations in Table 7. Their “conventional” method is the density-based turbulent topology optimization formulation without the proposed implicit wall-functions: solid regions are imposed through Brinkman/body-force penalization, including damping of k and ε toward zero in solid material, on the same

diffuse optimization mesh. It is therefore distinct from the body-fitted re-simulations with explicit wall-functions that they use for validation. Since the present thesis uses a wall-resolved SA model instead of $k\text{-}\varepsilon$, the conventional values from Bayat et al. are used only as scale comparisons in the individual Bayat benchmark discussions below. For the diffuser benchmark, the reference comparison is qualitative because Yoon [47] reports the laminar and turbulent diffuser layouts and post-processing behavior rather than a directly matching J_D table for the diffuser.

To bridge this gap, the primary validation metric is the physical viscous dissipation integrated over the full exported CFD domain, including the non-design inlet, outlet, and fixed regions that are retained in the final geometry,

$$\Phi_D = \int_{\Omega} \frac{1}{2} \mu_{\text{eff}} (\nabla \mathbf{u} + \nabla \mathbf{u}^T) : (\nabla \mathbf{u} + \nabla \mathbf{u}^T) d\Omega, \quad (6.1)$$

where Ω denotes the complete sharp-geometry flow domain exported to Ansys. The effective viscosity is $\mu_{\text{eff}} = \mu + \mu_t$ for turbulent evaluations and $\mu_{\text{eff}} = \mu$ for laminar reference solves. Furthermore, because pressure drop is the most direct and practical engineering interpretation of these flow losses, Δp is also reported from the Ansys validation runs. All two-dimensional dissipation quantities are evaluated per unit depth.

6.3 Pipe Bend Benchmark

The pipe-bend benchmark introduces a direct comparison with the turbulent topology optimization results of Bayat et al. [11]. While the reference study employs a standard $k\text{-}\varepsilon$ RANS model with wall functions, the present framework tackles the same problem using a wall-resolved SA model under the frozen-turbulence assumption. This comparison tests whether the fundamental topological solution, a streamlined and separation-avoiding bend, remains consistent across different turbulence closures.

6.3.1 Geometry and Boundary Conditions

The boundary conditions and geometric constraints replicate the setup shown in Fig. 10 of Bayat et al. [11], specifically targeting the $V_f = 0.25$ volume-fraction case. The inlet is placed on the left side of the non-design extension, the outlet is placed along the lower boundary, and the outlet pressure is fixed to zero. The optimization minimizes the average pressure on this prescribed inlet boundary required to drive the flow through the evolving bend. The local benchmark geometry used for the present implementation is shown in Figure 6.1.

6.3.2 Sensitivity Verification

Before interpreting the optimized topology, the frozen adjoint is checked directly on the pipe-bend configuration. This is the relevant local test for the final thesis setup because it uses the same topology-dependent SA and wall-distance penalties,




FIGURE 6.1: Geometry and boundary conditions for the pipe-bend benchmark from Bayat et al. [11]. The present evaluation targets the $V_f = 0.25$ case using the SA model.

inlet-pressure objective, and frozen-turbulence adjoint used during the optimization. No full turbulence-adjoint or wall-distance-adjoint comparison is made; the purpose is to compare the derivative supplied to MMA with objective changes obtained by re-solving the perturbed velocity-pressure state.

The check is performed at the initial pipe-bend design state, with base objective $J_0 = 1.8363$. In every perturbation, the SA working variable $\tilde{\nu}^0$ and reciprocal wall-distance field G^0 are held fixed. This matches the frozen-turbulence assumption used by the adjoint: the flow state is re-solved, but the turbulence and wall-distance updates are not differentiated.

Five active DG_0 design cells are sampled deterministically from seed 13. These sampled control volumes are labelled CV_1 – CV_5 in Table 6.2. For sampled cell i , the raw density vector γ is perturbed by $\pm \Delta h \mathbf{e}_i$, where $\mathbf{e}_i$ is the unit vector for that active design variable and $\Delta h = 10^{-6}$. The central finite-difference derivative is

$$S_i^{\text{FD}} = \frac{J_p(\gamma + \Delta h \mathbf{e}_i; \mathbf{u}^+, p^+, \tilde{\nu}^0, G^0) - J_p(\gamma - \Delta h \mathbf{e}_i; \mathbf{u}^-, p^-, \tilde{\nu}^0, G^0)}{2\Delta h}. \quad (6.2)$$

The frozen-adjoint value for the same cell is denoted by S_i^{fr} ; the superscripts FD and fr denote finite-difference and frozen-adjoint quantities, respectively. The absolute

and relative discrepancies are computed as

$$\Delta_i^{\text{abs}} = |S_i^{\text{fr}} - S_i^{\text{FD}}|, \quad \varepsilon_i^{\text{fr}} = \frac{\Delta_i^{\text{abs}}}{\max(|S_i^{\text{fr}}|, |S_i^{\text{FD}}|, 10^{-30})}. \quad (6.3)$$

TABLE 6.2: Coordinate sensitivity check for the frozen pipe-bend configuration. Values are taken from `SensitivityVerificationTable.tsv`.

Cell	Finite difference S_i^{FD}	Frozen adjoint S_i^{fr}	Absolute error Δ_i^{abs}	Relative error $\varepsilon_i^{\text{fr}}$
CV_1	-2.7041×10^{-5}	-2.7042×10^{-5}	4.1371×10^{-10}	1.5299×10^{-5}
CV_2	-2.1494×10^{-6}	-2.1491×10^{-6}	3.1679×10^{-10}	1.4738×10^{-4}
CV_3	-1.6641×10^{-4}	-1.6641×10^{-4}	3.6170×10^{-10}	2.1735×10^{-6}
CV_4	-1.7281×10^{-5}	-1.7281×10^{-5}	1.5118×10^{-12}	8.7488×10^{-8}
CV_5	2.3808×10^{-5}	2.3808×10^{-5}	8.1272×10^{-11}	3.4136×10^{-6}

The coordinate check shows close agreement between the frozen-adjoint derivative and the re-solved central finite differences. The largest absolute discrepancy is 4.14×10^{-10} , and the largest relative discrepancy is 1.47×10^{-4} for CV_2 . This relative error remains small and is amplified by the small magnitude of the CV_2 derivative, approximately 2.15×10^{-6} , which places the finite-difference objective change generated by a 10^{-6} perturbation close to roundoff.

TABLE 6.3: Directional Taylor check for the frozen pipe-bend configuration. Values are taken from `TaylorCheckLog.tsv`. The same five sampled cells as in Table 6.2 are perturbed together.

Step	Adjoint directional derivative	Central finite-difference directional derivative	Relative error	Symmetric Taylor remainder	Observed order
1.0×10^{-3}	9.6135×10^{-5}	9.6135×10^{-5}	1.5795×10^{-8}	2.1515×10^{-11}	—
3.0×10^{-4}	9.6135×10^{-5}	9.6135×10^{-5}	3.8379×10^{-9}	1.9358×10^{-12}	2.0003
1.0×10^{-4}	9.6135×10^{-5}	9.6135×10^{-5}	1.5386×10^{-8}	2.1494×10^{-13}	2.0006
3.0×10^{-5}	9.6135×10^{-5}	9.6135×10^{-5}	3.4657×10^{-8}	1.9096×10^{-14}	2.0107

The Taylor check provides a complementary directional test. Instead of perturbing one cell at a time, the test perturbs the same five sampled cells together along one fixed direction and checks whether the objective change follows the adjoint’s first-order prediction as the step is reduced. The direction signs for CV_1 – CV_5 are $(+, +, -, +, -)$, defining a combined perturbation of the coordinate-check cells rather than a separate set of locations. Table 6.3 shows that the central finite-difference directional derivatives agree with the adjoint value to a maximum relative error of 3.47×10^{-8} over the tested steps. The Taylor remainders decrease with approximately second-order rates, which is the expected behavior when the implemented gradient captures the local first-order objective change.

These tests therefore verify the implemented frozen-turbulence derivative for the pipe-bend setup used in the optimization. The conclusion is deliberately limited to that derivative: the SA working variable and reciprocal wall-distance field are held

fixed during the checks, matching the optimization adjoint; consequently, the result should not be interpreted as validation of a fully coupled turbulence sensitivity. Repeating the same diagnostic on the U-bend would test the same algebraic implementation path rather than a fundamentally different adjoint formulation. Since there is no independent reference field for this SA frozen derivative, a spatial $dJ_p/d\gamma$ contour is not used as a verification metric; the relevant evidence is the coordinate finite-difference agreement and the second-order Taylor behavior reported above.

6.3.3 Optimization Results

Figure 6.2 illustrates the projected density field $\bar{\gamma}$ and the corresponding velocity magnitude. Despite the differing turbulence models, the FEniCS-SA topology is expected to follow the same broad routing strategy as the k - ε reference: the primary high-speed path should bypass the sharpest corners to minimize momentum loss, while boundary-layer treatment can still affect the exact gray-scale density distribution near the walls.




FIGURE 6.2: FEniCS-SA optimization result for the pipe-bend benchmark at $V_f = 0.25$. Final objective function value $J_p = 0.2032$ Pa.

For the same $V_f = 0.25$, $Re = 5,000$ operating point, Bayat et al. [11] report $J_p = 0.4023$ Pa for their conventional k - ε topology optimization. This value is used only as a reference scale, not as a direct validation target, because the turbulence closure and near-wall treatment differ from the present wall-resolved SA formulation.

Figure 6.3 reports the convergence behavior that accompanies the final topology. The pressure objective is normalized by its initial value J_0 , while the volume constraint is plotted as a relative residual rather than as the raw volume fraction. Denoting the realized fluid volume fraction at MMA iteration k by $v_f^{(k)}$ and the prescribed target fraction by V_f , the plotted residual is

$$r_V^{(k)} = 100 \frac{v_f^{(k)} - V_f}{V_f}. \quad (6.4)$$

With this sign convention, the dashed line at $r_V = 0$ represents the active volume constraint. Values below this line are feasible and correspond to a design using slightly less fluid volume than the prescribed limit.




FIGURE 6.3: Optimization convergence history for the pipe-bend benchmark. The blue curve shows the normalized pressure objective J/J_0 , while the orange curve shows the relative volume-fraction residual r_V from Eq. (6.4).

The objective decreases rapidly during the first part of the optimization, reaching roughly one tenth of its initial value by the end of the run. The subsequent oscillations are expected for the continuation-based MMA solve, since changes in the interpolation and projection parameters alter the effective optimization landscape. The volume residual starts at zero because the initial uniform density is set equal to the prescribed volume fraction, $V_f = 0.25$. It then briefly becomes negative as the optimizer removes fluid while reducing the pressure objective, before returning close to the active constraint. The small negative peaks in the residual occur immediately after new continuation stages are introduced, most clearly near global MMA iterations 110 and 190. At these iterations the continuation update changes the mapping between the design variables and the projected material distribution, so the same design can momentarily sit below the target volume fraction until MMA rebalances the objective and constraint in the following iterations.

6.3.4 CFD Validation

The structural robustness of the extracted pipe-bend topology is verified in Table 6.4. Because the SST $k - \omega$ model is notably more sensitive to adverse pressure gradients

than SA, evaluating the extracted geometry under both closures serves as a rigorous stress test for the optimized shape.

TABLE 6.4: CFD validation metrics for the extracted pipe-bend topology.

Solver	Model	Φ_D [W/m]	Δp [Pa]	Comment
FEniCS diffuse design	SA	0.0147	0.2032	Optimization state
Ansys sharp geometry	SA	0.0153	0.1457	Cross-code comparison
Ansys sharp geometry	SST $k - \omega$	0.0166	0.1522	Turbulence-model sensitivity

Using the Ansys-SA result as the reference, the SST $k - \omega$ evaluation increases the pressure drop by approximately 4.5% and the viscous dissipation by approximately 8.5%. The SST $k - \omega$ model therefore predicts slightly higher losses, as expected for a closure that is more sensitive to adverse pressure gradients, but the changes remain modest. This suggests that the optimized pipe-bend geometry is robust under a change of turbulence closure, rather than being tuned only to the SA model.


6.3.5 Comparison with Laminar Design

The laminar reference uses the same pipe-bend mesh, design-domain restrictions, inlet and outlet treatment, objective type, and volume fraction, but replaces the SA turbulent state equations by the laminar flow formulation and is optimized at $Re = 1$. The expected distinction is that a laminar bend can minimize wall length and viscous shear in the creeping-flow limit, but may produce a sharper or less separation-aware turn when evaluated at the target turbulent Reynolds number. The comparison is therefore interpreted under identical turbulent operating conditions, rather than by comparing the laminar and turbulent optimization objectives directly.

Figure 6.4 compares the laminar and turbulent optimized topologies before sharp-geometry extraction. The laminar-design mesh-independence study supporting the SST $k - \omega$ result is reported in Table C.1 of Appendix C.

Figure 6.5 shows the corresponding SST $k - \omega$ velocity contours through the extracted laminar and turbulent pipe-bend designs. These contours provide the flow-field evidence behind the scalar metrics: sharper turns and corners in the laminar design are expected to promote stronger local separation and momentum loss when the geometry is operated at the turbulent Reynolds number.


The final SST $k - \omega$ comparison in Table 6.5 evaluates the extracted laminar and turbulent pipe-bend designs under identical turbulent operating conditions. A lower pressure drop and lower viscous dissipation for the turbulent design would support the need for a turbulent topology optimization framework rather than a laminar surrogate.



(A) Projected density field $\bar{\gamma}$: laminar-optimized pipe-bend topology ($Re = 1$).

(B) Projected density field $\bar{\gamma}$: turbulent-optimized pipe-bend topology.

FIGURE 6.4: Laminar and turbulent topology comparison for the pipe-bend benchmark.



(A) Velocity magnitude field u : laminar-optimized design.

(B) Velocity magnitude field u : turbulent-optimized design.

FIGURE 6.5: Velocity-contour comparison for the extracted pipe-bend designs under identical turbulent operating conditions (SST $k - \omega$ model).

TABLE 6.5: Performance comparison of the extracted pipe-bend topologies under identical turbulent operating conditions.

Design origin	Model	Φ_D [W/m]	Δp [Pa]	Relative performance
Laminar TO, $Re = 1$	SST $k - \omega$	0.0313	0.5444	Baseline
Turbulent TO, SA	SST $k - \omega$	0.0166	0.1522	47.0% lower Φ_D ; 72.0% lower Δp

6.4 U-Bend Benchmark

The U-bend is the companion benchmark to the pipe bend. It uses the same frozen-SA topology optimization framework, but the geometry forces a 180° return path between two ports on the same side of the domain. This makes the case useful for

checking whether the optimizer can form a smooth turbulent return duct without relying on the shorter 90° routing available in the pipe-bend case.

6.4.1 Geometry and Boundary Conditions

The geometry and boundary conditions follow Fig. 11 and Table 4 of Bayat et al. [11]. The square design region has side length L , with a non-design lead duct of length $0.2L$ on the left. The inlet and outlet ports both have height $0.2L$; the top port acts as the inlet and the lower port acts as the pressure outlet. A fixed internal separator bar of thickness $0.10L$ extends from the lead duct into the design domain, forcing the flow to turn around its rounded tip. The target volume fraction is $V_f = 0.27$, and the Reynolds number is $Re = 5,000$ based on the port height and maximum inlet velocity. Figure 6.6 shows the corresponding local geometry and boundary-condition layout used for this benchmark.




FIGURE 6.6: Geometry and boundary conditions for the U-bend benchmark from Bayat et al. [11].

6.4.2 Optimization Results

Figure 6.7 reports the final projected density and velocity-magnitude field for the U-bend. The main feature to assess is whether the design forms a rounded return passage around the separator tip, rather than a short path with tight local curvature. Because the objective is average inlet pressure, a successful layout reduces the

pressure required to maintain the imposed inlet flow while respecting the prescribed volume fraction.




FIGURE 6.7: FEniCS-SA optimization result for the U-bend benchmark at $V_f = 0.27$. Final objective function value $J_p = [TBD]$.

For the matching $V_f = 0.27$, $Re = 5,000$ U-bend case, the conventional $k-\varepsilon$ result reported by Bayat et al. [11] is $J_p = 0.8327$. As for the pipe bend, this value provides a reference scale only; the comparison remains qualitative because the present solver uses a different turbulence model and near-wall treatment.

The U-bend convergence histories are collected in Figure 6.8; the two panels separate the inlet-pressure objective history in Figure 6.8a from the enforced volume-fraction history in Figure 6.8b.




FIGURE 6.8: Optimization convergence histories for the U-bend benchmark.

6.4.3 CFD Validation

The extracted U-bend topology is validated with the same sharp-geometry workflow as the pipe bend. The Ansys SST $k - \omega$ solve tests whether the design remains

6. TURBULENT TOPOLOGY OPTIMIZATION RESULTS

robust under a two-equation model that is more sensitive to separation and adverse pressure gradients, while the Ansys-SA solve checks cross-code consistency for the same turbulence closure family. The resulting validation metrics are summarized in Table 6.6.

TABLE 6.6: CFD validation metrics for the extracted U-bend topology.

Solver	Model	Φ_D [W/m]	Δp [Pa]	Comment
FEniCS diffuse design	SA	[TBD]	[TBD]	Optimization state
Ansys sharp geometry	SA	[TBD]	[TBD]	Cross-code comparison
Ansys sharp geometry	SST $k - \omega$	[TBD]	[TBD]	Turbulence-model sensitivity

6.4.4 Comparison with Laminar Design

The laminar reference uses the same U-bend mesh, inlet and outlet ports, separator geometry, objective type, and volume fraction, but it is optimized at $Re = 1$ without the SA turbulence model. The extracted laminar and turbulent layouts are then evaluated under the same target turbulent operating conditions. This comparison is expected to show whether the laminar optimizer creates an overly tight return path around the separator tip, while the turbulent optimizer allocates fluid volume to reduce separation and momentum loss through the 180° turn.

Figure 6.9 compares the laminar and turbulent optimized U-bend topologies. The laminar-design mesh-independence study supporting the SST $k - \omega$ result is reported in Table C.2 of Appendix C.





FIGURE 6.9: Laminar and turbulent topology comparison for the U-bend benchmark.

Figure 6.10 shows the SST $k - \omega$ velocity contours for the extracted U-bend designs. In this case the comparison is especially useful near the separator tip, where a laminar-optimized shortcut can create a tighter return path, stronger adverse-

pressure-gradient effects, and larger separated or low-momentum regions than the turbulent design.



(A) Velocity magnitude field u : laminar-optimized design.

(B) Velocity magnitude field u : turbulent-optimized design.

FIGURE 6.10: SST $k - \omega$ velocity-contour comparison for the extracted U-bend designs under identical turbulent operating conditions.

Table 6.7 compares the extracted laminar and turbulent U-bend topologies under identical turbulent operating conditions. If the turbulent design lowers Δp and Φ_D under the SST $k - \omega$ evaluation, the result supports the thesis claim that high-Reynolds topology optimization should include turbulence directly in the design loop.

TABLE 6.7: Performance comparison of the extracted U-bend topologies under identical turbulent operating conditions.

Design origin	Model	Φ_D [W/m]	Δp [Pa]	Relative performance
Laminar TO, $Re = 1$	SST $k - \omega$	[TBD]	[TBD]	Baseline
Turbulent TO, SA	SST $k - \omega$	[TBD]	[TBD]	[TBD]% vs. laminar

6.5 Diffuser Benchmark

The diffuser benchmark follows the third numerical example of Yoon [47]. It differs from the Bayat benchmarks in two important ways. First, the inlet and outlet are both prescribed by parabolic velocity profiles, with the outlet placed only on the center third of the right boundary. Second, the optimization objective is the dissipation objective J_D , not the average inlet-pressure objective J_p . The benchmark therefore tests whether the frozen-SA implementation can reproduce the expected turbulent diffuser behavior while optimizing the same energy-loss functional used in the reference formulation.

6.5.1 Geometry and Boundary Conditions

The geometry and boundary conditions follow Fig. 20 of Yoon [47]. The design domain is a square of side length $L = 1$ m. A parabolic inflow enters through the full left boundary with maximum speed $U_{\max,\text{in}} = 3$ m/s, while a parabolic outflow is prescribed on the center third of the right boundary with maximum speed $U_{\max,\text{out}} = 9$ m/s. The remaining external boundaries are treated as no-slip walls. The Reynolds number is $Re = 3,000$ based on $U_{\max,\text{in}}L\rho/\mu$, and the present density convention converts Yoon’s 70% solid mass usage into a maximum fluid volume fraction of $V_f = 0.30$. Figure 6.11 shows the corresponding local geometry and boundary-condition layout used for this benchmark.




FIGURE 6.11: Geometry and boundary conditions for the diffuser benchmark from Yoon [47].

6.5.2 Optimization Results

Figure 6.12 reports the final projected density and velocity-magnitude field for the diffuser. The main feature to assess is whether the turbulent design expands the flow near the inlet and the contracted outlet, rather than forming only the simple cone-like diffuser associated with the laminar reference in Yoon’s Fig. 21. Because the objective is J_D , the optimizer directly penalizes viscous deformation loss and residual Brinkman drag in the diffuse design domain.

Yoon [47] reports that the turbulent diffuser design develops enlarged inlet and outlet regions, while the laminar solution forms a more cone-shaped diffuser. The present FEniCS-SA result is interpreted against this qualitative behavior rather than against a direct numerical objective target, because the present implementation uses a wall-resolved frozen-SA solve, active-mask filtering, and Heaviside projection.




FIGURE 6.12: FEniCS-SA optimization result for the diffuser benchmark at $V_f = 0.30$. Final objective function value $J_D = 6.2274 \times 10^5$ W/m.

The diffuser convergence histories are collected in Figure 6.13; the two panels separate the dissipation objective history in Figure 6.13a from the enforced volume-fraction history in Figure 6.13b.




FIGURE 6.13: Optimization convergence histories for the diffuser benchmark.

6.5.3 CFD Validation

The extracted diffuser topology is validated with the same sharp-geometry workflow as the pipe bend and U-bend. The Ansys SST $k - \omega$ solve tests whether the diffuser remains robust under a two-equation model, while the Ansys-SA solve checks cross-code consistency for the same turbulence closure family. Because the optimization objective is J_D , the validation emphasizes both physical viscous dissipation and static pressure drop. The resulting validation metrics are summarized in Table 6.8.

6. TURBULENT TOPOLOGY OPTIMIZATION RESULTS

TABLE 6.8: CFD validation metrics for the extracted diffuser topology.

Solver	Model	Φ_D [W/m]	Δp [Pa]	Comment
FEniCS diffuse design	SA	TBD	TBD	Diffuse diagnostic; $J_D = TBD$ W/m
Ansys sharp geometry	SA	[TBD]	[TBD]	Cross-code comparison
Ansys sharp geometry	SST $k - \omega$	[TBD]	[TBD]	Turbulence-model sensitivity

6.5.4 Comparison with Laminar Design

The laminar reference uses the same diffuser mesh, inlet and outlet profile shapes, objective type, and volume fraction, but it is optimized at $Re = 1$ without the SA turbulence model. The extracted laminar and turbulent layouts are then evaluated under the same target turbulent operating conditions. This comparison is expected to show whether the laminar cone-like diffuser remains competitive once turbulent eddy viscosity and the outlet contraction are considered.

Figure 6.14 compares the laminar and turbulent optimized diffuser topologies. The expected distinction follows Yoon [47]: the laminar design tends toward a cone-like diffuser, while the turbulent design should allocate additional fluid area near the inlet and outlet where eddy-viscosity effects increase the local loss. The laminar-design mesh-independence study supporting the SST $k - \omega$ result is reported in Table C.3 of Appendix C.






FIGURE 6.14: Laminar and turbulent topology comparison for the diffuser benchmark.

Figure 6.15 shows the corresponding SST $k - \omega$ velocity contours through the extracted diffuser designs. The contour comparison complements the topology plot by indicating whether the cone-like laminar layout accelerates or separates near the contraction, while the turbulent layout distributes the expansion more smoothly near the inlet and outlet.

Table 6.9 compares the extracted laminar and turbulent diffuser topologies under

(A) Velocity magnitude field u : laminar-optimized design.(B) Velocity magnitude field u : turbulent-optimized design.FIGURE 6.15: SST $k - \omega$ velocity-contour comparison for the extracted diffuser designs under identical turbulent operating conditions.

identical turbulent operating conditions. Because the diffuser was optimized with J_D , lower Φ_D is the primary performance indicator, while the static pressure drop remains the most direct hydraulic interpretation of the extracted sharp geometry.

TABLE 6.9: Performance comparison of the extracted diffuser topologies under identical turbulent operating conditions.

Design origin	Model	Φ_D [W/m]	Δp [Pa]	Relative performance	perfor- mance
Laminar TO, $Re = 1$	SST $k - \omega$	1.1831×10^4	3.4299×10^4	Baseline	
Turbulent TO, SA	SST $k - \omega$	[TBD]	[TBD]	[TBD]% vs. laminar	

6.6 Conclusion

This chapter centers the final thesis evidence on the pipe-bend and U-bend benchmarks from Bayat et al. [11] and the diffuser benchmark from Yoon [47]. The sensitivity verification is performed on the same penalized wall-distance, topology-penalized SA, frozen-turbulence formulation used for the optimization; therefore, the coordinate finite-difference and Taylor checks directly test the derivative supplied to MMA. The pipe-bend and U-bend results test the same framework on two distinct curved-flow layouts, while the diffuser tests the dissipation objective J_D under prescribed inlet and outlet velocity profiles.

The independent CFD validation using Ansys Fluent is therefore interpreted as a sharp-geometry check of the optimized designs, not merely as another internal FEniCS result. Testing the extracted topologies with both SA and SST $k - \omega$ quantifies how much of the optimized performance survives a change in turbulence closure. Together with the laminar comparisons, this chapter supports the thesis

6. TURBULENT TOPOLOGY OPTIMIZATION RESULTS

argument that high-Reynolds turbulence physics must be included directly in the optimization loop when designing viable internal-flow topologies.

Chapter 7

Conclusion and Future Work

7.1 Conclusion

The primary objective of this thesis was to develop a stable, continuous turbulent topology optimization framework within FEniCS, verify its forward-flow and sensitivity behavior, and critically evaluate the physical credibility of its resulting designs. The final thesis structure deliberately separates methodology from evidence: Chapters 3 and 4 define the solver and optimization methods, Chapter 5 verifies the standalone SA flow solver, and Chapter 6 evaluates the pipe-bend and U-bend topology optimization benchmarks from Bayat et al. [11] with independent CFD re-analysis.

Addressing the **first research question**, the theoretical limitations and computational advantages of the Spalart-Allmaras (SA) turbulence model were assessed. While its one-equation, wall-resolved formulation makes it highly tractable for continuous adjoint frameworks, its reliance on the linear Boussinesq approximation fundamentally limits its accuracy in resolving anisotropic Reynolds stresses—particularly in regimes dominated by high streamline curvature and separation.

Regarding the **second research question**, the numerical integration of this highly non-linear SA model into the FEniCS finite-element environment was successfully achieved. By deploying a decoupled Picard iteration scheme, an Incremental Pressure Correction Scheme (IPCS) for the Navier-Stokes equations, and a relaxed Eikonal equation for wall-distance computation, the forward fluid state was robustly driven to equilibrium. The channel and U-bend verification cases in Chapter 5 show that this implementation reproduces the relevant Ansys-SA flow behavior with sufficient accuracy for the subsequent optimization campaign.

Building on this foundation, the mathematical framework for topology optimization was constructed. The implementation confirmed that the "frozen turbulence" adjoint assumption, which treats the turbulent eddy viscosity and reciprocal wall-distance fields as fixed coefficients during gradient assembly, provides a stable and computationally efficient optimization engine. This assumption is approximate, and the pipe-bend finite-difference check in Section 6.3.2 of Chapter 6 provides a Dilgen-style coordinate diagnostic of the implemented frozen derivative for the configuration

used in the final optimization campaign.

Finally, answering the **third research question**, the framework was subjected to strict cross-code CFD validation. The pipe bend served as the primary verification and optimization case, while the U-bend provided a companion curved-flow topology benchmark with stronger flow reversal. In both cases, the diffuse FEniCS-SA topology is extracted as a sharp geometry and re-analyzed in Ansys Fluent. The laminar-design comparisons further support the engineering motivation for turbulent optimization by showing why laminar-optimized layouts can perform poorly at high Reynolds numbers.

Re-evaluating the extracted geometries using the Ansys SST $k - \omega$ model quantifies the sensitivity of the optimized designs to the underlying turbulence physics. This confirms that topologies optimized under isotropic, one-equation assumptions must be checked carefully when subjected to secondary flows, curvature effects, and anisotropic shear layers captured more strongly by two-equation models.

7.2 Limitations

An objective evaluation of this methodology reveals several important limitations that must be acknowledged:

1. **The Frozen Adjoint Approximation:** By ignoring the derivative of the eddy viscosity with respect to the design field, the optimizer cannot mathematically anticipate future turbulence production. In flows dominated by massive separation, this can blind the optimizer to the downstream hydrodynamic consequences of upstream geometric changes, potentially leading to sub-optimal local minima.
2. **Geometry Extraction Uncertainties:** The transition from a continuous, Brinkman-penalized density field ($\gamma \in [0, 1]$) to a Boolean CAD geometry introduces inherent ambiguity. The specific iso-value threshold chosen to extract the fluid volume dictates the final wall placement. Consequently, the performance discrepancies observed between FEniCS and Ansys are inextricably linked to this manual geometry extraction process, making it difficult to isolate pure solver physics from thresholding effects.
3. **Two-Dimensional Simplification:** The benchmark geometries evaluated in this thesis were restricted to 2D planar domains. Because turbulent phenomena—particularly curvature-induced secondary flows like Dean vortices and turbulent energy cascades—are inherently three-dimensional, the 2D assumption artificially suppresses complex mixing mechanisms that would inevitably occur in physical prototypes.
4. **Lack of Experimental Validation:** While independent cross-code verification against commercial finite-volume solvers strongly corroborates the numerical implementation, it is not a substitute for physical validation. The true operational performance of these algorithmically generated turbulent manifolds

requires future validation against experimental pressure-drop and flow-field measurements.

7.3 Future Work

The verified turbulent fluid mechanics framework developed in this thesis establishes a mathematically stable foundation for several advanced avenues of future research:

7.3.1 Conjugate Heat Transfer (CHT) and Thermal Management

The ultimate engineering application for high-Reynolds topology optimization is not merely fluid routing, but active thermal management. Historically, the immense computational cost of resolving Conjugate Heat Transfer (CHT) has forced optimization frameworks to rely on laminar approximations (e.g., natural convection as demonstrated by Alexandersen et al. [2]) or simplified flow proxies such as Darcy models [49]. Even when high-resolution, density-based solvers are successfully deployed for microelectronics cooling, they typically circumvent the complexities of turbulent momentum transport [10].

However, the necessity of capturing true turbulent convection in high-performance heat sinks has driven recent advancements in the literature. Emerging efforts have introduced turbulence into thermal topology optimization using extended finite element (XFEM) level-set approaches with the SA model [29], computationally reduced multi-layer thermofluid approximations [48], and standard density-based frameworks coupled to two-equation $k - \epsilon$ closures [38].

With the turbulent momentum and SA transport equations now stably implemented and verified within the current FEniCS architecture, this framework is perfectly positioned to contribute to this cutting-edge research front. Extending the current methodology to CHT will require implementing a turbulent Prandtl number (Pr_t) approach to model the turbulent heat flux vector, alongside deriving the corresponding thermal adjoint sensitivities to minimize peak temperatures or maximize overall heat transfer coefficients. Furthermore, the capacity for optimizing two-fluid heat exchangers—currently dominated by laminar assumptions [24]—could be significantly enhanced by adapting this turbulent solver.

7.3.2 Three-Dimensional Topologies

As noted in the limitations, turbulent secondary flows are inherently 3D phenomena. To fully evaluate the geometric limits of the turbulence models and design physically realizable manifolds, the current 2D FEniCS implementation must be extended to three dimensions. While the governing equations and adjoint derivations remain mathematically identical, the computational cost of resolving the boundary layers ($y^+ \leq 1$) across a 3D topology optimization grid will require significant advancements in parallel processing and highly efficient linear solvers (e.g., Algebraic Multigrid preconditioners).

7.3.3 Adjoint Model Hierarchies

This thesis deliberately uses the frozen-turbulence adjoint as the robust baseline. Future frameworks should investigate stabilized ways of reintroducing selected turbulence-production sensitivities, especially in separated flows where the Constant Eddy Viscosity assumption is least accurate. Higher-fidelity two-equation models, such as $k - \omega$, should also be investigated directly within the continuous adjoint topology optimization loop to mitigate the curvature limitations identified in Chapter 6.

Appendices

Appendix A

Configuration and Grid Independence of the Spalart-Allmaras Verification Cases

To ensure complete scientific reproducibility and to prevent the disruption of the narrative flow in Chapter 3, all comprehensive configuration matrices and spatial grid independence studies for the fluid benchmarks are documented in this appendix.

A.1 Verification Benchmark 1: Channel Flow

A.1.1 Physical and Geometric Parameters

The first verification benchmark consists of a pressure-driven periodic 2D channel. To align with fundamental turbulence research conventions, the characteristic length is defined as the full channel height (H), and the Reynolds number (Re_H) is computed strictly based on this geometric parameter rather than the hydraulic diameter. The complete physical setup is detailed in Table A.1.

A.1.2 Numerical Simulation Parameters

The channel flow was resolved using an Incremental Pressure Correction Scheme (IPCS) embedded within a Picard fixed-point iteration loop. The exact numerical solver configurations, convergence tolerances, and relaxation factors utilized in FEniCS are summarized in Table A.2.

A.1.3 Spatial Verification and Grid Independence

To guarantee that the numerical results were independent of the spatial discretization, a formal grid independence study was conducted. Three distinct boundary-layer inflated meshes (Coarse, Medium, and Fine) were evaluated. The geometric

TABLE A.1: Geometric, physical, and boundary-condition parameters for the periodic channel flow benchmark.

Parameter	Channel flow setting
Benchmark type	Pressure-driven, streamwise-periodic 2D channel
Domain dimensions	Length $L = 2.0$ m; full channel height $H = 2.0$ m
Characteristic length	Full channel height $H = 2.0$ m
Reference velocity	$U_{\text{ref}} = 18.5$ m/s, utilized to define the target Reynolds number
Kinematic viscosity	$\nu = 1.81818 \times 10^{-3}$ m ² /s
Reynolds number	$Re_H \approx 2.03 \times 10^4$
Body force	$\mathbf{f} = (0, 0)$ m/s ²
Driving condition	Fixed pressure drop from $p = 2$ Pa to $p = 0$ Pa across the periodic streamwise direction
Velocity boundary conditions	Periodic velocity between inlet/outlet planes; no-slip on top and bottom walls
Pressure boundary conditions	$p = 2$ Pa at the upstream pressure plane; $p = 0$ Pa at the downstream pressure plane
SA boundary conditions	Periodic $\tilde{\nu}$ between inlet/outlet planes; $\tilde{\nu} = 0$ at walls
Initial fields	$\mathbf{u}_0 = (18.5, 0)$ m/s, $p_0 = 2$ Pa, $\tilde{\nu}_0 = 5.0 \times 10^{-3}$ m ² /s
Verification mesh	Fine_FEniCS ; first layer $\Delta y_1 = 1.709 \times 10^{-3}$ m

TABLE A.2: Numerical simulation parameters for the FEniCS steady Spalart-Allmaras channel benchmark.

Parameter	Symbol	Setting
Velocity space	$\mathbf{u}$	Vector CG2, periodic in streamwise direction
Pressure space	p	Scalar CG1
SA working-variable space	$\tilde{\nu}$	Scalar CG1, periodic in streamwise direction
Quadrature degree	–	2
Maximum outer Picard steps	$N_{\text{Picard,max}}$	200
SA solves per Picard step	–	1
Outer velocity tolerance	$\epsilon_{\mathbf{u}}$	1.0×10^{-6}
Outer pressure tolerance	ϵ_p	1.0×10^{-6}
Outer SA tolerance	$\epsilon_{\tilde{\nu}}$	1.0×10^{-6}
IPCS pseudo-time step	Δt	5.0×10^{-3} s
Maximum IPCS steps per Picard step	–	500
IPCS velocity tolerance	–	1.0×10^{-6}
IPCS pressure tolerance	–	1.0×10^{-6}
Velocity relaxation factor	α_u	0.3
Pressure relaxation factor	α_p	0.3
SA relaxation factor	$\alpha_{\tilde{\nu}}$	0.3
SA lower bound	$\tilde{\nu}_{\text{min}}$	1.0×10^{-12}
Wall-distance equation	–	Relaxed Eikonal, $\sigma_w = 0.1$, $G_0 = 20$
Wall-distance Newton settings	–	Relative tolerance 1.0×10^{-8} ; absolute tolerance 1.0×10^{-10} ; maximum 80 iterations; relaxation 0.5
Linear solver strategy	–	IPCS blocks: direct MUMPS; SA transport: default solver

A. CONFIGURATION AND GRID INDEPENDENCE OF THE SPALART-ALLMARAS VERIFICATION CASES

properties, node counts, and resolution parameters of this mesh hierarchy are explicitly detailed in Table A.3. Crucially, the first-layer height was kept constant across all grids to strictly enforce the $y^+ \approx 1$ requirement, while the bulk and inflation-layer resolutions were systematically refined.

TABLE A.3: Wall-resolved channel mesh hierarchy used for the FEniCS grid-independence study. The first-layer height is kept fixed to maintain $y^+ \approx 1$, while the streamwise resolution, wall-normal resolution, and inflation-layer resolution are refined.

Parameter	Coarse_FEniCS	Medium_FEniCS	Fine_FEniCS
Streamwise cells	10	20	40
Wall-normal cells	42	51	68
First-layer height Δy_1 [m]	1.709×10^{-3}	1.709×10^{-3}	1.709×10^{-3}
Minimum wall-normal spacing [m]	1.709×10^{-3}	1.709×10^{-3}	1.709×10^{-3}
Maximum wall-normal spacing [m]	7.707×10^{-2}	8.454×10^{-2}	7.589×10^{-2}
Inflation layers per wall	10	14	18
Growth ratio	1.45	1.35	1.25
Vertices	473	1,092	2,829
Triangular elements	840	2,040	5,440
Bulk velocity U_{bulk} [m/s]	16.923	18.489	18.472
Relative change U_{bulk} , Coarse–Medium [%]	–	8.47	–
Relative change U_{bulk} , Medium–Fine [%]	–	–	0.09

Visual confirmation of this spatial convergence is provided by the wall-normal velocity profiles plotted in Figure A.1. In each case, the profile is extracted at the streamwise mid-plane, $x = 0.5L$, with the wall-normal coordinate y varied across the full channel height. The **Fine_FEniCS** grid was consequently selected for the final verification against the Ansys Fluent baseline, while the coarser grids demonstrate that the wall-normal velocity profile is already asymptotically converged.

Furthermore, to ensure a rigorous cross-code validation, an equivalent grid independence study was performed for the commercial Ansys Fluent baseline. The Ansys domain was discretized using a comparable mapped mesh hierarchy, maintaining a strict $y^+ \approx 1$ resolution at the walls. In addition to the residual monitors, the bulk velocity was monitored and required to change by less than 1.0×10^{-3} over 250 iterations before accepting a converged solution. The continuity residual criterion was set to 1.0×10^{-5} , while the momentum and SA working-variable residual criteria were set to 1.0×10^{-6} . Mass conservation was also checked from the inlet and outlet mass-flow reports. The properties of these meshes are documented in Table A.4. The wall-normal velocity profiles extracted from the Ansys solver at the same mid-channel location, illustrated in Figure A.2, exhibited identical asymptotic convergence, confirming that the baseline reference data is mathematically robust and completely independent of spatial discretization errors.




FIGURE A.1: FEniCS wall-normal velocity profiles extracted at $x = 0.5L$ for the Coarse, Medium, and Fine grids in the periodic channel flow benchmark. The overlapping profiles demonstrate asymptotic spatial convergence.

TABLE A.4: Wall-resolved channel mesh hierarchy used for the Ansys grid-independence study.

Parameter	Coarse_Ansys	Medium_Ansys	Fine_Ansys
Mesh method	MultiZone	MultiZone	MultiZone
Free face mesh type	All quad	All quad	All quad
Streamwise edge divisions	80	160	240
Wall-normal edge divisions	96	144	192
Wall-normal bias type	Two-sided	Two-sided	Two-sided
Bias factor	47	28	19
Approx. first-layer height Δy_1 [m]	1.700×10^{-3}	1.700×10^{-3}	1.700×10^{-3}
Nodes	10,500	26,448	49,551
Elements	10,296	26,123	49,104
Bulk velocity U_{bulk} [m/s]	18.120	17.267	<i>TBD</i>
Relative change U_{bulk} , Coarse–Medium [%]	–	4.94	–
Relative change U_{bulk} , Medium–Fine [%]	–	–	<i>TBD</i>




FIGURE A.2: Wall-normal velocity profiles extracted at $x = 0.5L$ for the Coarse, Medium, and Fine grids in the Ansys periodic channel flow benchmark, demonstrating spatial convergence of the baseline data.

A.2 Verification Benchmark 2: U-Bend Geometry

A.2.1 Physical and Geometric Parameters

The second verification benchmark evaluates a highly convective U-bend. Because this FEniCS domain is a 2D planar approximation of the symmetry mid-plane of the 3D Ansys VMFL048 case, the 3D physical pipe diameter (D) is explicitly retained as the characteristic length scale. This ensures the Reynolds (Re_D) and Dean (De) numbers remain non-dimensionally consistent with the commercial baseline. The straight inlet and outlet legs are shorter than in the original VMFL048 Ansys mesh to reduce meshing and computational effort; both the FEniCS and Ansys runs use this same short-leg U-bend geometry. The configuration is detailed in Table A.5.

A.2.2 Numerical Simulation Parameters

Due to the adverse pressure gradients and secondary flow separation induced by the streamline curvature, the U-bend represents a significantly more demanding numerical challenge than the channel. Consequently, the FEniCS solver required a smaller pseudo-time step, tighter under-relaxation, and advanced iterative preconditioners to maintain stability, as detailed in Table A.6.

A.2.3 Spatial Verification and Grid Independence

Managing the highly stretched, non-orthogonal inflation layers around the curved arcs of the U-bend required custom mesh generation via Gmsh. A spatial grid independence study was conducted across three refined unstructured grids. The node counts, element densities, and inflation layer parameters for these generated meshes are explicitly documented in Table A.7. Similar to the channel benchmark, the wall-adjacent element height was strictly maintained at 1.20×10^{-5} m to satisfy the low-

TABLE A.5: Geometric, physical, and boundary-condition parameters for the U-bend benchmark.

Parameter	U-bend setting
Benchmark type	Short-leg 2D symmetry-plane U-bend based on Ansys VMFL048
Domain dimensions	Pipe diameter $D = 0.028$ m; centreline bend radius $R_c = 0.125$ m; straight-leg length $30D = 0.840$ m
Characteristic length	Physical pipe diameter $D = 0.028$ m
Reference / inlet velocity	Uniform prescribed inlet velocity $\mathbf{u}_{in} = (0, -1.42)$ m/s
Kinematic viscosity	$\nu = 8.9 \times 10^{-7}$ m ² /s
Reynolds number	$Re_D \approx 4.47 \times 10^4$
Dean number	$De \approx 1.50 \times 10^4$
Body force	$\mathbf{f} = (0, 0)$ m/s ²
Driving condition	Prescribed inlet velocity and zero relative static pressure at the outlet
Velocity boundary conditions	No-slip on all curved and straight pipe walls; natural velocity condition at pressure outlet
Pressure boundary conditions	$p = 0$ Pa at outlet; pressure left free at inlet and walls
SA boundary conditions	$\tilde{\nu}_{in} = 9.34 \times 10^{-5}$ m ² /s at inlet; $\tilde{\nu} = 0$ at walls; natural condition at outlet
Initial fields	$\mathbf{u}_0 = (0, 0)$ m/s, $p_0 = 0$ Pa, $\tilde{\nu}_0 = \tilde{\nu}_{in}$
Verification mesh	Medium_FEniCS; first layer $\Delta y_1 = 1.20 \times 10^{-5}$ m

Reynolds SA wall constraints, while the internal element sizing was systematically refined.

The reported FEniCS static pressure-drop metric is computed from boundary-measure-weighted mean pressures over the inlet and outlet sections, rather than from local point samples. Let Γ_{in} and Γ_{out} denote the inlet and outlet boundaries, and let p_k be the kinematic pressure solved in FEniCS. The physical static pressure drop reported in pascals is

$$\Delta p = \rho (\bar{p}_{k,in} - \bar{p}_{k,out}) = \rho \left[\frac{1}{|\Gamma_{in}|} \int_{\Gamma_{in}} p_k d\Gamma - \frac{1}{|\Gamma_{out}|} \int_{\Gamma_{out}} p_k d\Gamma \right], \quad (\text{A.1})$$

where $|\Gamma| = \int_{\Gamma} 1 d\Gamma$ is the FEniCS boundary measure and $\rho = 997$ kg/m³. In the implementation this corresponds to integrating over the marked inlet and outlet facets with `ds(2)` and `ds(3)`, respectively. The density factor is required because the IPCS formulation stores pressure in kinematic form, while the benchmark table reports the static pressure drop in Pa. Although the outlet pressure is prescribed as zero relative pressure, the outlet term is still evaluated as a boundary average so that the diagnostic follows the same inlet-minus-outlet definition on every grid.

Convergence was visually verified by extracting the velocity profiles directly at the bend exit across the pipe diameter $D = 0.028$ m, a highly sensitive region for curvature-induced separation, as illustrated in Figure A.3. The chosen Medium_FEniCS grid perfectly captured the velocity gradients near the wall and

A. CONFIGURATION AND GRID INDEPENDENCE OF THE SPALART-ALLMARAS VERIFICATION CASES

TABLE A.6: Numerical simulation parameters for the highly convective FEniCS U-bend benchmark.

Parameter	Symbol	Setting
Velocity space	$\mathbf{u}$	Vector CG2
Pressure space	p	Scalar CG1
SA working-variable space	$\tilde{v}$	Scalar CG1
Quadrature degree	–	4
Maximum outer Picard steps	$N_{\text{Picard,max}}$	180
SA solves per Picard step	–	1
Outer velocity tolerance	$\epsilon_{\mathbf{u}}$	1.0×10^{-4}
Outer pressure tolerance	ϵ_p	1.0×10^{-4}
Outer SA tolerance	$\epsilon_{\tilde{v}}$	1.0×10^{-6}
IPCS pseudo-time step	Δt	1.0×10^{-4} s
Maximum IPCS steps per Picard step	–	180
IPCS velocity tolerance	–	1.0×10^{-4}
IPCS pressure tolerance	–	2.0×10^{-3}
Velocity relaxation factor	α_u	0.3
Pressure relaxation factor	α_p	0.1
SA relaxation factor	$\alpha_{\tilde{v}}$	0.01
SA lower bound	$\tilde{v}_{\min}$	1.0×10^{-12}
Wall-distance equation	–	Relaxed Eikonal, $\sigma_w = 0.1$, $G_0 = 20$
Wall-distance Newton settings	–	Relative tolerance 1.0×10^{-8} ; absolute tolerance 1.0×10^{-10} ; maximum 80 iterations; relaxation 0.5
Linear solver strategy	–	IPCS blocks: GMRES with ILU velocity/correction preconditioning and Hypre-AMG pressure preconditioning; SA transport: default solver

TABLE A.7: Wall-resolved U-bend mesh hierarchy used for the FEniCS grid-independence study. The first-layer height is kept fixed to maintain $y^+ \approx 1$, while the bulk element size and inflation-layer resolution are refined.

Parameter	Coarse_FEniCS	Medium_FEniCS	Fine_FEniCS
Bulk element size [m]	2.80×10^{-3}	1.40×10^{-3}	8.00×10^{-4}
First-layer height Δy_1 [m]	1.20×10^{-5}	1.20×10^{-5}	1.20×10^{-5}
Inflation layers	10	14	18
Growth ratio	1.45	1.30	1.25
Inflation thickness [m]	1.07×10^{-3}	1.53×10^{-3}	2.62×10^{-3}
Vertices	23,751	73,978	182,450
Triangular elements	45,998	144,952	359,646
Static pressure drop Δp [Pa]	<i>TBD</i>	<i>TBD</i>	<i>TBD</i>
Relative change Δp , Coarse–Medium [%]	–	<i>TBD</i>	–
Relative change Δp , Medium–Fine [%]	–	–	<i>TBD</i>

provided sufficient isotropic resolution in the bulk flow to resolve the separation bubbles without numerical instability.




FIGURE A.3: Velocity profiles at the U-bend exit, extracted across the pipe diameter $D = 0.028$ m, for the Coarse, Medium, and Fine grids in FEniCS. The profiles demonstrate spatial convergence, justifying the selection of the Medium grid for the final verification.

Similarly, to guarantee the integrity of the cross-code validation in this highly convective regime, an equivalent mesh independence study was executed for the Ansys Fluent baseline. Using a parallel unstructured meshing strategy with strict $y^+ \approx 1$ boundary layer inflation (documented in Table A.8), the Ansys velocity profiles at the bend exit across the same pipe-diameter line, shown in Figure A.4, were verified to be spatially converged. For each Ansys solve, convergence was accepted only after the scaled continuity, momentum, and Spalart–Allmaras residuals reached 1.0×10^{-6} and the monitored static pressure drop changed by less than a relative tolerance of 5.0×10^{-4} over the final 100 iterations. The inlet and outlet mass-flow rates were also compared for each grid to verify the global mass balance before accepting the pressure-drop values. The **Fine_Ansys** grid is therefore selected as the commercial verification baseline, ensuring that the reference data is independent of grid resolution.

A. CONFIGURATION AND GRID INDEPENDENCE OF THE SPALART-ALLMARAS VERIFICATION CASES

TABLE A.8: Wall-resolved U-bend mesh hierarchy used for the Ansys grid-independence study.

Parameter	Coarse_Ansys	Medium_Ansys	Fine_Ansys
Mesh method	Quad-dominant	Quad-dominant	Quad-dominant
Face element size [m]	2.00×10^{-3}	1.25×10^{-3}	8.00×10^{-4}
Curvature angle	12°	8°	5°
Inflation option	First layer	First layer	First layer
First-layer height Δy_1 [m]	1.20×10^{-5}	1.20×10^{-5}	1.20×10^{-5}
Inflation layers	16	22	28
Growth rate	1.22	1.16	1.12
Inflation thickness [m]	1.59×10^{-3}	2.17×10^{-3}	2.29×10^{-3}
Nodes	47,707	106,460	223,089
Elements	46,640	104,763	220,457
Static pressure drop Δp [Pa]	813.350	803.326	797.972
Relative change Δp , Coarse-Medium [%]	—	1.25	—
Relative change Δp , Medium-Fine [%]	—	—	0.67




FIGURE A.4: Velocity profiles at the U-bend exit, extracted across the pipe diameter $D = 0.028$ m, for the Coarse, Medium, and Fine Ansys grids, demonstrating spatial convergence of the commercial baseline.

Appendix B

Configuration and Grid Independence of the Topology Optimization Benchmarks

This appendix collects the detailed setup information for the topology optimization benchmarks used in Chapter 6. The pipe-bend benchmark, U-bend benchmark, and diffuser benchmark are included. For each case, the appendix records the geometry and physical parameters, the FEniCS simulation and optimization parameters, and the Ansys Fluent mesh-independence tables used for sharp-geometry CFD validation.

B.1 Validation and Meshing Strategy

The FEniCS meshes used during topology optimization are treated as fixed optimization discretizations rather than as a separate mesh-convergence hierarchy. This distinction is important because refining the FEniCS mesh would also change the DG_0 design space, the filter stencil, the number of MMA design variables, and the diffuse Brinkman interface location. The pipe-bend optimization uses a rectangular-coordinate triangular mesh with near-wall clustering based on a smooth-wall $y^+ \approx 1$ estimate at the prescribed inlet Reynolds number. The U-bend optimization uses the corresponding wall-resolved triangular Gmsh mesh, with a near-wall size field based on the same $y^+ \approx 1$ estimate and a far-field target size of $h_{\max} = 0.007L$ used to set the density-filter scale. The diffuser benchmark from Yoon [47] uses a wall-resolved rectangular-coordinate triangular mesh over the square diffuser domain, refined around the full-height inlet, the center-third outlet, and the no-slip walls. In all cases, the mesh is chosen to provide sufficient wall-resolved SA and design resolution for the optimization problem; quantitative grid-convergence evidence is instead reported for the extracted sharp geometries in Ansys Fluent.

For each extracted turbulent topology optimization design, the Ansys validation mesh is refined through coarse, medium, and fine resolutions. Both SA and SST $k - \omega$ validation runs use the same extracted geometry, inlet condition, outlet condi-

tion, and fluid properties. The mesh-independence comparison is reported for both turbulence models because near-wall behavior and separated-flow predictions may respond differently to grid refinement.

All Fluent mesh-independence runs in this appendix use the same convergence and acceptance criteria. The static pressure drop must change by less than 5×10^{-4} over the final 100 iterations, while the viscous dissipation must change by less than 10^{-3} over the final 100 iterations. The fine mesh is accepted only when the medium–fine relative difference is below 1% for static pressure drop and below 2% for viscous dissipation for both SA and SST $k - \omega$. The reported refinement differences are computed as $|\phi_{\text{medium}} - \phi_{\text{coarse}}|/|\phi_{\text{coarse}}|$ and $|\phi_{\text{fine}} - \phi_{\text{medium}}|/|\phi_{\text{medium}}|$ for each validation metric.

B.2 Shared Pipe-Bend and U-Bend Topology Optimization Conventions

The pipe-bend and U-bend benchmarks use the same turbulence-topology coupling conventions. Topology-created solids act as walls for the SA working-variable penalty and the reciprocal wall-distance solve. The design density source is therefore **design**, the wall-distance mode is **reciprocal_penalized**, and both the SA and wall-distance penalties use the power-law solid indicator defined in Chapter 4.

The pipe-bend finite-difference sensitivity verification in Chapter 6 is not assigned a separate configuration table. It uses the same geometry, physical parameters, wall-distance convention, penalty settings, objective, filter, and solver setup as the pipe-bend topology optimization run. The only additional settings are the finite-difference sampling parameters listed in Table B.3.

TABLE B.1: Shared wall-distance and SA topology-penalty conventions for the pipe-bend and U-bend benchmarks.

Parameter	Setting
SA wall-distance mode	reciprocal_penalized
SA wall-density source	design
Solid threshold for wall source	$\bar{\gamma} = 0.50$
SA working-variable penalty	Power-law solid indicator with amplitude 10^3 and exponent $n = 4$
Wall-distance penalty	Power-law solid indicator with amplitude 10^3 and exponent $n = 4$
Wall-distance recovery	Reciprocal recovery from the relaxed Eikonal variable G
SA inlet turbulence estimate	$I = 7.5\%$, $\ell/H = 0.05$, eddy-viscosity ratio capped at 50
Optimization objective	Average inlet pressure with fixed outlet pressure $p = 0$

B.3 Pipe Bend Benchmark

The pipe-bend case follows the setup of Bayat et al. [11]. The reference paper uses a standard k - ε model with wall functions; the present thesis uses the wall-resolved SA implementation and evaluates the $V_f = 0.25$ design.

TABLE B.2: Geometric, physical, and boundary-condition parameters for the pipe-bend benchmark.

Parameter	Setting
Reference case	Pipe-bend benchmark from Bayat et al. [11]
Reference figure for setup	Bayat et al. Fig. 10
Reference velocity figure	Bayat et al. Fig. 15
Reference turbulence model	Standard k - ε with wall functions
Turbulence model in this thesis	Spalart-Allmaras
Design-domain dimensions	Square design region $L \times L$ with $L = 1$
Full computational domain	$x/L \in [-0.2, 1]$, $y/L \in [-0.2, 1]$
Non-design inlet duct	Length $0.2L$, height $0.2L$, inlet over $y/L \in [0.7, 0.9]$
Non-design outlet duct	Width $0.2L$, length $0.2L$, outlet over $x/L \in [0.7, 0.9]$ on the lower boundary
Volume fraction	$V_f = 0.25$
Reynolds number	$Re = 5,000$, based on $U_{in}H_{in}\rho/\mu$ with $H_{in} = 0.2L$
Density and viscosity	$\rho = 1.0$, $\mu = 4.0 \times 10^{-5}$
Inlet condition	Uniform inlet velocity $\mathbf{u} = (1, 0)$; SA inlet estimated from the shared turbulence-intensity settings in Table B.1
Outlet condition	Pressure outlet $p = 0$ on the lower outlet, with the normal/-tangential component constraint used by the FEniCS configuration
Wall conditions	No-slip velocity and $\tilde{\nu} = 0$ on external walls; topology-created solids enter through the design-density penalties
Validation metrics	Average inlet pressure, pressure drop, Φ_D , and relative changes between FEniCS-SA, Ansys-SA, and Ansys-SST k - ω

The extracted turbulent pipe-bend geometry is meshed in Ansys Fluent with a patch-conforming triangular bulk mesh and wall-resolved inflation on all no-slip walls. The first-layer height is kept fixed from the smooth-wall $y^+ \approx 1$ estimate, while the mesh hierarchy refines the bulk face size, curvature angle, and inflation distribution.

B. CONFIGURATION AND GRID INDEPENDENCE OF THE TOPOLOGY OPTIMIZATION BENCHMARKS

TABLE B.3: Numerical FEniCS solver, topology-optimization, and sensitivity-verification parameters for the pipe-bend benchmark.

Parameter	Symbol / option	Setting
Flow finite elements	$(\mathbf{u}, p)$	Taylor–Hood P_2/P_1 continuous Galerkin elements
SA working-variable space	$\tilde{\nu}$	P_1 continuous Galerkin
Design-density space	γ	DG_0 cellwise constant density
hh	–	<code>mesh_yplus1.xdmf</code> ; 39,128 vertices, 77,334 triangles, $h_{\max} = 0.007L$
Objective	J_p	Average inlet pressure
Initial design density	γ_0	$V_f = 0.25$, matched to the filtered volume
Brinkman penalty	$\alpha(\bar{\gamma})$	RAMP form with $\alpha_{\text{fluid}} = 0$, $\alpha_{\text{solid}} = 100$
RAMP continuation	q	Code schedule $q = 1/q_\alpha = [1/150, 1/75, 1/30, 1/15]$
Projection continuation	β	$[4, 6, 9, 13]$ with threshold $\eta = 0.50$
MMA move limits	–	$[0.05, 0.04, 0.03, 0.02]$
Maximum optimization iterations	–	Four continuation stages with 25 MMA iterations each; stop if objective change is below 10^{-5} for five iterations
Filter radius	r_f	$4h_{\max} = 0.028L$
Forward flow solver	–	IPCS flow solves for frozen-SA Picard updates, followed by a warm-started SNES Newton line-search final flow solve
Frozen NS-SA coupling	–	Three Picard updates per design iteration with SA relaxation factor 0.20
Linear solver	–	MUMPS for SNES/state and adjoint solves; BiCGSTAB with ILU preconditioning for IPCS warm-start solves
Adjoint treatment	–	Frozen-turbulence adjoint in $(\mathbf{u}^*, p^*)$; $\tilde{\nu}$ and G are frozen during gradient assembly
Sensitivity verification samples	–	Five active DG_0 cells, seed 13
Sensitivity verification step	Δh	10^{-6} , central finite difference
Sensitivity verification fields held fixed	–	SA working variable $\tilde{\nu}^0$ and reciprocal wall-distance field G^0

TABLE B.4: Ansys Fluent mesh-independence table for the extracted turbulent pipe-bend benchmark geometry.

Parameter	Coarse_Ansys	Medium_Ansys	Fine_Ansys
Mesh method	Triangles	Triangles	Triangles
Bulk face element size [m]	1.10×10^{-2}	5.00×10^{-3}	3.25×10^{-3}
Curvature angle	15°	7°	4.5°
Near-wall treatment	Wall-resolved inflation	Wall-resolved inflation	Wall-resolved inflation
First-layer height Δy_1 [m]	6.00×10^{-4}	6.00×10^{-4}	6.00×10^{-4}
Inflation layers	8	14	18
Growth rate	1.15	1.10	1.08
Inflation thickness [m]	8.24×10^{-3}	1.68×10^{-2}	2.25×10^{-2}
Nodes	5,147	21,343	44,042
Elements	7,551	32,780	68,743
Viscous dissipation Φ_D^{SA} [W/m]	0.0143	0.0151	0.0153
Relative change Φ_D^{SA} , Coarse–Medium [%]	–	5.38	–
Relative change Φ_D^{SA} , Medium–Fine [%]	–	–	1.32
Viscous dissipation $\Phi_D^{SST\ k-\omega}$ [W/m]	0.0150	0.0164	0.0166
Relative change $\Phi_D^{SST\ k-\omega}$, Coarse–Medium [%]	–	9.04	–
Relative change $\Phi_D^{SST\ k-\omega}$, Medium–Fine [%]	–	–	1.22
Static pressure drop Δp^{SA} [Pa]	0.1423	0.1449	0.1457
Relative change Δp^{SA} , Coarse–Medium [%]	–	1.84	–
Relative change Δp^{SA} , Medium–Fine [%]	–	–	0.55
Static pressure drop $\Delta p^{SST\ k-\omega}$ [Pa]	0.1481	0.1516	0.1522
Relative change $\Delta p^{SST\ k-\omega}$, Coarse–Medium [%]	–	2.36	–
Relative change $\Delta p^{SST\ k-\omega}$, Medium–Fine [%]	–	–	0.40
Selected validation mesh		Fine_Ansys	

B.4 U-Bend Benchmark

The U-bend case follows Fig. 11 and Table 4 of Bayat et al. [11]. It uses the same turbulent topology optimization framework as the pipe bend, but forces a 180° return path around a fixed internal separator.

B. CONFIGURATION AND GRID INDEPENDENCE OF THE TOPOLOGY OPTIMIZATION BENCHMARKS

TABLE B.5: Geometric, physical, and boundary-condition parameters for the U-bend benchmark.

Parameter	Setting
Reference case	U-bend benchmark from Bayat et al. [11]
Reference figure/table for setup	Bayat et al. Fig. 11 and Table 4
Reference turbulence model	Standard k - ε with wall functions
Turbulence model in this thesis	Spalart-Allmaras
Design-domain dimensions	Square design region $L \times L$ with $L = 1$
Full computational domain	$x/L \in [-0.2, 1]$, $y/L \in [0, 1]$
Port geometry	Two left-side ports of height $0.2L$; inlet over $y/L \in [0.55, 0.75]$, outlet over $y/L \in [0.25, 0.45]$
Fixed separator	Bar thickness $0.10L$, rounded-tip radius $0.05L$, extending from $x/L = -0.2$ to tip at $x/L = 0.5$
Volume fraction	$V_f = 0.27$
Reynolds number	$Re = 5,000$, based on $U_{\text{in}} H_{\text{port}} \rho / \mu$ with $H_{\text{port}} = 0.2L$
Density and viscosity	$\rho = 1.0$, $\mu = 4.0 \times 10^{-5}$
Inlet condition	Uniform inlet velocity $\mathbf{u} = (1, 0)$; SA inlet estimated from the shared turbulence-intensity settings in Table B.1
Outlet condition	Pressure outlet $p = 0$ on the lower left port, with the component constraint used by the FEniCS configuration
Wall conditions	No-slip velocity and $\tilde{\nu} = 0$ on external walls and the fixed separator; topology-created solids enter through the design-density penalties
Validation metrics	Average inlet pressure, pressure drop, Φ_D , and relative changes between FEniCS-SA, Ansys-SA, and Ansys-SST k - ω

TABLE B.6: Numerical FEniCS solver and topology-optimization parameters for the U-bend benchmark.

Parameter	Symbol / option	Setting
Flow finite elements	$(\mathbf{u}, p)$	Taylor-Hood P_2/P_1 continuous Galerkin elements
SA working-variable space	$\tilde{v}$	P_1 continuous Galerkin
Design-density space	γ	DG_0 cellwise constant density
Reference FEniCS mesh	–	<code>mesh_yplus1.xdmf</code> ; wall-resolved triangular Gmsh mesh with 45,280 vertices, 85,780 triangles, near-wall refinement based on $y^+ \approx 1$, and far-field $h_{\max} = 0.007L$
Objective	J_p	Average inlet pressure
Initial design density	γ_0	$V_f = 0.27$, matched to the filtered volume
Brinkman penalty	$\alpha(\bar{\gamma})$	RAMP form with $\alpha_{\text{fluid}} = 0$, $\alpha_{\text{solid}} = 100$
RAMP continuation	q	Code schedule $q = 1/q_\alpha = [1/150, 1/75, 1/35, 1/12]$
Projection continuation	β	$[8, 10, 18, 18]$ with threshold $\eta = 0.50$
MMA move limits	–	$[0.05, 0.04, 0.03, 0.02]$
Maximum optimization iterations	–	Four continuation stages with 25 MMA iterations each; stop if objective change is below 10^{-5} for five iterations
Filter radius	r_f	$4h_{\max} = 0.028L$
Forward flow solver	–	IPCS flow solves for frozen-SA Picard updates, followed by a warm-started SNES Newton line-search final flow solve
Frozen NS-SA coupling	–	Three Picard updates per design iteration with SA relaxation factor 0.20
Linear solver	–	MUMPS for SNES/state and adjoint solves; BiCGSTAB with ILU preconditioning for IPCS warm-start solves
Adjoint treatment	–	Frozen-turbulence adjoint in $(\mathbf{u}^*, p^*)$; $\tilde{v}$ and G are frozen during gradient assembly

B. CONFIGURATION AND GRID INDEPENDENCE OF THE TOPOLOGY OPTIMIZATION BENCHMARKS

TABLE B.7: Ansys Fluent mesh-independence table for the extracted turbulent U-bend benchmark geometry.

Parameter	Coarse	Medium	Fine
Bulk mesh type	[TBD]	[TBD]	[TBD]
Near-wall treatment	[TBD]	[TBD]	[TBD]
First-layer height Δy_1 [m]	[TBD]	[TBD]	[TBD]
Inflation layers	[TBD]	[TBD]	[TBD]
Growth rate	[TBD]	[TBD]	[TBD]
Elements / cells	[TBD]	[TBD]	[TBD]
Maximum y^+ , SA	[TBD]	[TBD]	[TBD]
Maximum y^+ , SST $k - \omega$	[TBD]	[TBD]	[TBD]
Φ_D^{SA} [W/m]	[TBD]	[TBD]	[TBD]
$\Phi_D^{\text{SST } k-\omega}$ [W/m]	[TBD]	[TBD]	[TBD]
Δp^{SA} [Pa]	[TBD]	[TBD]	[TBD]
$\Delta p^{\text{SST } k-\omega}$ [Pa]	[TBD]	[TBD]	[TBD]
Relative change, Medium–Fine [%]		[TBD]	
Selected validation mesh		[TBD]	

B.5 Diffuser Benchmark

The diffuser case follows Fig. 20–22 of Yoon [47]. Unlike the pipe-bend and U-bend benchmarks, the diffuser prescribes parabolic velocity profiles at both the inlet and the outlet and uses the dissipation objective J_D .

TABLE B.8: Geometric, physical, and boundary-condition parameters for the diffuser benchmark.

Parameter	Setting
Reference case	Diffuser benchmark from Yoon [47]
Reference figures for setup and topology	Yoon Fig. 20–22
Reference turbulence model	Spalart-Allmaras
Turbulence model in this thesis	Spalart-Allmaras with frozen-turbulence adjoint
Design-domain dimensions	Square design region $L \times L$ with $L = 1$ m
Full computational domain	$x/L \in [0, 1]$, $y/L \in [0, 1]$
Inlet geometry	Exported Fluent inlet segment from $(1.16 \times 10^{-3}, 1.38 \times 10^{-2})$ to $(1.16 \times 10^{-3}, 0.979)$, with $H_{\text{in}} = 0.96558$ m
Outlet geometry	Exported Fluent outlet segment from $(0.999, 0.336)$ to $(0.999, 0.670)$, with $H_{\text{out}} = 0.33333$ m
Volume fraction	$V_f = 0.30$ fluid, corresponding to Yoon’s 70% solid mass usage
Reynolds number	$Re = 3,000$, based on $U_{\text{max,in}} L \rho / \mu$
Density and viscosity	$\rho = 1000$ kg/m ³ , $\mu = 1.0$ Pa s
Inlet condition	Shifted parabolic velocity profile with $U_{\text{max,in}} = 3$ m/s; $\tilde{\nu} = 1$ for the FEniCS-SA solve
Outlet condition	Shifted parabolic outflow profile with the maximum velocity set by incompressible flow-rate conservation; $\tilde{\nu} = 1$ for the FEniCS-SA solve; pressure pinned at $(0, 0)$ for gauge fixing
Wall conditions	No-slip velocity and $\tilde{\nu} = 0$ on the top and bottom walls and on the right boundary outside the outlet
Brinkman solid resistance	$\alpha_{\text{solid}} = 10^9$, $\alpha_{\text{fluid}} = 0$
Optimization objective	Dissipation objective J_D , including physical viscous dissipation and Brinkman drag over the design domain
Validation metrics	J_D , pressure drop, Φ_D , and relative changes between FEniCS-SA, Ansys-SA, and Ansys-SST $k - \omega$

For the body-fitted Fluent evaluations, the parabolic inlet and outlet velocity profiles are shifted to the extracted inlet and outlet segments. The listed heights are used as the profile heights, while the lower vertices define the local profile origins. The inlet maximum velocity is kept at 3 m/s so that the nominal Reynolds number remains $Re = 3,000$; the outlet maximum velocity is adjusted by incompressible flow-rate conservation.

The extracted turbulent diffuser geometry is meshed with the same wall-resolved triangular strategy as the laminar diffuser comparison in Table C.3, but with a

B. CONFIGURATION AND GRID INDEPENDENCE OF THE TOPOLOGY
OPTIMIZATION BENCHMARKS

TABLE B.9: Numerical FEniCS solver and topology-optimization parameters for the diffuser benchmark.

Parameter	Symbol / option	Setting
Flow finite elements	$(\mathbf{u}, p)$	Taylor-Hood P_2/P_1 continuous Galerkin elements
SA working-variable space	$\tilde{v}$	P_1 continuous Galerkin
Design-density space	γ	DG_0 cellwise constant density
Reference FEniCS mesh	–	<code>mesh_yoon_yplus1.xdmf</code> ; wall-resolved triangular mesh with 38,745 vertices, 76,704 triangles, far-field $h_{\max} = L/180 = 5.56 \times 10^{-3}L$, and first-cell width $1.82 \times 10^{-3}L$
Objective	J_D	Dissipation objective over the design domain
Initial design density	γ_0	$V_f = 0.30$, matched to the filtered volume
Brinkman penalty	$\alpha(\bar{\gamma})$	RAMP form with $\alpha_{\text{fluid}} = 0$, $\alpha_{\text{solid}} = 10^9$
RAMP continuation	q	[0.01, 0.02, 0.04, 0.08]
Projection continuation	β	[1, 2, 4, 8] with threshold $\eta = 0.50$
MMA move limits	–	[0.05, 0.04, 0.03, 0.02]
Maximum optimization iterations	–	Four continuation stages with [50, 60, 80, 100] MMA iterations; stop if objective change is below 10^{-5} for five iterations
Filter radius	r_f	$4h_{\max} = 2.22 \times 10^{-2}L$
SA working-variable penalty	–	Power-law solid indicator with amplitude 10^5 and exponent $n = 3$
Wall-distance penalty	–	Power-law solid indicator with amplitude 10^5 and exponent $n = 3$
Wall-distance equation	–	Reciprocal penalized relaxed Eikonal equation with $\sigma_w = 0.1$ and $G_0 = 20$
Forward flow solver	–	IPCS flow solves for frozen-SA Picard updates, followed by a warm-started SNES Newton line-search final flow solve
Frozen NS-SA coupling	–	Three Picard updates per design iteration with SA relaxation factor 0.15
Linear solver	–	MUMPS for SNES/state and adjoint solves; BiCGSTAB with ILU preconditioning for IPCS warm-start solves
Adjoint treatment	–	Frozen-turbulence adjoint in $(\mathbf{u}^*, p^*)$; $\tilde{v}$ and G are frozen during gradient assembly

distinct coarse–medium–fine hierarchy. The first-layer height is kept fixed at $\Delta y_1 = 1.50 \times 10^{-3}$ m to match the laminar-design diffuser meshes, while the bulk size, curvature angle, and inflation-layer distribution are refined independently for the turbulent topology.

TABLE B.10: Ansys Fluent mesh-independence table for the extracted turbulent diffuser benchmark geometry.

Parameter	Coarse_Ansys	Medium_Ansys	Fine_Ansys
Mesh method	Triangles	Triangles	Triangles
Near-wall treatment	Wall-resolved inflation	Wall-resolved inflation	Wall-resolved inflation
Bulk face element size [m]	7.50×10^{-3}	5.00×10^{-3}	3.50×10^{-3}
Curvature angle	9°	6°	4°
First-layer height Δy_1 [m]	1.50×10^{-3}	1.50×10^{-3}	1.50×10^{-3}
Inflation layers	10	14	18
Growth rate	1.11	1.09	1.07
Inflation thickness [m]	2.51×10^{-2}	3.90×10^{-2}	5.10×10^{-2}
Nodes	8,206	[TBD]	[TBD]
Elements	12,453	[TBD]	[TBD]
Viscous dissipation Φ_D^{SA} [W/m]	[TBD]	[TBD]	[TBD]
Relative change Φ_D^{SA} , Coarse–Medium [%]	–	[TBD]	–
Relative change Φ_D^{SA} , Medium–Fine [%]	–	–	[TBD]
Viscous dissipation $\Phi_D^{\text{SST } k-\omega}$ [W/m]	[TBD]	[TBD]	[TBD]
Relative change $\Phi_D^{\text{SST } k-\omega}$, Coarse–Medium [%]	–	[TBD]	–
Relative change $\Phi_D^{\text{SST } k-\omega}$, Medium–Fine [%]	–	–	[TBD]
Static pressure drop Δp^{SA} [Pa]	[TBD]	[TBD]	[TBD]
Relative change Δp^{SA} , Coarse–Medium [%]	–	[TBD]	–
Relative change Δp^{SA} , Medium–Fine [%]	–	–	[TBD]
Static pressure drop $\Delta p^{\text{SST } k-\omega}$ [Pa]	[TBD]	[TBD]	[TBD]
Relative change $\Delta p^{\text{SST } k-\omega}$, Coarse–Medium [%]	–	[TBD]	–
Relative change $\Delta p^{\text{SST } k-\omega}$, Medium–Fine [%]	–	–	[TBD]
Selected validation mesh	Fine_Ansys, pending metric confirmation		

Appendix C

Laminar-Design Mesh Independence for Benchmark Topologies

This appendix supports the laminar-design comparisons in Chapter 6. It defines the common evaluation pipeline and reports the Ansys mesh-independence studies for the extracted laminar geometries. The topology figures and final laminar-versus-turbulent SST $k - \omega$ metric tables are placed in the corresponding Chapter 6 benchmark sections so that the main comparison can be read together with each optimization result.

C.1 Methodology and Evaluation Pipeline

Each benchmark is solved twice. The laminar reference design is optimized at $Re = 1$, where the optimizer primarily responds to creeping-flow viscous losses. The turbulent design is optimized at the target high-Reynolds operating condition using the SA topology optimization framework described in Chapter 4. Apart from replacing the SA turbulent state equations by the laminar flow formulation and setting $Re = 1$, the laminar topology-optimization runs use the same benchmark-specific FEniCS setup as the turbulent runs: identical mesh, design-domain restrictions, non-design regions, volume constraint, prescribed inlet and outlet treatment, objective type, Brinkman penalty limits, filter radius, projection threshold, continuation schedules, MMA move limits, and maximum MMA iterations. Thus, the pipe-bend and U-bend comparisons retain the inlet-pressure objective J_p , while the diffuser comparison retains the dissipation objective J_D from Yoon [47]. This keeps the laminar-vs-turbulent comparison focused on the flow-physics model rather than on optimization or boundary-condition choices.

Because these FEniCS physical and numerical settings are inherited directly from the turbulent benchmark definitions in Appendix B, they are not repeated as separate configuration tables in this appendix. The additional tables reported

here are therefore limited to the Ansys mesh-independence studies for the extracted laminar geometries.

After optimization, both projected density fields are thresholded and extracted as sharp geometries. The extracted laminar and turbulent designs are then evaluated in Ansys Fluent under the same turbulent operating condition. Because Chapter 6 already validates the extracted turbulent topologies with both Ansys-SA and Ansys-SST $k-\omega$, the laminar-design comparison uses the higher-fidelity SST $k-\omega$ closure. This keeps the comparison focused on topology performance under an independent two-equation turbulence model, rather than repeating the SA cross-code validation for the laminar reference geometries. The key performance metrics are viscous dissipation and pressure drop.

All Fluent mesh-independence studies in this appendix use the same convergence and acceptance criteria as Appendix B, but only for the SST $k-\omega$ model. The static pressure drop must change by less than 5×10^{-4} over the final 100 iterations, while the viscous dissipation must change by less than 10^{-3} over the final 100 iterations. The fine mesh is accepted only when the medium-fine relative difference is below 1% for static pressure drop and below 2% for viscous dissipation. Mass conservation is additionally checked from the inlet and outlet mass-flow reports.

For fairness, the extracted laminar geometries receive their own Ansys mesh-independence checks. The turbulent-design grid studies are reported in Table B.4 of Section B.3, Table B.7 of Section B.4, and Table B.10 of Section B.5 in Appendix B; this appendix adds the corresponding laminar-design mesh tables.

C.2 Pipe Bend Benchmark

The pipe bend tests whether a topology optimized under turbulent conditions produces a more separation-aware 90° turning path than a laminar topology optimized at $Re = 1$. The turbulent case in Chapter 6 uses the $V_f = 0.25$ design. The topology-level comparison and final SST $k-\omega$ performance table are reported in Section 6.3.5.

C.2.1 Laminar-Design Mesh Independence

Table C.1 reports the Ansys mesh-independence study for the extracted laminar pipe-bend design. This table is separate from the turbulent-design mesh study in Table B.4 of Section B.3 because the laminar and turbulent topologies generally produce different wall shapes and near-wall mesh requirements.

C. LAMINAR-DESIGN MESH INDEPENDENCE FOR BENCHMARK TOPOLOGIES

TABLE C.1: Ansys mesh-independence table for the extracted laminar pipe-bend design.

Parameter	Coarse_Ansys	Medium_Ansys	Fine_Ansys
Mesh method	Triangles	Triangles	Triangles
Bulk face element size [m]	4.50×10^{-3}	2.60×10^{-3}	2.00×10^{-3}
Curvature angle	8°	4°	3°
First-layer height Δy_1 [m]	1.20×10^{-3}	1.20×10^{-3}	1.20×10^{-3}
Inflation layers	12	15	16
Growth rate	1.06	1.045	1.04
Inflation thickness [m]	2.02×10^{-2}	2.49×10^{-2}	2.62×10^{-2}
Nodes	22,494	56,468	89,804
Elements	35,948	93,828	153,163
Viscous dissipation $\Phi_D^{\text{SST } k-\omega}$ [W/m]	0.0293	0.0310	0.0313
Relative change $\Phi_D^{\text{SST } k-\omega}$, Coarse–Medium [%]	–	5.82	–
Relative change $\Phi_D^{\text{SST } k-\omega}$, Medium–Fine [%]	–	–	0.95
Static pressure drop $\Delta p^{\text{SST } k-\omega}$ [Pa]	0.5229	0.5420	0.5444
Relative change $\Delta p^{\text{SST } k-\omega}$, Coarse–Medium [%]	–	3.65	–
Relative change $\Delta p^{\text{SST } k-\omega}$, Medium–Fine [%]	–	–	0.45
Selected comparison mesh		Fine_Ansys	

C.3 U-Bend Benchmark

The U-bend tests whether the turbulent optimizer remains advantageous when the flow must complete a 180° return path around a fixed separator. The turbulent case in Chapter 6 uses the $V_f = 0.27$ design. The topology-level comparison and final SST $k - \omega$ performance table are reported in Section 6.4.4.

C.3.1 Laminar-Design Mesh Independence

Table C.2 reports the Ansys mesh-independence study for the extracted laminar U-bend design. The turbulent U-bend mesh study is reported separately in Table B.7 of Section B.4. For the $Re = 5,000$ U-bend operating point, the inlet-port height is $0.2L$ and $U_{\text{in}} = 1 \text{ m s}^{-1}$, giving $\nu = 4.0 \times 10^{-5} \text{ m}^2 \text{ s}^{-1}$. Using the same smooth-wall estimate as the FEniCS mesh generator gives $u_\tau \approx 6.59 \times 10^{-2} \text{ m s}^{-1}$ and a wall distance of $y_+ = 1$ at $6.07 \times 10^{-4} \text{ m}$. The Ansys mesh hierarchy therefore keeps the first layer fixed at $6.00 \times 10^{-4} \text{ m}$ and refines the bulk triangular mesh, curvature resolution, and inflation-layer distribution around the separator tip and 180° return wall.

TABLE C.2: Ansys mesh-independence table for the extracted laminar U-bend design.

Parameter	Coarse_Ansys	Medium_Ansys	Fine_Ansys
Mesh method	Triangles	Triangles	Triangles
Bulk face element size [m]	6.00×10^{-3}	4.00×10^{-3}	2.75×10^{-3}
Curvature angle	8°	6°	4°
First-layer height Δy_1 [m]	6.00×10^{-4}	6.00×10^{-4}	6.00×10^{-4}
Inflation layers	12	16	20
Growth rate	1.10	1.08	1.06
Inflation thickness [m]	1.28×10^{-2}	1.82×10^{-2}	2.21×10^{-2}
Nodes	16,106	32,580	62,622
Elements	24,460	50,067	98,084
Viscous dissipation $\Phi_D^{\text{SST } k-\omega}$ [W/m]	0.101	0.105	0.107
Relative change $\Phi_D^{\text{SST } k-\omega}$, Coarse–Medium [%]	–	3.96	–
Relative change $\Phi_D^{\text{SST } k-\omega}$, Medium–Fine [%]	–	–	1.90
Static pressure drop $\Delta p^{\text{SST } k-\omega}$ [Pa]	1.125	1.141	1.150
Relative change $\Delta p^{\text{SST } k-\omega}$, Coarse–Medium [%]	–	1.42	–
Relative change $\Delta p^{\text{SST } k-\omega}$, Medium–Fine [%]	–	–	0.79
Selected comparison mesh		Fine_Ansys	

C.4 Diffuser Benchmark

The diffuser benchmark tests whether the turbulent optimizer remains advantageous when the design must diffuse a full-height inflow into a center-third outflow while minimizing the dissipation objective J_D . The turbulent case in Chapter 6 uses the $V_f = 0.30$, $Re = 3,000$ design from Yoon [47]. The topology-level comparison and final SST $k - \omega$ performance table are reported in Section 6.5.4.

C.4.1 Laminar-Design Mesh Independence

Table C.3 reports the Ansys mesh-independence study for the extracted laminar diffuser design. The turbulent diffuser mesh study is reported separately in Table B.10 of Section B.5. The Fluent solves use the same shifted inlet and outlet parabolic velocity-profile convention as the turbulent diffuser validation setup.

C. LAMINAR-DESIGN MESH INDEPENDENCE FOR BENCHMARK TOPOLOGIES

TABLE C.3: Ansys mesh-independence table for the extracted laminar diffuser design.

Parameter	Coarse_Ansys	Medium_Ansys	Fine_Ansys
Mesh method	Triangles	Triangles	Triangles
Near-wall treatment	Wall-resolved inflation	Wall-resolved inflation	Wall-resolved inflation
Bulk face element size [m]	1.20×10^{-2}	5.50×10^{-3}	3.75×10^{-3}
Curvature angle	15°	7°	5°
First-layer height Δy_1 [m]	1.50×10^{-3}	1.50×10^{-3}	1.50×10^{-3}
Inflation layers	6	12	16
Growth rate	1.15	1.10	1.08
Inflation thickness [m]	1.31×10^{-2}	3.19×10^{-2}	4.55×10^{-2}
Nodes	3,300	13,490	25,414
Elements	5,074	21,036	39,561
Viscous dissipation $\Phi_D^{\text{SST } k-\omega}$ [W/m]	1.1426×10^4	1.1749×10^4	1.1831×10^4
Relative change $\Phi_D^{\text{SST } k-\omega}$, Coarse–Medium [%]	–	2.83	–
Relative change $\Phi_D^{\text{SST } k-\omega}$, Medium–Fine [%]	–	–	0.70
Static pressure drop $\Delta p^{\text{SST } k-\omega}$ [Pa]	3.4547×10^4	3.4513×10^4	3.4299×10^4
Relative change $\Delta p^{\text{SST } k-\omega}$, Coarse–Medium [%]	–	0.10	–
Relative change $\Delta p^{\text{SST } k-\omega}$, Medium–Fine [%]	–	–	0.62
Selected comparison mesh		Fine_Ansys	

Bibliography

- [1] A. A. Abuhatira, S. M. Salim, and J. B. Vorstius. Numerical simulation of turbulent pipe flow with 90-degree elbow using wall y^+ approach. In *Volume 10: Fluids Engineering*, ASME 2021 International Mechanical Engineering Congress and Exposition. American Society of Mechanical Engineers, 11 2021.
- [2] J. Alexandersen, N. Aage, C. S. Andreasen, and O. Sigmund. Topology optimisation for natural convection problems. *International Journal for Numerical Methods in Fluids*, 76:699–721, 2014.
- [3] J. Alexandersen and C. S. Andreasen. A review of topology optimisation for fluid-based problems. *Fluids*, 5(1):29, 2020.
- [4] M. S. Alnæs, J. Blechta, J. Hake, A. Johansson, B. Kehlet, A. Logg, C. Richardson, J. Ring, M. E. Rognes, and G. N. Wells. The fenics project version 1.5. *Archive of Numerical Software*, 3(100), 2015.
- [5] D. H. Alonso, L. F. Garcia Rodriguez, and E. C. N. Silva. Flexible framework for fluid topology optimization with openfoam and finite element-based high-level discrete adjoint method (fenics/dolfin-adjoint). *Structural and Multidisciplinary Optimization*, 64(6):4409–4440, 2021.
- [6] Altair Engineering, Inc. *AcuSolve Turbulence Modeling Training Manual*, 2023. Available online: <https://help.altair.com/hwcfdsolvers/acusolve>.
- [7] ANSYS, Inc., Canonsburg, PA. *Ansys Fluid Dynamics Verification Manual, Release 2024 R1*, January 2024.
- [8] R. K. Apalowo. Numerical study of different models for turbulent flow in 90° pipe bend. *U.Porto Journal of Engineering*, 8(2):121–134, 2022.
- [9] D. N. Arnold, F. Brezzi, B. Cockburn, and L. D. Marini. Unified analysis of discontinuous galerkin methods for elliptic problems. *SIAM Journal on Numerical Analysis*, 39(5):1749–1779, 2002.
- [10] B. Barroca, M. Blommaert, G. Buckinx, M. Schevenels, and M. Baelmans. High-resolution topology optimization of heat sinks for microelectronics and benchmark against conventional designs. In *17th International Heat Transfer Conference, IHTC-17*, Cape Town, 2023.

- [11] A. Bayat, H. Li, and J. Alexandersen. Density-based topology optimization for turbulent fluid flow using the standard $k - \varepsilon$ RANS model with wall-functions imposed through an implicit wall penalty formulation, Jan. 2026.
- [12] M. P. Bendsøe and N. Kikuchi. Generating optimal topologies in structural design using a homogenization method. *Computer methods in applied mechanics and engineering*, 71(2):197–224, 1988.
- [13] T. Borrvall and J. Petersson. Topology optimization of fluids in stokes flow. *International journal for numerical methods in fluids*, 41(1):77–107, 2003.
- [14] A. J. Chorin. Numerical solution of the navier-stokes equations. *Mathematics of computation*, 22(104):745–762, 1968.
- [15] Y. Deng, Z. Liu, and Y. Wu. Topology optimization of steady and unsteady incompressible navier-stokes flows driven by body forces. *Structural and Multidisciplinary Optimization*, 47:555–570, 2013.
- [16] C. B. Dilgen, S. B. Dilgen, D. R. Fuhrman, O. Sigmund, and B. S. Lazarov. Topology optimization of turbulent flows. *Computer Methods in Applied Mechanics and Engineering*, 331:363–393, 2018.
- [17] M. M. Enayet, M. M. Gibson, A. M. K. P. Taylor, and M. Yianneskis. Laser-doppler measurements of laminar and turbulent flow in a pipe bend. *International Journal of Heat and Fluid Flow*, 3(4):213–219, 1982.
- [18] A. Fawaz, Y. Hua, S. Le Corre, Y. Fan, and L. Luo. Topology optimization of heat exchangers: A review. *Energy*, 258:124053, 2022.
- [19] J. H. Ferziger, M. Perić, and R. L. Street. *Computational Methods for Fluid Dynamics*. Springer, Berlin, Heidelberg, 4th edition, 2012.
- [20] A. Gersborg-Hansen, M. P. Bendsøe, and O. Sigmund. Topology optimization of channel flow problems. *Structural and multidisciplinary optimization*, 30(3):181–192, 2005.
- [21] C. Geuzaine and J.-F. Remacle. Gmsh: A 3-d finite element mesh generator with built-in pre-and post-processing facilities. *International journal for numerical methods in engineering*, 79(11):1309–1331, 2009.
- [22] J. K. Guest, J. H. Prévost, and T. Belytschko. Achieving minimum length scale in topology optimization using nodal independent density variables and projection functions. *International journal for numerical methods in engineering*, 61(2):238–254, 2004.
- [23] J. S. Hesthaven and T. Warburton. *Nodal Discontinuous Galerkin Methods: Algorithms, Analysis, and Applications*. Springer Science & Business Media, New York, NY, 2007.

-
- [24] L. C. Høghøj, D. R. Nørhave, J. Alexandersen, O. Sigmund, and C. S. Andreasen. Topology optimization of two fluid heat exchangers. *International Journal of Heat and Mass Transfer*, 163:120543, 2020.
- [25] H. Kobayashi, K. Yaji, S. Yamasaki, and K. Fujita. Topology design of two-fluid heat exchange. *Structural and Multidisciplinary Optimization*, 63:821–834, 2021.
- [26] H. P. Langtangen and A. Logg. *Solving PDEs in Python: The FEniCS Tutorial I*. Springer Nature, 2016.
- [27] B. S. Lazarov and O. Sigmund. Filters in topology optimization based on helmholtz-type differential equations. *International Journal for Numerical Methods in Engineering*, 86(6):765–781, 2011.
- [28] A. Logg, K.-A. Mardal, and G. N. Wells. *Automated Solution of Differential Equations by the Finite Element Method: The FEniCS Book*, volume 84. Springer Science & Business Media, 2012.
- [29] L. Noël and K. Maute. Xfem level set-based topology optimization for turbulent conjugate heat transfer problems. *Structural and Multidisciplinary Optimization*, 66(1), 2023.
- [30] C. Othmer. Adjoint methods for car aerodynamics. *Journal of Mathematics in Industry*, 4(1):1–25, 2014.
- [31] E. M. Papoutsis-Kiachagias and K. C. Giannakoglou. Continuous adjoint methods for turbulent flows, applied to shape and topology optimization: industrial applications. *Archives of Computational Methods in Engineering*, 23(2):255–299, 2016.
- [32] S. B. Pope. *Turbulent Flows*. Cambridge University Press, 2000.
- [33] C. L. Rumsey et al. Turbulence modeling resource. NASA Langley Research Center, <https://turbmodels.larc.nasa.gov>, 2023. Accessed: 2024-05-15.
- [34] M. Schramm, B. Stoevesandt, and J. Peinke. Optimization of airfoils using the adjoint approach and the influence of adjoint turbulent viscosity. *Computation*, 6(1):10, 2018.
- [35] P. Spalart and S. Allmaras. A one-equation turbulence model for aerodynamic flows. *30th aerospace sciences meeting and exhibit*, page 439, 1992.
- [36] P. R. Spalart and S. R. Allmaras. A one-equation turbulence model for aerodynamic flows. *La Recherche Aéronautique*, 1:5–21, 1994.
- [37] M. Stolpe and K. Svanberg. An alternative interpolation scheme for minimum compliance topology optimization. *Structural and Multidisciplinary Optimization*, 22(2):116–124, 2001.

- [38] Y. Sun, M. Hao, and Z. Wang. Topology optimization of turbulent flow cooling structures based on the k - ε model. *Entropy*, 25:1061, 2023.
- [39] K. Svanberg. The method of moving asymptotes—a new method for structural optimization. *International journal for numerical methods in engineering*, 24(2):359–373, 1987.
- [40] T. Takamasa and M. Kondo. Measurement of gas-liquid two-phase flow in a c-shaped pipe. *Nuclear Engineering and Design*, 154(2):147–156, 1995.
- [41] R. Temam. Sur l’approximation de la solution des équations de navier-stokes par la méthode des pas fractionnaires (i). *Archive for Rational Mechanics and Analysis*, 32(2):135–153, 1969.
- [42] P. Tucker. Differential equation-based wall distance computation for des and rans. *Journal of Computational Physics*, 190(1):229–248, 2003.
- [43] H. K. Versteeg and W. Malalasekera. *An introduction to computational fluid dynamics: the finite volume method*. Pearson education, 2007.
- [44] F. Wang, B. S. Lazarov, and O. Sigmund. On projection methods, convergence and robust formulations in topology optimization. *Structural and Multidisciplinary Optimization*, 43(6):767–784, 2011.
- [45] F. M. White and J. Majdalani. *Viscous fluid flow*, volume 3. McGraw-Hill New York, 2006.
- [46] H. Wu, D. Wang, and X. Huang. Effect of constant eddy viscosity assumption on optimization using a discrete adjoint method. *Chinese Journal of Aeronautics*, 36(11):102–118, 2023.
- [47] G. H. Yoon. Topology optimization for turbulent flow with spalart-allmaras model. *Computer Methods in Applied Mechanics and Engineering*, 303:288–311, 2016.
- [48] J. Zhao, M. Zhang, Y. Zhu, R. Cheng, and L. Wang. Topology optimization of turbulent forced convective heat sinks using a multi-layer thermofluid model. *Structural and Multidisciplinary Optimization*, 64(6):3835–3859, 2021.
- [49] X. Zhao, M. Zhou, O. Sigmund, and C. S. Andreasen. A “poor man’s approach” to topology optimization of cooling channels based on a darcy flow model. *International Journal of Heat and Mass Transfer*, 116:1108–1123, 2018.
- [50] T. Zhou, T. Liu, Y. Deng, L. Chen, S. Qian, and Z. Liu. Design of microfluidic channel networks with specified output flow rates using the cfd-based optimization method. *Microfluidics and Nanofluidics*, 21(1), 2017.
- [51] A. Zymaris, D. Papadimitriou, K. Giannakoglou, and C. Othmer. Continuous adjoint approach to the spalart–allmaras turbulence model for incompressible flows. *Computers & Fluids*, 38(8):1528–1538, 2009.

Code of conduct and transparency statement on the use of GenAI for KU Leuven-students (academic year 2025–2026)

Generative AI (GenAI) assistance tools can be used to generate various types of content, including text, images, code, video, music, or combinations thereof. Common examples of such tools include ChatGPT, Google Gemini, Microsoft Copilot, Midjourney, Claude.ai, Perplexity.ai, and DALL·E, among others.

This code of conduct is a tool that helps students to be transparent about the use of GenAI and fits within the university's principles on academic integrity.

Important guidelines and remarks

- **Sensitive or personal data:** Some GenAI tools protect your input and sensitive or personal data better than others. There is often no transparency on what the owners of the AI applications do with the data entered. Therefore, **do not enter sensitive or personal data in free GenAI tools**. More info about what sensitive and personal data are, is described in the [three confidentiality levels of the KU Leuven data classification model](#) (non-confidential, confidential, strictly confidential). For confidential data you should preferably use **M365 Copilot Chat with your KU Leuven account**. If you use another GenAI tool, you must be **absolutely certain** it does **not store or reuse the data you enter**. **Try to avoid entering these data. Do so only if strictly necessary, and only to the extent required. For personal data, work with anonymous or pseudonymized data.** Additionally, in case of strictly confidential data this should first be discussed with the teaching staff of the course or your thesis supervisor.
- **Copyrighted materials:** For lawfully obtained copyrighted material you should preferably use **M365 Copilot Chat with your KU Leuven**

account. If you use another GenAI tool, you must be **absolutely certain** it does **not store or reuse the data you enter**.

- GenAI assistance may not be used for data or topics covered by a **Non-Disclosure Agreement (NDA)**. Even in the absence of an NDA, certain information may still need to be treated as confidential—for example, due to regulatory requirements or the risk of significant harm to the university if disclosed. In case of doubt, check with your teaching staff or supervisor.
- If your master’s thesis is under **embargo**, you should first discuss with your supervisor whether the use of GenAI is permitted.
- Before using a GenAI tool, always consider whether its use is responsible, including from a **sustainability perspective** (e.g. when using GenAI as a search engine, language assistant, ...).
- **Take a scientific and critical attitude** when interacting with GenAI assistance and interpreting its output, that may not always be correct.
- As a student you are responsible for complying with Article 84 of the Regulations on Education and Examinations: your report or thesis should reflect your own knowledge, understanding and skills. Be aware that plagiarism rules also apply to (work that is the result of) the use of GenAI assistance tools.

Exam Regulations Article 84: *“Every conduct individual students display with which they (partially) inhibit or attempt to inhibit a correct judgement of their own knowledge, understanding and/or skills or those of other students, is considered an irregularity which may result in a suitable penalty. A special type of irregularity is plagiarism, i.e. copying the work (ideas, texts, structures, designs, images, plans, codes , ...) of others or prior personal work in an exact or slightly modified way without adequately acknowledging the sources. Every possession of prohibited resources during an examination (see article 65) is considered an irregularity.”*

- In order to maintain academic integrity and avoid plagiarism, **more information about being transparent on the use of GenAI assistance and about correctly citing and referencing GenAI** can be found on this website for students (**Dutch/English**).
- **Additional reading: KU Leuven guidelines on responsible use of Generative AI tools, and other information** (**Dutch/English**)

A few final words

If you are uncertain whether or not you should declare your use of GenAI tools, we suggest that you discuss this with your instructor or supervisor. It is always safer to declare GenAI use, even when it is not strictly required. However, declaring GenAI use does not entail that its use is

allowed; the right column in the table below provides more detailed instructions in this regard (code of conduct).

Moreover, advanced AI tools are evolving rapidly, and their capabilities have expanded significantly in a short period of time. As a result, we do not yet have all the answers about their responsible use. Finally, it is important to follow-up on the most recent evolutions in AI technologies, to have an open mind but also to be a bit cautious, to communicate with instructors, teaching assistants, supervisors and peers, to be as transparent as we can, and to learn together as we move along.

Student name:

Student number:

Please indicate with "X" whether it relates to a course assignment or to the Bachelor's or Master's thesis:

☐ This form is related to a **course assignment**.

Course name: **Course number:**

This form is related to my ☐ **Bachelor's** or ☐ **Master's thesis**.

Title Bachelor's or Master's thesis: **Supervisor:**

Daily supervisor:

Please indicate with "X":

☐ I **did not use** any GenAI assistance tool.

☐ I **did use** GenAI Assistance. In this case **specify which ones** (e.g. ChatGPT, M365 Copilot, ...):

GenAI assistance used as/for:	Name of the GenAI tool(s) used If helpful, also describe in which way you were using GenAI related to what is specified as code of conduct.	Code of conduct: For each of the categories below, always take into account the important guidelines and remarks mentioned above (e.g. copyrighted data, sensitive or personal data,...).
As a language assistant for reviewing or improving texts I wrote myself	<i>add text here</i>	This use is similar to using spelling and grammar check tools. In general, you do not have to refer to such kind of GenAI use in the text. However, be careful: When using GenAI tools on texts you did not write yourself to improve the text, you have to refer to the original source or author, otherwise you are committing plagiarism and thus an irregularity.
As a paraphrasing tool	<i>add text here</i>	This use is allowed except when it is prohibited by the teacher or the program of study. You may paraphrase your own text or texts by an author other than yourself and take inspiration from what a GenAI tool or another tool suggests (unless it is not allowed). In general, you do not have to refer to such kind of GenAI use or other paraphrasing tools in the text. However, be careful: If it entails text by an author other than yourself, you are not allowed to include that paraphrased text without reference to the original source or author. Without such reference, you would be committing plagiarism and thus an irregularity.

For translation aid to improve texts I wrote myself or to better understand text from others	<i>add text here</i>	This use is allowed except when it is prohibited by the teacher or the program of study. It is similar to using translation tools (Google translate, DeepL, ...). In general, you do not have to refer to such kind of GenAI use in the text. However, be careful: • You are not allowed to include that translated text without reference to the original source or author. Without such reference, you would be committing plagiarism and thus an irregularity. • Always check the translated text for correctness and meaning.
As a search engine to get information on a topic or to search for existing research on the topic	<i>add text here</i>	This use is similar to e.g. a Google search or checking Wikipedia. If you write your own text based on this information, you do not have to refer to the use of GenAI in the text. You only have to refer to the existing research and references you have checked and used (without such references you would be committing plagiarism and thus committing an irregularity). However, be careful: • Be aware that the output of the GenAI tool cannot be guaranteed as a 100% reliable source of information. The output may not be entirely correct and/or be limited due to the databases it uses. Moreover, knowledge evolves and may change over time; therefore, the database of the GenAI tool may not be up to date. Therefore, verify the information and do not just copy-paste it as you should understand and critically process everything you are writing.

For literature search	<i>add text here</i>	This use is comparable to e.g. a Google Scholar search. You do not have to refer to such kind of GenAI use; you only have to refer to the literature references you have checked and used (without such references you would be committing plagiarism and thus committing an irregularity). However, be careful: • Be aware that the search output is restricted to the database the GenAI tool is built on. After this initial search, look for scientific sources and conduct your own analysis of the source documents. Interpret, analyse and process the information you obtained; verify it and do not just copy-paste it as you should understand and critically process everything you are writing. • Be aware that some GenAI tools may output no or wrong references. As a student you are responsible for further checking and verifying the absence or correctness of references; do not just copy-paste it.
For generating programming code	<i>add text here</i>	Use of GenAI for coding is allowed except when it is prohibited by the teacher or the program of study. If used for coding, correctly mention the use of GenAI assistance in accordance with the instructions on the page on being transparent about the use of GenAI.

For new ideas generating (research)	<i>add text here</i>	This use of GenAI is allowed except when it is prohibited by the teacher or the program of study. Correctly mention the use of GenAI assistance in accordance with the instructions on the page on being transparent about the use of GenAI. Be careful: • Further verify in this case whether the idea is novel or not. It is likely that it is related to existing work. If so, that existing work should be correctly referenced in the text (without such reference you would be committing plagiarism and thus committing an irregularity).
For generating synthetic data	<i>add text here</i>	Use of GenAI for generating synthetic data is allowed, provided that it is methodologically and ethically justifiable, except when it is prohibited by the teacher or the program of study. Always correctly mention the use of GenAI assistance in accordance with the instructions on the page on being transparent about the use of GenAI. Be careful: • Always carefully evaluate the generated synthetic data for quality and possible bias since the output is highly dependent on the quality of the data on which the models are trained.

For generating blocks of text (other than the allowed use without referencing mentioned above)		According to Article 84 of the Regulations on Education and Examinations your text should allow to correctly and properly assess your own knowledge, understanding and skills. Therefore, inserting blocks of text without quotes and a reference to GenAI assistance in your work is not allowed. Be careful: • If it is really needed to insert a block of text from a GenAI tool, for instance because of the nature of your assignment, mention it as a citation by using quotes and correctly mention the use of GenAI assistance in accordance with the instructions page on being transparent about the use of GenAI. • However, in general, such GenAI use should be kept to an absolute minimum; you should always check the original sources.
For generating visuals, video or audio	<i>add text here</i>	Use of GenAI for generating visuals, audio or video is allowed except when it is prohibited by the teacher or the program of study. Be careful: • If used, refer to GenAI for visuals according to the usual referencing style, following the instructions on the page on referencing GenAI. • If you work with existing visuals, audio or video, that existing work should be correctly referenced in the text (without such reference you would be committing plagiarism and thus an irregularity). • Explain the usage in the methods section (if there is one) and optionally attach (or link to) the prompts with the full output (or history).

Other use (specify here; this may also include a combination of types of use mentioned above):	<i>add text here</i>	To make sure other use of GenAI is allowed within the course or thesis, and if so, the conditions that may apply, contact the teaching staff of the course or the supervisor of the thesis beforehand and explain the intended GenAI use. Also inform the programme director. Motivate how you would comply with Article 84 of the exam regulations. Explain the use and the added value of the AI tool you consider to use and how it is in accordance with the assignment or thesis and the KU Leuven guidelines on responsible use of Generative AI tools. Depending on the kind of GenAI use, it may be needed to properly reference it in the text, in accordance with the instructions page on being transparent about the use of GenAI.
---	----------------------	--